

The Coherence-Delay Trade-off in Tracking Under Drift

A Cube-Root Law for Finite-Memory Estimation

Vinicius Parreira 

Independent Researcher

`figueiredo@pm.me`

May 2026

Abstract

This paper develops two linked claims about tracking in non-stationary environments. First, finite-memory tracking under drift obeys a structural coherence-delay trade-off. When observations are aggregated over time, estimation variance decreases with memory length, but delay-induced bias increases as older data becomes less representative of the present environment. We formalize this trade-off and show that, in the finite-memory setting considered here, it induces a **cube-root scaling law**: the **minimum achievable tracking error** scales as $\zeta^{1/3}$ with drift rate ζ , with an **optimal effective memory** that scales as $\zeta^{-2/3}$. Empirically, we recover slopes of 0.316 for sliding windows and 0.317 for EWMA, closely matching the theoretical prediction.

Second, in action-coupled systems, low observed error can be maintained by forcing the environment toward a stale model rather than by genuine adaptation. We call this coercive masking and introduce an **effort-aware coherence diagnostic** to separate genuine adaptation from action-induced alignment. On passive benchmarks, generic changepoint detectors remain the stronger default alarms. On active synthetic and logged recommendation benchmarks, however, the effort-aware diagnostic exposes control-induced masking that passive detectors do not encode under the same calibration, identifying masking onset with median delay 5.0 steps in the particle benchmark and improving KuaiRand bubble detection from 0.458 to 0.692. These results show that tracking under drift is constrained by temporal staleness, and that evaluating adaptive systems under feedback requires accounting for intervention effort rather than relying on raw fit alone.

Keywords: tracking under drift; finite-memory estimation; Wasserstein distance; Sinkhorn divergence; concept drift; online compression; coherence score; effort-aware coherence; coercive masking.

MSC 2020: 93D05, 93D30, 68P99.

Contents

1	Introduction	3
2	Mathematical Definitions	4
2.1	State Spaces and Information State	4
2.2	ISS Definitions	5
3	Why Action Enters the Update in Closed-Loop Tracking	5
3.1	External Characterization: Drift-Coherence	5
3.2	Belief Update With Feedback	6
4	A Minimal Stability Template for Closed-Loop Tracking	7
4.1	Hypotheses	7
4.2	Main Stability Template	7
5	Diagnostic Failure Modes	8
6	Coherence Diagnostics	9
6.1	Estimable Score	10
6.2	Architecture and Behavior Are Not Equivalent	15
7	Empirical Validation	15
7.1	Formalization	16
7.2	Experiment A: Component Ablation	16
7.3	Experiment B: Sinkhorn Sample Complexity and the Coherence-Delay U-Curve	19
7.4	Experiment C: ELEC2 Early-Warning Lead Time	21
7.5	Experiment D: Dynamic n_t^* Optimization on ELEC2	23
7.6	Experiment E: Flagship Real-Data Masking on KuaiRand	25
7.7	Experiment F: Feedback Masking in the Particle Tracker	26
7.8	Experiment G: Bikes Early-Warning Lead Time	32
7.9	Passive and Active Diagnostics Are Complementary	33
7.10	Discussion: Limitations of the Empirical Validation	33
8	Relation to Algorithmic Regulation	34
8.1	Coercive Masking	34
9	Related Work	36
10	Limitations and Open Directions	37
11	Conclusion	38
A	Code Repository	42
B	KuaiRand Logged Benchmark Protocol	42

1 Introduction

Modern deployed models operate in environments that are not merely noisy but continuously changing. In recommender systems, adaptive control, and streaming prediction, the system must estimate a time-varying state of the world from finite data while continuing to act on that same world. Those actions can change future observations, so the system faces a joint problem of inference, staleness, and feedback. This paper focuses on the two resulting questions: what finite memory can achieve under drift, and how to detect when apparent coherence is being maintained by intervention rather than by adaptation.

A finite-memory estimator benefits from using more observations, because more data reduces statistical uncertainty. But larger windows also make the estimate older. Under drift, that temporal lag produces model staleness: the data used to build the estimate becomes progressively less representative of the present environment. The central question is not whether one can estimate accurately in a stationary regime, but how estimation accuracy is limited once memory and non-stationarity are both unavoidable.

This tension admits a simple structure. The total tracking error decomposes into a variance term that falls with window size and a lag-bias term that rises with the drift rate and the effective age of the data. Optimizing this trade-off yields a cube-root law for the finite-memory estimator class analyzed here: $\mathcal{E}_{\min} = \frac{3}{2}C_K^{2/3}\zeta^{1/3}$ with $n^* = (C_K/\zeta)^{2/3}$. This law can be read as a limit on *online compression of a non-stationary process*: memory reduces variance, but it also extends the interval over which a compressed model must remain valid while the world keeps moving.

This finite-memory baseline is also the right reference point for evaluating closed-loop diagnostics. In action-coupled systems, lower observed error need not mean better estimation, because interventions can reduce error by steering the environment toward a stale model. The coherence score $\mathcal{C}$ measures the apparent state-update bottleneck, while $\mathcal{C}^E$ is the main diagnostic for the active masking regime because it discounts alignment purchased through intervention rather than adaptation.

In closed-loop settings, raw tracking error is not always enough. A system may appear well aligned with its environment because its actions are pushing the environment toward a stale model rather than because the model remains current. This is the phenomenon of *feedback masking*. In the synthetic control regime studied below, the same effect is called coercive masking. A useful decomposition has three functional roles: a representation of the world ($\mathcal{P}$), an actuation rule ($\mathcal{A}$), and an adaptation operator (Φ). It is motivated by action-effective tracking systems rather than by a claim about information in general.

The decomposition yields two diagnostics. The first is the coherence score $\mathcal{C}$, a bottleneck-sensitive summary that combines model fidelity, actuation reproducibility, and adaptation quality. The second is an effort-corrected version, $\mathcal{C}^E$, designed for regimes where low observed error may be purchased through intervention rather than through adaptation. The distinction separates degradation from its cause and controls for action-induced confounding.

Section 6 derives the coherence-delay trade-off and the finite-memory cube-root law. Section 3 shows that action-effective environments require the update to depend on the current state, the most recent action, and the incoming signal, which motivates the decomposition $(\mathcal{P}, \mathcal{A}, \Phi)$. Sections 6 and 8 introduces $\mathcal{C}$ and the

effort-aware $\mathcal{C}^E$. Section 7 validates both claims: passive benchmarks test the finite-memory baseline against generic drift alarms, while active and logged benchmarks test whether effort-aware coherence can surface coercive masking.

The cube-root law holds for the finite-memory averaging class analyzed here. The decomposition is motivated only for *action-effective* systems; passive observers are a genuine dyadic boundary case. The stability results remain conditional on explicit ISS hypotheses, with the Gaussian tracker serving as the constructive reference instance, and the diagnostics describe failure modes rather than every possible adaptive architecture.

2 Mathematical Definitions

2.1 State Spaces and Information State

Definition 2.1 (State space). Let S be a non-empty measurable space (the *signal space*). The *information state space* $\mathcal{I}$ is the space of all probability measures on S equipped with a metric $d_{\mathcal{I}}$ compatible with weak convergence. Unless otherwise specified, $d_{\mathcal{I}}$ denotes the L^2 -Wasserstein distance W_2 .

Definition 2.2 (Potentiality). The *Potentiality* at time t is a measurable map

$$\mathcal{P}(S, t) \in \mathcal{I},$$

representing the system’s current probability measure over possible world states. $\mathcal{P}$ is ε -coherent at time t if $W_2(\mathcal{P}(S, t), \mathcal{P}^*(S, t)) \leq \varepsilon$, where $\mathcal{P}^*(S, t)$ denotes the true marginal distribution of the environment at time t .

Remark 2.3 (Computational realization of W_2). The L^2 -Wasserstein distance W_2 is chosen as the metric on $\mathcal{I}$ for its geometric faithfulness: unlike f -divergences such as KL divergence, W_2 inherits the geometry of the underlying signal space S and behaves well under spatial translations of the distribution support, a desirable property when concept drift shifts the support of $\mathcal{P}^*$. In high-dimensional signal spaces, exact computation of W_2 costs $O(n^3 \log n)$ via network-flow solvers. For streaming applications requiring real-time Φ_t updates, the entropic regularization of Cuturi [2013] replaces W_2 with the Sinkhorn divergence S_ε , computable in $O(n^2/\varepsilon^2)$ via parallelizable matrix-scaling iterations [Altschuler et al., 2017]. S_ε converges to W_2^2 as $\varepsilon \rightarrow 0$ and breaks the curse of dimensionality of raw W_2 at the cost of a bias proportional to ε [Genevay et al., 2019]. The theoretical results hold for W_2 ; their computational realization in streaming contexts is demonstrated in the Gaussian-tracker experiments (section 7; code at appendix A).

Definition 2.4 (Actuation). The *Actuation operator* is a measurable map

$$\mathcal{A}_t: \mathcal{I} \rightarrow \mathcal{O},$$

where $\mathcal{O}$ is the output space (distinct from lowercase o , which denotes a specific output element $o \in \mathcal{O}$). $\mathcal{A}_t$ is δ -reproducible if for any $p \in \mathcal{I}$ and any two executions ω, ω' , $d_{\mathcal{O}}(\mathcal{A}_t(p; \omega), \mathcal{A}_t(p; \omega')) \leq \delta$. A deterministic actuation is 0-reproducible.

Definition 2.5 (Convergence operator). The *Convergence operator* is a measurable map

$$\Phi_t: \mathcal{I} \times \mathcal{O} \times S \rightarrow \mathcal{I},$$

mapping the current information state, the most recent actuation, and an incoming signal to an updated information state. Φ_t is *coherence-restoring* if there exist $\beta \in \mathcal{KL}$ and $\gamma \in \mathcal{K}$ such that for all $t \geq 0$,

$$W_2(\mathcal{P}(S, t), \mathcal{P}^*(S, t)) \leq \beta(W_2(\mathcal{P}(S, 0), \mathcal{P}^*(S, 0)), t) + \gamma\left(\sup_{s \leq t} \|u(s)\|\right),$$

where $u(s)$ represents the environmental drift input at time s .

Definition 2.6 (Three-Role System). An information system $I(S, t)$ is a *three-role system* if it is equipped with all three operators $(\mathcal{P}, \mathcal{A}, \Phi)$ satisfying definitions 2.2, 2.4 and 2.5 respectively. The composite information state is

$$I(S, t) := (\mathcal{P}(S, t), \mathcal{A}_t, \Phi_t) \in \mathcal{I} \times \mathcal{F}_O \times \mathcal{F}_I.$$

2.2 ISS Definitions

Definition 2.7 (Class $\mathcal{K}$, $\mathcal{KL}$). A function $\gamma: \mathbb{R}_{\geq 0} \rightarrow \mathbb{R}_{\geq 0}$ is class $\mathcal{K}$ if it is continuous, strictly increasing, and $\gamma(0) = 0$. A function $\beta: \mathbb{R}_{\geq 0}^2 \rightarrow \mathbb{R}_{\geq 0}$ is class $\mathcal{KL}$ if $\beta(\cdot, t) \in \mathcal{K}$ for each t and $\beta(r, \cdot) \rightarrow 0$ as $t \rightarrow \infty$ for each $r > 0$.

Definition 2.8 (Input-to-State Stability). A dynamical system $\dot{x} = f(x, u)$ is *input-to-state stable* (ISS) if there exist $\beta \in \mathcal{KL}$ and $\gamma \in \mathcal{K}$ such that for every initial state $x(0)$, every essentially bounded input u , and all $t \geq 0$,

$$\|x(t)\| \leq \beta(\|x(0)\|, t) + \gamma(\|u\|_\infty).$$

3 Why Action Enters the Update in Closed-Loop Tracking

The coherence-delay law in section 6 concerns estimation with finite memory. In open-loop settings, the estimator only needs to combine past state information and new observations. In closed-loop settings, however, the system's own actions alter future observations. The update therefore needs to keep track of what the system did, not only of what it saw.

3.1 External Characterization: Drift-Coherence

Definition 3.1 (Drift-Coherence). An information system $I(S, t)$ is *drift-coherent* if there exist $\beta \in \mathcal{KL}$ and $\gamma \in \mathcal{K}$ such that for all $t \geq 0$ and all environmental inputs with $\text{ess sup}_{s \leq t} \|u(s)\| \leq M < \infty$:

$$W_2(\mathcal{P}(S, t), \mathcal{P}^*(S, t)) \leq \beta(W_2(\mathcal{P}(S, 0), \mathcal{P}^*(S, 0)), t) + \gamma(M).$$

The system maintains bounded tracking error under bounded drift.

Definition 3.2 (Action-Effectiveness). The actuation operator $\mathcal{A}_t$ is *action-effective* in environment $\mathcal{X}$ if there exist distinct outputs $o, o' \in \mathcal{O}$ and a state $s \in S$ such that

$$D_{\text{TV}}\left(P(\cdot \mid s_t = s, o) \parallel P(\cdot \mid s_t = s, o')\right) > 0.$$

If this never occurs, the system is *passive*: its actions do not affect the environmental transition law.

3.2 Belief Update With Feedback

Lemma 3.3 (Belief Update Lemma). *Let $I(S, t)$ be a drift-coherent, action-effective information system. Any update map U whose repeated application preserves drift-coherence in this class must depend on three arguments:*

- (i) *the current internal state $p = \mathcal{P}(S, t)$,*
- (ii) *the most recent action $o = \mathcal{A}_t(\mathcal{P})$, and*
- (iii) *the incoming environmental signal s_{t+1} .*

Each argument cannot be dropped without leaving the class of drift-coherent updates considered here.

Proof. If U omits p , the update becomes memoryless and discards the sufficient summary of the past; under temporally structured drift this prevents bounded tracking error. If U omits o , then in an action-effective environment the update cannot distinguish exogenous drift from changes induced by the system’s own actuation. This is precisely the credit-assignment problem that appears in POMDP filtering and belief update semantics [Kaelbling et al., 1998, Katsuno and Mendelzon, 1992, Bonanno, 2024]. If U omits s_{t+1} , no new environmental evidence enters the state estimate and the gap between $\mathcal{P}$ and $\mathcal{P}^*$ grows with drift. In each case, the drift-coherence bound fails. $\square$

Remark 3.4 (Passive observer systems). When actions do not influence transitions, the update can drop the actuation argument and the system is genuinely dyadic. The three-role decomposition is therefore motivated only for action-effective settings.

Corollary 3.5 (Ternary Update Structure). *For action-effective tracking systems, any drift-coherent update map has the form*

$$U: \mathcal{I} \times \mathcal{O} \times S \rightarrow \mathcal{I}.$$

Any ostensibly smaller description either hides one of these arguments inside a state variable or fails drift-coherence.

Proof. This is immediate from lemma 3.3. $\square$

Theorem 3.6 (Minimal Functional Roles). *Let $I(S, t)$ be a drift-coherent, action-effective information system. Any functional decomposition that preserves corollary 3.5 in this class must realize three roles: a stored state, an action channel, and an adaptation rule that maps the state, action, and new observation to an updated state. The tuple $(\mathcal{P}, \mathcal{A}, \Phi)$ is the explicit decomposition used here.*

Proof. By corollary 3.5, the update requires all three arguments (p, o, s_{t+1}) . A decomposition that omits any one of the corresponding roles must either re-encode it implicitly inside another component or lose the update structure required for drift-coherence. The explicit realization used here is $\mathcal{P}$ for stored state, $\mathcal{A}$ for the action channel, and Φ for the update map. $\square$

4 A Minimal Stability Template for Closed-Loop Tracking

The necessity argument in section 3 identifies the three inputs a closed-loop update must use. To turn that decomposition into a usable diagnostic, we need a stability template under which each component admits a coherence score. We use an ISS-based cascade model for that purpose.

4.1 Hypotheses

The following assumptions are explicit and application-dependent.

- (H-ISS-P) The representation subsystem admits an ISS-Lyapunov function $V_P: \mathcal{I} \rightarrow \mathbb{R}_{\geq 0}$.
- (H-ISS-A) The actuation subsystem admits an ISS-Lyapunov function V_A .
- (H-ISS- Φ) The adaptation subsystem admits an ISS-Lyapunov function V_Φ .
- (H-ISS-Env) The environmental drift input is essentially bounded and evolves slowly enough for Φ_t to remain contractive on the induced state trajectory.

Remark 4.1 (Measure-valued state spaces). When $\mathcal{P}(S, t)$ is measure-valued, the natural potential function is $V_P = \frac{1}{2}W_2(\mathcal{P}(S, t), \mathcal{P}^*(S, t))^2$. In Gaussian instances this reduces to an ordinary quadratic tracking error; in more general settings the construction follows the Wasserstein gradient-flow viewpoint [Ambrosio et al., 2005, Mironchenko and Prieur, 2020].

4.2 Main Stability Template

Theorem 4.2 (Conditional Borromean Stability Template). *Under hypotheses (H-ISS-P), (H-ISS-A), (H-ISS- Φ), and (H-ISS-Env), the composite system $I(S, t) = (\mathcal{P}, \mathcal{A}, \Phi)$ with feedback structure*

$$\Phi_t \xrightarrow{\text{updates}} \mathcal{P}(S, t) \xrightarrow{\text{drives}} \mathcal{A}_t \xrightarrow{\text{feeds back}} \Phi_t$$

admits the composite Lyapunov functional $V_{\text{total}} := V_P + V_A + V_\Phi$, which is non-increasing in expectation along trajectories for bounded environmental inputs. Removing any one of the three roles destroys this guarantee in general.

Proof. The cascade $(\Phi \rightarrow \mathcal{P} \rightarrow \mathcal{A})$ is ISS by repeated application of standard cascade arguments [Sontag, 1989, 1995, Dashkovskiy et al., 2010]. The subsystem Lyapunov functions combine into a composite functional whose cross terms are absorbed by the ISS gains. If any one of the three roles is removed, at least one of

the required input channels disappears: without Φ the state is not corrected, without $\mathcal{P}$ the action has no bounded state input, and without $\mathcal{A}$ the feedback loop to adaptation is broken. In each case the monotone decrease of V_{total} is no longer guaranteed. $\square$

Remark 4.3 (Scope of the stability template). Theorem 4.2 is a conditional stability template, not a universal theorem about all three-role architectures. Its force is constructive in cases where the subsystem Lyapunov functions can be written down and their gains checked directly. In this paper that verification is carried out only for the Gaussian tracker. For non-linear or black-box systems, the template should be read as a checklist of sufficient conditions rather than as an already verified property of the architecture itself.

Example 4.4 (Concrete Gaussian instantiation). Let $\mathcal{P}^*(S, t) \sim \mathcal{N}(\mu_t, \Sigma)$ with

$$\mu_t = \mu_0 + \int_0^t \eta(s) ds, \quad \|\eta\|_\infty < M.$$

Let $\mathcal{P}(S, t) = \mathcal{N}(\hat{\mu}_t, \Sigma)$ be the Kalman estimate, $\mathcal{A}_t(\mathcal{P}) = \hat{\mu}_t$, and Φ_t the Kalman update. Then $V_P = W_2(\mathcal{P}, \mathcal{P}^*)^2 = \|\hat{\mu}_t - \mu_t\|^2$, together with the standard actuation and innovation terms, yields a concrete instance satisfying the ISS template [Jazwinski, 1970]. This Gaussian case is the paper’s constructive reference instance for theorem 4.2.

5 Diagnostic Failure Modes

The decomposition $(\mathcal{P}, \mathcal{A}, \Phi)$ is useful only if it identifies failure regimes that matter empirically. This section records the three diagnostic failure modes used throughout the experiments and the coherence scores.

Proposition 5.1 (Failure Mode 1: Static- $\mathcal{P}$ collapse). *Suppose Φ_t ceases to update $\mathcal{P}$, so $\mathcal{P}(S, t) = \mathcal{P}(S, t_0)$ for all $t > t_0$. If the environment undergoes distributional shift $\mathcal{P}^*(S, t) \neq \mathcal{P}^*(S, t_0)$ for large t , then $W_2(\mathcal{P}(S, t), \mathcal{P}^*(S, t)) \rightarrow \infty$ and the expected actuation error $\mathbb{E}[d_{\mathcal{O}}(\mathcal{A}_t(\mathcal{P}), \mathcal{A}_t^*)]$ is unbounded.*

Proof. $\mathcal{P}(S, t) = \mathcal{P}(S, t_0)$ is fixed while $\mathcal{P}^*$ drifts, so $W_2 \rightarrow \infty$ by definition. Lipschitz continuity of $\mathcal{A}_t$ then propagates the divergence to actuation error. $\square$

Remark 5.2 (Interpretation). FM-1 is the stale-model regime: the representation no longer tracks the world, so downstream actions inherit an increasingly obsolete state estimate.

Proposition 5.3 (Failure Mode 2: Non-deterministic- $\mathcal{A}$ divergence). *If $\mathcal{A}_t$ is not δ -reproducible for any $\delta < \infty$, then $H(\mathcal{A}_t \mid \mathcal{P}(S, t)) > 0$ for all t : actuation carries information not determined by the state model. Distributed replicas of the system diverge, and the composite resists auditing and replication.*

Proof. Non- δ -reproducibility implies $\exists \omega \neq \omega'$ with $\mathcal{A}_t(p; \omega) \neq \mathcal{A}_t(p; \omega')$, giving positive conditional entropy. Auditability requires reconstructing actuation from state; this fails when conditional entropy is positive. $\square$

Remark 5.4 (Interpretation). FM-2 is the non-reproducible-action regime: the model may still be accurate, but the realized action can no longer be reconstructed from the represented state, so replicas and audits diverge.

Proposition 5.5 (Failure Mode 3: Absent- Φ monopolization). *Let $\Phi_t \equiv \text{id}$ (no update). In any closed-loop deployment, actuation feedback shifts $\mathcal{P}^*$; since $\mathcal{P}$ cannot track, the distributional mismatch grows at a rate lower-bounded by the rate of environmental response to actuation.*

Proof. The closed-loop case follows from $\mathcal{P}^*(t+1) \leftarrow f(\mathcal{P}^*(t), \mathcal{A}_t(\mathcal{P}(t)))$ with $\mathcal{P}$ fixed, so the gap $W_2(\mathcal{P}, \mathcal{P}^*)$ grows at least as fast as the variation of f in its second argument whenever $\mathcal{A}_t(\mathcal{P}) \neq \mathcal{A}_t^*$. $\square$

Remark 5.6 (Interpretation). FM-3 is the missing-adaptation regime. In passive settings it appears as a failure to track external drift; in closed-loop settings it can amplify the system’s own feedback and produce self-reinforcing misalignment.

6 Coherence Diagnostics

The coherence-delay law constrains the representation term of the finite-memory averaging systems studied here, but closed-loop performance also depends on whether actions remain reproducible and whether adaptation remains effective. The ISS template of section 4 provides one way to score those three ingredients. This section packages them into a diagnostic built from the minimal decomposition $(\mathcal{P}, \mathcal{A}, \Phi)$ used throughout the paper.

Definition 6.1 (Component Coherence Scores). Let $I(S, t) = (\mathcal{P}, \mathcal{A}, \Phi)$ be an information system with ISS-Lyapunov functionals V_P, V_A, V_Φ as in section 4. Define the three *component coherence scores* at time t as:

$$\sigma_P(t) := \frac{1}{1 + V_P(t)} = \frac{1}{1 + \frac{1}{2}W_2(\mathcal{P}(S, t), \mathcal{P}^*(S, t))^2}, \quad (1)$$

$$\sigma_A(t) := \exp(-H(\mathcal{A}_t \mid \mathcal{P}(S, t))), \quad (2)$$

$$\sigma_\Phi(t) := \frac{1}{1 + V_\Phi(t)}. \quad (3)$$

Each score takes values in $(0, 1]$, equals 1 when the corresponding component is at its reference state, and approaches 0 as the component degrades.

Definition 6.2 (Coherence Score). The *coherence score* of a system $I(S, t)$ at time t is

$$\mathcal{C}(I(S, t), t) := \min(\sigma_P(t), \sigma_A(t), \sigma_\Phi(t)) \in (0, 1].$$

The min aggregation is the correct choice given the Borromean property (theorem 4.2): system coherence is bounded by the weakest component, and no strength in two components compensates for failure in the third. An arithmetic mean would obscure localized failures; the min preserves the diagnostic identity of the bottleneck.

Remark 6.3 (Bottleneck sensitivity and effort correction). The use of $\min(\sigma_P, \sigma_A, \sigma_\Phi)$ is bottleneck-sensitive: it is a diagnostic of the weakest active subsystem, not a smooth summary of overall system quality. Aggregators such as arithmetic or geometric means may be reported as descriptive summaries, but they are not valid replacements for the structural diagnostic given by definition 6.2.

In action-effective environments, however, high actuation capacity can keep σ_P artificially high by forcing the world toward a stale model. When such feedback masking is possible, a conservative extension is the *effort-corrected score*

$$\mathcal{C}^E(I(S, t), t) := \min(\sigma_P(t)e^{-\lambda E_t/E_0}, \sigma_A(t), \sigma_\Phi(t)),$$

where E_t is any monotone proxy for corrective actuation burden, $E_0 > 0$ is a reference scale, and $\lambda > 0$ controls sensitivity. It discounts apparent improvements in model-world alignment when those gains come from action-induced confounding rather than better state estimation. In the synthetic particle tracker, E_t is the realized coercive effort $\alpha|a_t - s_t^{\text{free}}|$; in deployed systems, action increments, energy expenditure, or control-surface usage can play the same role. For logged recommendation benchmarks, the preferred instantiation is the per-user KL divergence between the current exposure-tag distribution and the healthy random-exposure baseline distribution, aggregated sequentially over time and then batched across users.

Remark 6.4 (Finite-memory interpretation). The scaling law behind proposition 6.9 is not an artefact of the sliding-window estimator alone. The two terms in the bound, $n^{-1/2}$ and $\zeta\tau$, encode the structural trade-off between variance reduction and delay-induced bias for the finite-memory averaging family studied here: the first term reflects the statistical limit of estimating a distribution from finitely many samples under subgaussian noise, and the second reflects the mismatch induced by drift when the estimator necessarily looks at stale information.

Concretely, the paper proves the cube-root scaling for the canonical local averaging schemes used throughout the experiments, namely sliding windows and exponentially weighted averages. The same exponent reappears because these estimators expose the same variance-vs.-lag balance under drift. We do not claim here a universal minimax theorem over arbitrary adaptive estimators; a broader lower bound for all causal finite-memory kernels remains an open strengthening.

Within the averaging class studied here, the $\zeta^{1/3}$ law should be read as a robust operational regime under the stated regularity assumptions rather than as a peculiarity of a single windowing scheme.

6.1 Estimable Score

Definition 6.1 requires $\mathcal{P}^*(S, t)$, the true oracle distribution of the environment, to compute $\sigma_P(t)$. In any deployed system, $\mathcal{P}^*$ is unobservable by construction: a system with complete knowledge of $\mathcal{P}^*$ would have nothing to estimate. We therefore introduce an estimable version of σ_P using an empirical approximation of $\mathcal{P}^*$.

Definition 6.5 (Empirical Potentiality Estimator). Let $\{s_{t-n+1}, \dots, s_t\}$ be the n most recent environmental signals observed by the system. The *empirical potential-*

ity estimator is the empirical measure:

$$\hat{\mathcal{P}}^*(S, t) := \frac{1}{n} \sum_{k=t-n+1}^t \delta_{s_k},$$

where δ_{s_k} is the Dirac measure at observation s_k .

Definition 6.6 (Estimable Score). The *debiased Sinkhorn divergence* between measures μ, ν is

$$S_\varepsilon(\mu, \nu) := T_\varepsilon(\mu, \nu) - \frac{1}{2}T_\varepsilon(\mu, \mu) - \frac{1}{2}T_\varepsilon(\nu, \nu),$$

where T_ε is the raw entropically regularized optimal transport cost. Unlike T_ε , the debiased divergence satisfies $S_\varepsilon(\mu, \mu) = 0$, restoring the metric property and enabling the dimension-independent convergence rate of proposition 6.9.

The *estimable component* $\hat{\sigma}_P$ replaces $\mathcal{P}^*$ with $\hat{\mathcal{P}}^*$ in eq. (1) using S_ε :

$$\hat{\sigma}_P(t) := \frac{1}{1 + \frac{1}{2}S_\varepsilon\left(\mathcal{P}(S, t), \hat{\mathcal{P}}^*(S, t)\right)^2}.$$

The *estimable score* is

$$\hat{\mathcal{C}}(I(S, t), t) := \min(\hat{\sigma}_P(t), \sigma_A(t), \sigma_\Phi(t)).$$

The three components have different operational meanings. In the Gaussian tracker, σ_A is exact because the action channel is deterministic. In the particle tracker, σ_A is a controlled synthetic surrogate based on the actuation gap. In KuaiRand, σ_A is a logged-data surrogate based on normalized exposure entropy, available only because the benchmark exposes the necessary exposure counts. Likewise, $\sigma_\Phi = 1/(1 + V_\Phi)$ is exact only when the convergence subsystem's Lyapunov functional is available or constructively specified.

We therefore treat action observability as a three-level ladder: ‘exact’ when the action channel is deterministic or constructively known, ‘identified’ when the channel is logged richly enough to estimate the target quantity with a direct estimator, and ‘proxy’ when the benchmark only supports an operational surrogate. The Gaussian tracker sits at the exact level, the synthetic particle benchmark sits at the identified/surrogate boundary, and KuaiRand sits at the proxy level.

Assumption 6.7 (W_2 -Lipschitz Drift). The oracle trajectory $t \mapsto \mathcal{P}^*(S, t)$ is W_2 -Lipschitz with drift rate $\zeta \geq 0$: for all $s, t \geq 0$,

$$W_2(\mathcal{P}^*(S, s), \mathcal{P}^*(S, t)) \leq \zeta |t - s|.$$

Remark 6.8 (Verification in the Gaussian tracker). For the Gaussian family $\mathcal{P}^*(S, t) = \mathcal{N}(\mu^*(t), \Sigma)$ with common covariance Σ , the Wasserstein distance reduces to translation of the mean:

$$W_2(\mathcal{N}(\mu_s, \Sigma), \mathcal{N}(\mu_t, \Sigma)) = \|\mu_s - \mu_t\|_2.$$

Hence the Gaussian tracker satisfies assumption 6.7 whenever $\|\dot{\mu}^*\|_\infty \leq \zeta$.

Proposition 6.9 (Estimation Error Bound for $\hat{\sigma}_P$). *Let assumption 6.7 hold, and assume each oracle marginal $\mathcal{P}^*(S, t)$ is K -subgaussian. Let $\hat{\mathcal{P}}^*$ be the sliding-window empirical estimator with window n (definition 6.5), and let $\tau = (n - 1)/2$ be the effective lag of the centered window. Then under the entropic regularization of remark 2.3,*

$$\mathbb{E} |\sigma_P(t) - \hat{\sigma}_P(t)| \leq C_K n^{-1/2} + \zeta \tau,$$

where C_K absorbs the finite-sample Sinkhorn constant, the fixed regularization ε , and the uniform subgaussian envelope. This is a finite-memory upper bound for the sliding-window class, not a minimax statement over all adaptive estimators.

Proof. Let

$$\bar{\mathcal{P}}^*(S, t) := \frac{1}{n} \sum_{k=t-n+1}^t \mathcal{P}^*(S, k)$$

be the oracle window-average law. The finite-sample term is the usual subgaussian Sinkhorn estimation error: by Genevay et al. [2019, Thm. 3], the debiased divergence between $\bar{\mathcal{P}}^*(S, t)$ and its empirical estimator is of order $n^{-1/2}$, with the regularization dependence absorbed into C_K .

The additional error comes from temporal staleness of the window. By convexity of W_2^2 in its second argument,

$$W_2(\mathcal{P}^*(S, t), \bar{\mathcal{P}}^*(S, t)) \leq \frac{1}{n} \sum_{k=t-n+1}^t W_2(\mathcal{P}^*(S, t), \mathcal{P}^*(S, k)).$$

Applying assumption 6.7,

$$\begin{aligned} W_2(\mathcal{P}^*(S, t), \bar{\mathcal{P}}^*(S, t)) &\leq \frac{\zeta}{n} \sum_{k=t-n+1}^t |t - k| \\ &= \zeta \frac{n-1}{2} = \zeta \tau. \end{aligned}$$

This is the explicit lag-bias term. Combining it with the $C_K n^{-1/2}$ statistical term yields the stated decomposition. Finally, the score map $x \mapsto 1/(1 + \frac{1}{2}x^2)$ is globally Lipschitz on $\mathbb{R}_{\geq 0}$, so the same order transfers to $\mathbb{E} |\sigma_P(t) - \hat{\sigma}_P(t)|$. $\square$

Remark 6.10 (Tightness of the lag-bias term). The lag term is not merely an artifact of upper-bounding. For monotone translation trajectories $\mathcal{P}^*(S, t) = \mathcal{P}_0(S - vt)$ with $\|v\| = \zeta$, the current law and the centered window-average law are separated by exactly the mean age of the window, so the temporal bias is of order $\zeta \tau$. The cube-root balance therefore reflects a real finite-memory limitation rather than a proof-only looseness.

proposition 6.9 grounds $\hat{\mathcal{C}}$: the estimation error decreases as the window n grows and increases proportionally to drift rate ζ . The required window size for target error η is $n = O(C_{K,\varepsilon}^2/\eta^2)$, independent of dimension.

This characterization also yields the tracking limit used throughout the paper. When window lag is accounted for explicitly ($\tau = n/2$ for a centred window), the total estimation error decomposes as a sum of two opposing terms. Minimising over n yields the cube-root optimum for the averaging class, including the sliding-window and exponentially weighted schemes used in the paper.

Definition 6.11 (Lag-Inclusive Estimation Error). For a sliding window of size n centred at lag $\tau = n/2$, the total expected estimation error is:

$$\mathcal{E}(n) = C_K n^{-1/2} + \frac{1}{2} \zeta n,$$

where the first term is the *variance* (decreasing in n) and the second is the *lag-bias* (increasing in n).

Proposition 6.12 (Optimal Window Under Drift). *Under definition 6.11, the estimation error $\mathcal{E}(n)$ is minimised at the optimal window*

$$n^* = \left(\frac{C_K}{\zeta} \right)^{2/3},$$

yielding the minimum achievable tracking error

$$\mathcal{E}_{\min} = \frac{3}{2} C_K^{2/3} \zeta^{1/3}.$$

Proof. Setting $d\mathcal{E}/dn = 0$: $-\frac{1}{2}C_K n^{-3/2} + \frac{1}{2}\zeta = 0$, giving $n^* = (C_K/\zeta)^{2/3}$. Substituting back: $\mathcal{E}(n^*) = C_K(C_K/\zeta)^{-1/3} + \frac{1}{2}\zeta(C_K/\zeta)^{2/3} = C_K^{2/3}\zeta^{1/3} + \frac{1}{2}C_K^{2/3}\zeta^{1/3} = \frac{3}{2}C_K^{2/3}\zeta^{1/3}$. $\square$

Corollary 6.13 (EWMA Analogue). *Let the finite-memory estimator use exponential weights $w_j = \alpha(1-\alpha)^j$ for $j \geq 0$, so its effective lag is $\tau_\alpha = \sum_{j \geq 0} j w_j = (1-\alpha)/\alpha$ and its effective sample size is $n_{\text{eff}}(\alpha) = 1/\sum_{j \geq 0} w_j^2 = (2-\alpha)/\alpha$. Under assumption 6.7, the resulting finite-memory error obeys*

$$\mathcal{E}_{\text{EWMA}}(\alpha) \leq C_K \sqrt{\frac{\alpha}{2-\alpha}} + \zeta \frac{1-\alpha}{\alpha}.$$

Hence, in the small- α regime, $\alpha^* \propto (\zeta/C_K)^{2/3}$ and $\mathcal{E}_{\text{EWMA},\min} \propto C_K^{2/3} \zeta^{1/3}$.

Proof. The variance term is the sliding-window rate with n replaced by the effective sample size $n_{\text{eff}}(\alpha)$. The lag term is the W_2 -Lipschitz drift rate multiplied by the effective age τ_α . For small α this becomes $\mathcal{E}_{\text{EWMA}}(\alpha) = O(C_K \alpha^{1/2} + \zeta \alpha^{-1})$. Balancing the two exponents gives $\alpha^* = O((\zeta/C_K)^{2/3})$, and substitution yields the same $\zeta^{1/3}$ minimum-error law as the sliding-window case. $\square$

Definition 6.14 (Causal Averaging Kernel). A *causal averaging kernel* is a sequence of weights $w = (w_0, w_1, \dots)$ such that $w_j \geq 0$ and $\sum_{j \geq 0} w_j = 1$. Its effective size and effective lag are

$$n_{\text{eff}}(w) := \frac{1}{\sum_{j \geq 0} w_j^2}, \quad \tau_{\text{eff}}(w) := \sum_{j \geq 0} j w_j.$$

Proposition 6.15 (Kernel-Level Lower Bound). *Let $X_t = \mu_t + \xi_t$ with $\xi_t \sim \mathcal{N}(0, \sigma^2)$ i.i.d. and let the drift trajectory be the monotone translation model $\mu_t = \zeta t$ with $\zeta \geq 0$. For any causal averaging kernel w as in definition 6.14, the kernel estimator*

$$\hat{\mu}_t^{(w)} := \sum_{j \geq 0} w_j X_{t-j}$$

satisfies

$$\sqrt{\mathbb{E} (\hat{\mu}_t^{(w)} - \mu_t)^2} \geq c_1 \sigma n_{\text{eff}}(w)^{-1/2} \vee c_2 \zeta \tau_{\text{eff}}(w)$$

for universal constants $c_1, c_2 > 0$. Equivalently, the mean-squared error obeys the exact decomposition

$$\mathbb{E} (\hat{\mu}_t^{(w)} - \mu_t)^2 = \sigma^2 n_{\text{eff}}(w)^{-1} + \zeta^2 \tau_{\text{eff}}(w)^2,$$

and therefore the RMSE lower bound above. In particular, optimizing over the averaging family yields the same cube-root exponent as proposition 6.12.

Proof. The estimator error decomposes as

$$\hat{\mu}_t^{(w)} - \mu_t = \underbrace{\zeta \sum_{j \geq 0} j w_j}_{\text{bias}} + \underbrace{\sum_{j \geq 0} w_j \xi_{t-j}}_{\text{noise}}.$$

Hence the bias magnitude is exactly $\zeta \tau_{\text{eff}}(w)$ and the noise term is Gaussian with variance $\sigma^2 \sum_{j \geq 0} w_j^2 = \sigma^2 / n_{\text{eff}}(w)$. Therefore

$$\mathbb{E} (\hat{\mu}_t^{(w)} - \mu_t)^2 = \zeta^2 \tau_{\text{eff}}(w)^2 + \sigma^2 n_{\text{eff}}(w)^{-1},$$

and taking square roots gives the stated lower bound, since $\sqrt{a+b} \geq a^{1/2} \vee b^{1/2}$ for nonnegative a, b . Optimizing the resulting trade-off over the averaging family recovers the same 2/3 window scaling and 1/3 error exponent as the upper bound. $\square$

Remark 6.16 (Interpretation: a finite-memory floor). Proposition 6.12 establishes that perfect coherence ($V_P = 0$) is *impossible* under continuous drift for the finite-memory averaging class analyzed here. Every drift-coherent system must carry a residual tracking error at least $\mathcal{E}_{\min} \propto \zeta^{1/3}$. This can be read as a finite-memory limitation: the two sources of error (variance and lag) trade against each other with no joint optimum of zero. The cube-root scaling has a concrete engineering consequence: a fifty-fold increase in drift rate ζ raises the minimum error floor by only $50^{1/3} \approx 3.7\times$, while simultaneously forcing the optimal window n^* to contract by $50^{2/3} \approx 14\times$. The optimal window $n^* = (C_K/\zeta)^{2/3}$ provides a *self-tuning rule*: given a real-time estimate $\hat{\zeta}_t$ of the local drift rate, the system should maintain window size $n_t^* = (C_K/\hat{\zeta}_t)^{2/3}$ to operate at the minimum of the error U-curve at every step. This turns the coherence-delay law into a practical self-tuning rule for the windowed estimator class analyzed here.

For small relative calibration errors, the rule is predictably sublinear:

$$\frac{\Delta n_t^*}{n_t^*} \approx \frac{2}{3} \left(\frac{\Delta C_K}{C_K} - \frac{\Delta \hat{\zeta}_t}{\hat{\zeta}_t} \right), \quad \frac{\Delta \mathcal{E}_{\min}}{\mathcal{E}_{\min}} \approx \frac{2}{3} \frac{\Delta C_K}{C_K} + \frac{1}{3} \frac{\Delta \hat{\zeta}_t}{\hat{\zeta}_t}.$$

Thus moderate multiplicative misspecification in $\hat{\zeta}_t$ or C_K moves the operating point in the expected 2/3-power fashion rather than causing a linear blow-up. In practice, this implies that improvements in apparent tracking performance must be interpreted relative to this baseline, particularly in systems where actions influence observations.

Remark 6.17 (Practical estimation of V_Φ for black-box Φ_t). The estimable score requires $\sigma_\Phi = 1/(1 + V_\Phi)$, where V_Φ is the ISS-Lyapunov functional for the convergence subsystem. When Φ_t is a black box, with internal dynamics unknown and ISS comparison functions $(\alpha_{1\Phi}, \alpha_{2\Phi}, \alpha_{3\Phi}, \chi_\Phi)$ unavailable, V_Φ admits no closed-form derivation. One possible route is Counterexample-Guided Inductive Synthesis (CEGIS), in which a learner proposes a candidate $\hat{V}_\Phi$ from observed state trajectories and a verifier attempts to falsify its dissipation inequality on sampled trajectories. We have not implemented that pipeline in the present codebase, so it should be read as future work rather than as a validated part of the current diagnostic. For opaque adaptation maps, the present paper should therefore be interpreted as providing the template requirement on V_Φ , not an operational estimator for every black-box Φ_t .

6.2 Architecture and Behavior Are Not Equivalent

A structural asymmetry in the framework deserves explicit statement.

Drift-coherence implies the need for all three roles. By lemma 3.3 and theorem 3.6, any drift-coherent, action-effective system must implement the three functional roles $(\mathcal{P}, \mathcal{A}, \Phi)$. Drift-coherence is the stronger behavioral condition; the decomposition is its architectural consequence in action-effective settings.

Implementing the decomposition is not enough. A system may implement $(\mathcal{P}, \mathcal{A}, \Phi)$ structurally while failing to satisfy the ISS hypotheses (H-ISS-P/A/ Φ /Env) required for theorem 4.2. The components may be present but improperly tuned, have inadequate bandwidth, or lack sufficient expressiveness.

Proposition 6.18 (The Score Characterizes Component-Level Failure). *Let $I(S, t)$ be a three-role system under hypotheses (H-ISS-P/A/ Φ /Env). Then:*

- (i) $\mathcal{C}(I(S, t), t) = 1$ if and only if $V_P(t) = 0$, $H(\mathcal{A}_t \mid \mathcal{P}) = 0$, and $V_\Phi(t) = 0$ simultaneously.
- (ii) $\mathcal{C}(I(S, t), t) \rightarrow 0$ if and only if at least one of the three failure modes is active: $V_P(t) \rightarrow \infty$ (FM-1), $H(\mathcal{A}_t \mid \mathcal{P}) \rightarrow \infty$ (FM-2), or $V_\Phi(t) \rightarrow \infty$ (FM-3).
- (iii) The identity of $\arg \min(\sigma_P, \sigma_A, \sigma_\Phi)$ names the bottleneck component.

Proof. (i) Each score is a strictly decreasing function of its Lyapunov functional or conditional entropy: $\sigma_P = 1 \Leftrightarrow V_P = 0$, $\sigma_A = 1 \Leftrightarrow H(\mathcal{A}_t \mid \mathcal{P}) = 0$, $\sigma_\Phi = 1 \Leftrightarrow V_\Phi = 0$. The min equals 1 iff all three equal 1 simultaneously.

(ii) As $V_P(t) \rightarrow \infty$, $\sigma_P(t) \rightarrow 0$, so $\mathcal{C} \rightarrow 0$; similarly for $V_\Phi(t) \rightarrow \infty$. As $H(\mathcal{A}_t \mid \mathcal{P}) \rightarrow \infty$, $\sigma_A(t) = e^{-H} \rightarrow 0$, so $\mathcal{C} \rightarrow 0$. The converse holds because $\mathcal{C} \rightarrow 0$ requires at least one $\sigma_i \rightarrow 0$, which by the above requires the corresponding functional to diverge.

(iii) $\mathcal{C} = \min_i \sigma_i$; the argmin identifies which subsystem's functional is largest, i.e., which component is furthest from its reference state. $\square$

7 Empirical Validation

A Gaussian tracker provides an analytically tractable closed-loop tracking model that instantiates the three roles $(\mathcal{P}, \mathcal{A}, \Phi)$ and admits closed-form computation of all Lyapunov functionals.

The experiments verify the component-wise stability template by controlled ablation of each component under pure drift (section 7.2), validate the estimation error bound of proposition 6.9 numerically and characterize the lag-drift tradeoff (section 7.3), and demonstrate the score as a pre-failure early-warning diagnostic on a real concept-drift dataset (section 7.4). Code and reproduction instructions are provided in appendix A.

7.1 Formalization

The Gaussian tracker extends example 4.4. Let the environment produce observations from a drifting Gaussian:

$$s_t \sim \mathcal{N}(\mu^*(t), \sigma_{\text{obs}}^2), \quad \mu^*(t) = \mu_0 + \zeta t + \sigma_w B_t,$$

where B_t is standard Brownian motion. The tracker components are:

- $\mathcal{P}(S, t) = \mathcal{N}(\hat{\mu}_t, \hat{\sigma}_t^2)$: Gaussian belief state.
- $\mathcal{A}_t(\mathcal{P}) = \hat{\mu}_t$: deterministic mean extraction ($H(\mathcal{A}_t | \mathcal{P}) = 0$, $\delta = 0$).
- Φ_t : Kalman measurement update (see example 4.4).

All three Lyapunov functionals are available in closed form (example 4.4), making the tracker a convenient test bed for the trade-off and diagnostic claims studied here.

Lemma 7.1 (The Gaussian Tracker Satisfies the ISS Hypotheses). *The Gaussian tracker with $\sigma_w > 0$ and $\sigma_{\text{obs}} > 0$ satisfies all four ISS hypotheses (H-ISS-P/A/ Φ /Env) of theorem 4.2:*

- (i) **H-ISS-P**: $V_P(t) = \frac{1}{2}W_2(\mathcal{P}, \mathcal{P}^*)^2 = \frac{1}{2}(\hat{\mu}_t - \mu^*(t))^2$ is a valid ISS-Lyapunov functional; the Kalman update contracts V_P at rate $1 - K_t$ per step, where $K_t \in (0, 1)$ is the Kalman gain.
- (ii) **H-ISS-A**: $V_A = H(\mathcal{A}_t | \mathcal{P}) = 0$ since $\mathcal{A}_t = \hat{\mu}_t$ is deterministic; H-ISS-A is satisfied trivially.
- (iii) **H-ISS- Φ** : the Kalman recursion for $\hat{\sigma}_t^2$ converges to the steady-state variance $\hat{\sigma}_\infty^2 = \frac{1}{2}(\sigma_w^4 + \sqrt{\sigma_w^4 + 4\sigma_w^2\sigma_{\text{obs}}^2})$, a stable fixed point; $V_\Phi = |\hat{\sigma}_t^2 - \hat{\sigma}_\infty^2|$ decays geometrically with rate $1 - K_\infty$.
- (iv) **H-ISS-Env**: the environmental drift rate ζ is bounded by assumption; H-ISS-Env holds whenever $\zeta < (1 - K_\infty)\sigma_{\text{obs}}$.

Proof. Items (i)–(iii) follow directly from the Kalman filter’s convergence properties [Sontag, 1989, Mironchenko and Prieur, 2020]. Item (iv) is the H-ISS-Env condition of theorem 4.2 applied to the scalar Gaussian case; the contraction rate equals the steady-state Kalman gain $K_\infty = \hat{\sigma}_\infty^2/(\hat{\sigma}_\infty^2 + \sigma_{\text{obs}}^2)$. $\square$

7.2 Experiment A: Component Ablation

Protocol. We run $T = 500$ steps with $\zeta = 0.02$, $\sigma_w = 0.1$, $\sigma_{\text{obs}} = 1.0$, seed = 42, under four conditions:

C1. Full model: standard Gaussian tracker.

C2. FM-1 (static $\mathcal{P}$): $\hat{\mu}_t \equiv 0$ for all t ; Kalman update disabled.

Table 1: Gaussian-tracker experimental parameters. All experiments use `uv sync` from the repository for exact environment reproduction.

Parameter	Exp A	Exp B	Exp C/D
T (steps)	500	5,000	45,312 (ELEC2)
Seeds	42	0–49	fixed
σ_{obs}	1.0	1.0	1.0
ζ	0.02	0.001	estimated
σ_w	0.1	0.0 / 0.035	N/A
ε (Sinkhorn)	0.1	0.1	0.1
Window n	N/A	5–100	100
EMA α	N/A	0.005–0.5	0.05
PH threshold	N/A	N/A	84
PH δ	N/A	N/A	0.01
Warning level	N/A	N/A	0.295
Max lead gap	N/A	N/A	adaptive (166)

C3. FM-2 (non-det. $\mathcal{A}$): $\mathcal{A}_t = \hat{\mu}_t + \epsilon_t$, $\epsilon_t \sim \mathcal{N}(0, 0.25)$.

C4. FM-3 (absent Φ): Kalman update disabled; no feedback.

Environment: pure drift with no action feedback ($\alpha = 0$), so the world drifts independently of the system’s actuation. The effect of action-environment coupling is addressed in the Coercive Masking Corollary (corollary 8.2).

Results. Table 2 summarizes the tail-mean Lyapunov functionals and component coherence scores.

Table 2: Ablation results ($T = 500$, $\zeta = 0.02$, influence= 0, 42 seeds). Cells report mean $\pm$ standard deviation over the final 100 steps.

Condition	$\bar{V}_P$	$\bar{V}_A$	$\bar{V}_\Phi$	$\bar{\sigma}_P$	$\bar{\sigma}_A$	$\bar{\sigma}_\Phi$
Full model	0.051 $\pm$ 0.020	0.000 $\pm$ 0.000	0.010 $\pm$ 0.000	0.955 $\pm$ 0.017	1.000 $\pm$ 0.000	0.990 $\pm$ 0.000
FM-1 (static $\mathcal{P}$)	41.156 $\pm$ 16.005	0.000 $\pm$ 0.000	0.000 $\pm$ 0.000	0.034 $\pm$ 0.046	1.000 $\pm$ 0.000	1.000 $\pm$ 0.000
FM-2 (noisy $\mathcal{A}$)	0.051 $\pm$ 0.020	0.974 $\pm$ 0.139	0.010 $\pm$ 0.000	0.955 $\pm$ 0.017	0.579 $\pm$ 0.031	0.990 $\pm$ 0.000
FM-3 (absent Φ)	1.972 $\pm$ 3.680	0.000 $\pm$ 0.000	5.400 $\pm$ 0.000	0.573 $\pm$ 0.259	1.000 $\pm$ 0.000	0.157 $\pm$ 0.000

The results confirm the stability template of theorem 4.2: V_{total} is lowest under the full model. FM-1 and FM-3 both fail, but via different components: the score correctly identifies $\mathcal{P}$ as the bottleneck in FM-1 ($\bar{\sigma}_P = 0.040$) and Φ as the bottleneck in FM-3 ($\bar{\sigma}_\Phi = 0.157$), while FM-2 is selectively degraded through $\mathcal{A}$ ($\bar{\sigma}_A = 0.530$). The passive full-model run remains sharply concentrated near the equilibrium ($\bar{V}_{\text{total}} = 0.027$), whereas FM-1 rises to $\bar{V}_{\text{total}} = 24.259$, providing the diagnostic precision claimed in definition 6.2.

V_{total} is not sample-path monotone, consistent with theorem 4.2 which guarantees non-increasing *expectation* under the ISS hypotheses, not sample-path monotonicity of the stochastic process.




Figure 1: Component ablation: V_{total} (log scale, 20-step rolling mean, top) and $\mathcal{C}$ (bottom) under the four conditions. The full model is the only condition where $\mathcal{C}$ remains bounded away from 0. FM-1 and FM-3 both diverge in V_{total} but are distinguished by different bottleneck components: $\sigma_P \rightarrow 0$ for FM-1 (stale $\mathcal{P}$) and $\sigma_\Phi \rightarrow 0$ for FM-3 (absent Φ).

7.3 Experiment B: Sinkhorn Sample Complexity and the Coherence-Delay U-Curve

Protocol. We use the Gaussian tracker at stationary regime ($T = 5000$, 50 seeds, $\sigma_{\text{obs}} = 1$) and compute $\hat{\sigma}_P(t)$ using sliding-window estimators with $n \in \{5, 10, 25, 50, 100\}$. Mean absolute error (MAE) $|\sigma_P - \hat{\sigma}_P|$ is measured in the tail (last 40% of T).

The lag-drift tradeoff and the Coherence-Delay Principle. The estimation error has two competing terms (definition 6.11):

$$\mathcal{E}(n) \leq \underbrace{C_K n^{-1/2}}_{\text{variance term}} + \underbrace{\zeta \cdot \frac{n}{2}}_{\text{lag-bias term}}. \quad (4)$$

At large n or fast drift, the lag-bias term dominates and the total error increases with n , hiding the $O(n^{-1/2})$ variance rate.

To validate proposition 6.12 empirically, we additionally run a U-curve sweep: for each of four drift rates $\zeta \in \{0.001, 0.005, 0.01, 0.05\}$, we measure the raw tracking error $|\mu^*(t) - \bar{s}_{t-n:t}|$ across $n \in \{5, 10, 20, 50, 75, 100, 150, 200, 300, 500\}$ in a passive environment (influence = 0) over $T = 6,000$ steps and 20 seeds.




Figure 2: Coherence-Delay U-curve validation (proposition 6.12). *Left:* empirical tracking error vs. window size n for four drift rates (log scale; dotted lines are scaled theoretical curves; stars mark empirical n^*). Each curve shows the predicted U-shape: variance-dominated decay for small n , lag-dominated growth for large n . *Right:* $\mathcal{E}_{\min}$ vs. ζ on log-log scale. Fitted slope = 0.316 (95% CI [0.287, 0.345]); theoretical prediction = $1/3$.

Results. Figure 2 confirms the two predictions of proposition 6.12. The U-curves show the characteristic variance-lag tradeoff for each ζ , with empirical n^* shrinking as drift increases (from $n^* \approx 100$ at $\zeta = 0.001$ to $n^* \approx 10$ at $\zeta = 0.05$). The log-log slope of $\mathcal{E}_{\min}$ vs. ζ is **0.316** (95% CI [0.287, 0.345]), close to the theoretical $1/3$ scaling for this finite-window estimator class. Table 3 reports the extracted minima.

The ratio $\mathcal{E}_{\min}/\zeta^{1/3}$ remains close to constant (≈ 0.91 across the four drift levels), confirming the expected $\zeta^{1/3}$ scaling despite finite-window effects.

Table 3: Empirical validation of the Coherence-Delay bound ($T = 6,000$, 20 seeds, passive environment).

ζ	n_{obs}^*	$\mathcal{E}_{\min}$	$\mathcal{E}_{\min}/\zeta^{1/3}$
0.001	100	0.0924	0.9236
0.005	50	0.1566	0.9159
0.010	20	0.1973	0.9158
0.050	10	0.3177	0.8623

EWMA baseline and finite-memory universality. To test whether the observed cube-root law is specific to the uniform sliding-window estimator, we ran the same Gaussian drift sweep with an exponentially weighted moving average (EWMA) $\hat{\mu}_t = \alpha s_t + (1 - \alpha)\hat{\mu}_{t-1}$ over $\alpha \in \{0.005, 0.01, 0.02, 0.05, 0.1, 0.2, 0.3, 0.5\}$.

Table 4: EWMA baseline on the Gaussian tracker: optimal forgetting factor by drift rate.

ζ	α^*
0.001	0.02
0.005	0.05
0.010	0.10
0.050	0.20

The minimum-error EWMA slope is **0.3173** on a log-log fit of $\mathcal{E}_{\min}$ vs. ζ , essentially identical to the sliding-window slope **0.316** in fig. 2. The optimal forgetting factor follows log-log slope **0.6011**, close to the $2/3$ prediction from corollary 6.13. Empirically, the optimal forgetting factors increase with drift exactly as finite effective memory theory predicts, supporting the claim that the cube-root law is a property of the finite-memory class rather than an artifact of uniform windows.

For the score estimation rate, we use slow drift ($\zeta = 0.001$, $\sigma_w = 0.035$) to stay near the variance-dominated regime while retaining a visible lag term at $n = 100$. The observed log-log slope is -0.436 (95% CI $[-0.865, -0.008]$; target -0.5); the deviation reflects the onset of lag-bias at the largest window and residual autocorrelation induced by the slow-drift regime (table 5 and fig. 3).

Table 5: Sinkhorn estimation error vs window size ($\zeta = 0.001$, 50 seeds, $T = 5000$; mean $\pm$ standard deviation).

n	MAE $ \sigma_P - \hat{\sigma}_P $	std
5	0.0801	0.0036
10	0.0457	0.0032
25	0.0233	0.0026
50	0.0188	0.0029
100	0.0245	0.0061
Log-log slope: -0.436 (target: -0.5)		




Figure 3: Sinkhorn sample complexity in the variance-dominated regime ($\zeta = 0.001$, $n \leq 100$, 50 seeds, $T = 5000$). Empirical log-log slope = -0.436 (95% CI $[-0.865, -0.008]$); theoretical target -0.5 . This is the variance-only regime of the U-curve in fig. 2.

7.4 Experiment C: ELEC2 Early-Warning Lead Time

Dataset. The ELEC2 dataset [Harries, 1999] records 45,312 hourly observations of electricity demand in New South Wales, Australia, exhibiting documented concept drift due to market deregulation. We apply a 1D Gaussian tracker to the normalized demand stream.

Protocol. Failures are detected by the Page-Hinkley test [Harries, 1999, threshold = 84, $\delta = 0.01$], yielding 113 regime-shift events on the normalized NSW demand stream. $\hat{\mathcal{C}}$ is computed from a 1D Gaussian tracker on the normalized target stream, with all detector hyperparameters fixed before evaluation. The passive benchmark therefore uses no effort proxy and no adaptive threshold retuning on future events. CUSUM, forgetting-factor RLS, and scalar Kalman baselines are run on the same normalized stream, while ADWIN consumes the residual stream used in the original benchmark. The Fréchet baseline compares each current window to the healthy prefix in the same scalar representation space. $\hat{\mathcal{C}}$ warnings are recorded on each falling-edge crossing of $\hat{\sigma}_P < 0.295$. Because ELEC2 exposes no direct actuation or effort channel, this benchmark uses the estimable score rather than $\mathcal{C}^E$; the effort-corrected score is reserved for the active/coercive setting of section 7.7. Lead time is defined as the gap between a warning and the nearest subsequent Page-Hinkley event, using an adaptive maximum gap set to half the median inter-event interval (nearest-future matching; see appendix A for implementation).

Results. Across the 45,312-step stream, the fixed- $n = 100$ score issues 58 warnings (falling-edge crossings of $\hat{\sigma}_P < 0.295$), of which **18 are matched to a subsequent Page-Hinkley event** within the adaptive matching gap (half the median inter-event interval; 166 steps here). The median lead time is **94 steps**; mean is 96.9; minimum 45; maximum 157. For comparison, ADWIN with $\delta = 0.03$ on the same




Figure 4: ELEC2 early-warning diagnostic. *Top*: $\hat{\mathcal{C}}$ via $\hat{\sigma}_P(t)$ (sliding window $n = 100$) with warning threshold 0.295 (dashed red). Vertical lines mark score warnings (red), ADWIN warnings (purple), and Page-Hinkley failures (grey). *Bottom*: residual signal fed to ADWIN (50-step rolling mean).

residual stream issues 114 warnings, of which 54 are matched to a subsequent Page-Hinkley event, at 47.4% precision and median lead 84.5. At lower thresholds the detector becomes more conservative; at higher thresholds it produces more warnings but the same one-to-one matching rule reduces effective precision. On this passive benchmark ADWIN remains the strongest generic change detector under this calibration, while the score continues to serve as a tuneable structural diagnostic whose score semantics carry over to the active setting where $\mathcal{C}^E$ becomes the correct observable.

7.5 Experiment D: Dynamic n_t^* Optimization on ELEC2

Proposition 6.12 provides the self-tuning rule $n_t^* = (C_K/\hat{\zeta}_t)^{2/3}$. This experiment tests whether deploying that rule dynamically on ELEC2 improves on the fixed- $n = 100$ score baseline while remaining interpretable next to a strong passive detector.

Protocol. We estimate the local drift rate at each step using an exponentially weighted moving average (EMA, decay $\alpha = 0.05$) of the absolute change in mean between consecutive windows of $d = 50$ steps: $\hat{\zeta}_t = \text{EMA}_\alpha(|\bar{s}_{t-d:t} - \bar{s}_{t-2d:t-d}|/d)$. The constant C_K is calibrated once on the first 2,000 steps (near-stationary): $C_K \approx 1.25 |\bar{s}_t - \bar{s}| \cdot \sqrt{100}$. The resulting $n_t^* = (C_K/\hat{\zeta}_t)^{2/3}$ is clipped to $[30, 300]$, with $\hat{\zeta}_t$ estimated from consecutive 24-step means and EMA decay $\alpha = 0.03$. The full score sweep in the figure includes fixed $n = 50, 100$, and 300 as references. The main quantitative comparison reported in the table uses the validated fixed- $n = 100$ baseline, the dynamic n_t^* rule, and ADWIN on the same residual stream as an external passive baseline. Warnings and lead times are computed identically to section 7.4, using the same warning threshold 0.295 and one-to-one future matching.

Table 6: ELEC2 early-warning comparison: fixed windows, dynamic n_t^* , generic passive baselines, and ADWIN. 113 PH failures; warning level = 0.295; max gap adaptive (half the median inter-event interval; 166 steps here).

Strategy	Warnings	Leads	Precision	Median lead
Fixed $n = 50$	24	11	46%	91
Fixed $n = 100$	58	18	31%	94
Fixed $n = 300$	83	34	41%	44
Dynamic n_t^*	95	29	31%	101
CUSUM	793	96	12%	141.5
FF-RLS	151	45	30%	55
Kalman	96	24	25%	91
Fréchet	0	0	0%	NA
ADWIN ($\delta = 0.03$)	114	54	47%	84.5

Results. The results are qualitatively consistent with proposition 6.12, but the passive-stream nature of ELEC2 keeps the advantage modest. Relative to the validated fixed- $n = 100$ baseline, the dynamic rule increases matched leads from 18 to 29 and slightly improves median lead time from 94 to 101, but both remain well behind ADWIN on raw passive detection.




Figure 5: Dynamic n_t^* optimization on ELEC2 ($C_K \approx 1.25 \overline{|s_t - \bar{s}|} \sqrt{100}$, EMA $\alpha = 0.03$, drift window $d = 24$). *Top*: $\hat{\sigma}_P$ (100-step mean) for all three strategies; grey lines are PH regime shifts. *Middle*: dynamic window size n_t^* , contracting sharply near regime shifts and expanding during stable periods. *Bottom*: lead-time distributions, including ADWIN, CUSUM, forgetting-factor RLS, scalar Kalman, and Fréchet baselines.

The added passive baselines do not change that ordering. CUSUM is highly sensitive but imprecise; FF-RLS and Kalman sit between the fixed-window score and ADWIN, while the Fréchet baseline is too conservative on ELEC2 to fire at this calibration. None overtakes ADWIN on matched leads or precision. This is expected: ADWIN is designed for passive changepoint detection, whereas the score targets structural validity under feedback and extends naturally to the effort-corrected $\mathcal{C}^E$ used in the active masking experiment. The Bikes benchmark below shows the same passive-stream pattern on a longer real-world demand sequence.

The trade-off is structural and expected: the dynamic rule improves on the static score baseline without changing the basic passive-stream ordering, which is exactly the honest reading needed for the later active benchmarks.

The dynamic rule is also moderately robust to calibration error, as predicted by the 2/3 sensitivity of n_t^* . On ELEC2, scaling $\hat{\zeta}_t$ by $0.75\times$, $1.5\times$, and $2.0\times$ moves the mean dynamic window from 47.7 to 57.8, 36.7, and 31.9, while matched leads change from 29 to 22, 21, and 19. Perturbing C_K by the same factors gives mean windows 39.5, 62.5, and 75.7 with 25, 18, and 19 matched leads. The Bikes benchmark is more tolerant: $\hat{\zeta}_t$ perturbations yield 92, 76, and 74 matched leads, while scaling C_K upward to $1.5\times$ and $2.0\times$ actually increases matched leads to 98 and 110. These shifts are material but not catastrophic, which is the practical point of the sublinear sensitivity law.

7.6 Experiment E: Flagship Real-Data Masking on KuaiRand

Dataset and protocol. We treat KuaiRand-Pure as the real-data masking benchmark. The random-exposure slice is treated as healthy, the first half of the later standard-recommendation slice as bubble formation, and the second half as collapse. The benchmark is evaluated on 367 users with at least 20 interactions in each relevant phase. User signals are computed on 20-step windows over `watch_ratio` with video-tag exposure summaries; the effort term is the KL divergence between the current exposure-tag distribution and the user’s healthy random-exposure baseline. The KuaiRand implementation uses a surrogate σ_A based on normalized exposure entropy rather than on replicated action logs, so the benchmark should be read as an operational logged-data proxy for actuation reproducibility rather than as exact $e^{-H(A|\mathcal{P})}$. To keep the comparison honest, the score and $\mathcal{C}^E$ thresholds are calibrated only on healthy random-exposure windows, using the 20th percentile of the healthy score distribution, and the reported E_0 is the median healthy effort scale. We also report an EWMA-smoothed $\mathcal{C}^E$ baseline to distinguish genuine masking from simple temporal smoothing.

Results. The passive detectors behave unevenly once all of them see the same effort-adjusted signal. KSWIN is the most sensitive generic baseline here, while ADWIN and PageHinkley stay inert under this calibration. The effort-corrected score remains more sensitive than the plain score: bubble detection rises from 45.8% to 69.2%, and collapse detection rises from 54.8% to 70.0%. The EWMA-smoothed baseline sits in between, which is exactly what we want from a conservative baseline: it damps short fluctuations but does not match the raw effort-aware diagnostic. The key point of the benchmark is that the effort term exposes a form of control-induced



Figure 6: KuaiRand logged benchmark. The random-exposure phase is treated as healthy, the first half of the later standard-recommendation slice as bubble formation, and the second half as collapse. Thresholds for $\mathcal{C}$, $\mathcal{C}^E$, and EWMA-smoothed $\mathcal{C}^E$ are calibrated only on healthy windows.

coherence that a generic detector may or may not surface depending on its operating assumptions.

The KuaiRand λ sweep gives the same calibration message as the particle tracker. Bubble detection for $\mathcal{C}^E$ rises from 0.463 at $\lambda = 0.5$ to 0.575 at $\lambda = 2.0$, 0.692 at $\lambda = 3.0$, and 0.747 at $\lambda = 4.0$, while healthy false positives rise from 0.777 to 1.215, 1.387, and 1.632 per user. Together with the particle masking-gap sweep, this makes $\lambda \in [2, 3]$ the stable operating region and justifies the paper default $\lambda = 3$ as a compromise rather than a hand-tuned optimum.

7.7 Experiment F: Feedback Masking in the Particle Tracker

This experiment instantiates corollary 8.2 and definition 8.4 in the non-Gaussian, non-linear particle tracker. The environment evolves according to the scalar particle-filter simulator implemented in `experiments/particle/model.py`, but now with an explicit action-coupling parameter $\alpha \in [0, 1]$ so that the actuation can pull the latent state toward the stale controller output. In the synthetic setting we operationalise the effort proxy as the *realized coercive effort*

$$E_t^{\text{sim}} = \alpha |a_t - s_t^{\text{free}}|,$$

where s_t^{free} is the latent state that would have arisen at time t before applying control. This distinguishes mere action-world mismatch from actual world-forcing.

Protocol. We freeze $\mathcal{P}$ (FM-1) and compare two regimes: passive ($\alpha = 0$) and coercive ($\alpha > 0$). The main single-run diagnostic uses $\alpha = 0.3$ and effort penalty $\lambda = 3.0$. We then sweep $\alpha \in \{0.1, 0.2, 0.3, 0.5, 0.7, 0.9\}$ and $\lambda \in \{0.5, 1.0, 2.0, 3.0, 4.0\}$ over 8 seeds, exporting all raw and aggregated metrics from the experiment harness. To keep the main text readable, table 9 reports a five-regime masking matrix at the paper default $\lambda = 3.0$: passive, weak, mild, moderate, and strong coercive coupling.

Coercive masking: score stays high while effort-corrected score pays for effort




Figure 7: Synthetic feedback masking in the particle tracker. Left column: passive FM-1, where effort remains at zero and $\mathcal{C}$ and $\mathcal{C}^E$ coincide. Right column: coercive FM-1, where the action gap remains nonzero, realized coercive effort appears once influence is positive, and $\mathcal{C}^E$ drops once effort is penalized.

Table 7: KuaiRand logged benchmark summary on 367 users.

Phase	Detector	Detections	Rate	Median delay
Bubble onset	$\mathcal{C}$	168/367	0.458	27.0
Bubble onset	$\mathcal{C}^E$	254/367	0.692	31.8
Bubble onset	$\mathcal{C}^E$ -EWMA	194/367	0.529	35.5
Bubble onset	ADWIN	0/367	0.000	NA
Bubble onset	PageHinkley	0/367	0.000	NA
Bubble onset	KSWIN	225/367	0.613	23.0
Bubble onset	NoDrift	0/367	0.00	NA
Collapse onset	$\mathcal{C}$	201/367	0.548	37.0
Collapse onset	$\mathcal{C}^E$	257/367	0.700	45.0
Collapse onset	$\mathcal{C}^E$ -EWMA	208/367	0.567	41.8
Collapse onset	ADWIN	0/367	0.000	NA
Collapse onset	PageHinkley	0/367	0.000	NA
Collapse onset	KSWIN	357/367	0.973	26.5
Collapse onset	NoDrift	0/367	0.00	NA

Results. The matrix makes the masking mechanism explicit. As coupling increases, the stale FM-1 controller looks progressively healthier under the plain score: $\mathcal{C}$ rises from 0.795 in the passive regime to 0.999 under strong coercion. The effort-aware diagnostic moves in the opposite direction: $\mathcal{C}^E$ falls from 0.795 to 0.518, and the masking gap grows from 0 to 0.481. This is the closed-loop signature we care about: apparent coherence is improving while effort-aware coherence is deteriorating.




Figure 8: Multi-seed parameter sweep for feedback masking in the particle tracker (8 seeds). Left: mean standard $\mathcal{C}$. Middle: mean effort-corrected $\mathcal{C}^E$. Right: masking gap $M_t = \mathcal{C} - \mathcal{C}^E$. The standard score monotonically rises with influence, while the masking gap widens with both influence and effort penalty λ .

The sweep confirms that the correction is not an artifact of one parameter choice. At fixed $\lambda = 3.0$, the masking gap increases from 0.115 at $\alpha = 0.1$ to 0.481 at $\alpha = 0.9$, with a clear intermediate effect already at $\alpha = 0.2$ (gap = 0.194). At fixed $\alpha = 0.3$, increasing λ from 0.5 to 4.0 expands the gap from 0.051 to 0.313. Thus, stronger coercive capacity and stronger effort penalisation separate apparent from genuine coherence in the expected direction. The particle tracker closes the

Table 8: Calibration summaries for the dynamic window rule and effort penalty.

Setting	Perturbation / λ	Primary metric	Auxiliary metric
ELEC2 dyn. n_t^*	baseline	29 matched leads	mean window 47.7
ELEC2 dyn. n_t^*	$\hat{\zeta}_t \times 0.75$	22 matched leads	mean window 57.8
ELEC2 dyn. n_t^*	$\hat{\zeta}_t \times 1.5$	21 matched leads	mean window 36.7
ELEC2 dyn. n_t^*	$C_K \times 0.75$	25 matched leads	mean window 39.5
ELEC2 dyn. n_t^*	$C_K \times 1.5$	18 matched leads	mean window 62.5
Bikes dyn. n_t^*	baseline	85 matched leads	mean window 113.3
Bikes dyn. n_t^*	$\hat{\zeta}_t \times 0.75$	92 matched leads	mean window 136.1
Bikes dyn. n_t^*	$\hat{\zeta}_t \times 1.5$	76 matched leads	mean window 86.9
Bikes dyn. n_t^*	$C_K \times 1.5$	98 matched leads	mean window 146.3
Bikes dyn. n_t^*	$C_K \times 2.0$	110 matched leads	mean window 173.4
Particle ($\alpha = 0.3$)	$\lambda = 0.5$	$\mathcal{C}^E = 0.898$	$M_t = 0.051$
Particle ($\alpha = 0.3$)	$\lambda = 2.0$	$\mathcal{C}^E = 0.768$	$M_t = 0.181$
Particle ($\alpha = 0.3$)	$\lambda = 3.0$	$\mathcal{C}^E = 0.697$	$M_t = 0.252$
Particle ($\alpha = 0.3$)	$\lambda = 4.0$	$\mathcal{C}^E = 0.636$	$M_t = 0.313$
KuaiRand	$\lambda = 0.5$	bubble rate 0.463	healthy FP/u 0.777
KuaiRand	$\lambda = 2.0$	bubble rate 0.575	healthy FP/u 1.215
KuaiRand	$\lambda = 3.0$	bubble rate 0.692	healthy FP/u 1.387
KuaiRand	$\lambda = 4.0$	bubble rate 0.747	healthy FP/u 1.632

Table 9: Five-regime masking matrix for frozen- $\mathcal{P}$ FM-1 in the particle tracker at $\lambda = 3.0$ (8 seeds). Cells report tail means.

Regime	α	effort	$\mathcal{C}$	$\mathcal{C}^E$	M_t
Passive	0.0	0.000	0.795	0.795	0.000
Weak coercion	0.1	0.054	0.872	0.757	0.115
Mild coercion	0.2	0.090	0.920	0.726	0.194
Moderate coercion	0.3	0.118	0.949	0.697	0.252
Strong coercion	0.9	0.266	0.999	0.518	0.481

empirical loop on corollary 8.2: the standard score can be masked by control, whereas $\mathcal{C}^E$ remains sensitive to the intervention cost that creates the mask. Combined with the calibration grid in fig. 10, this supports the practical rule used throughout the paper: set E_0 to the healthy effort scale and choose λ in the range 2–3, with $\lambda = 3$ as a strong default.

To test whether this advantage survives comparison with a generic change detector, we build an active benchmark with two known regime transitions: healthy tracking for $t < 200$, feedback masking for $200 \leq t < 400$, and coercive collapse for $t \geq 400$. The masking phase freezes $\mathcal{P}$ while maintaining moderate coupling ($\alpha = 0.3$); the collapse phase sharply increases exogenous drift. We compare thresholded $\mathcal{C}$, thresholded $\mathcal{C}^E$, ADWIN, PageHinkley, KSWIN, and NoDrift on the raw observation stream, giving the generic detectors a fully observable signal rather than a model-derived residual.

The benchmark isolates the point that passive ELEC2 cannot address. ADWIN is stronger on passive streams because it is a generic changepoint detector.




Figure 9: Active benchmark in the particle tracker with two known phase changes. Orange line: masking onset at $t = 200$. Red line: collapse onset at $t = 400$. ADWIN receives the raw observation stream. $\mathcal{C}$ and $\mathcal{C}^E$ are thresholded at 0.8.

Table 10: Active benchmark summary in the particle tracker (12 seeds). Detection is measured against the known onset of feedback masking and the known onset of coercive collapse.

Phase	Detector	Detections	Rate	Median delay
Masking onset	$\mathcal{C}$	12/12	1.00	25.5
Masking onset	$\mathcal{C}^E$	12/12	1.00	5.0
Masking onset	ADWIN	7/12	0.58	279.0
Masking onset	PageHinkley	12/12	1.00	68.5
Masking onset	KSWIN	1/12	0.08	118.0
Masking onset	NoDrift	0/12	0.00	NA
Collapse onset	$\mathcal{C}$	12/12	1.00	2.0
Collapse onset	$\mathcal{C}^E$	12/12	1.00	0.0
Collapse onset	ADWIN	5/12	0.42	111.0
Collapse onset	PageHinkley	12/12	1.00	37.5
Collapse onset	KSWIN	0/12	0.00	NA
Collapse onset	NoDrift	0/12	0.00	NA

PageHinkley is reliable but slower, KSWIN is too conservative on this signal, and NoDrift is the sanity baseline. In the active setting, however, the relevant anomaly is not merely distributional change but coercive maintenance of an apparently healthy world-model fit. Here $\mathcal{C}^E$ detects masking onset in all 12 runs with median delay 5.0, whereas ADWIN detects only 7 of 12 masking onsets and does so with median delay 279.0. For the collapse transition the same ordering persists: ADWIN detects only 5 of 12 runs, while $\mathcal{C}$ and $\mathcal{C}^E$ detect all of them, with $\mathcal{C}^E$ firing earliest. This is the regime in which the effort-corrected index provides information that a generic passive detector does not encode.




Figure 10: Calibration grid for $\mathcal{C}^E$ over effort penalty λ and threshold. The rate panel is omitted because it saturates across most of the grid; the figure instead reports masking delay, collapse delay, and healthy false positives, which are the informative calibration axes. The paper setting ($\lambda = 3.0$, threshold 0.80) remains a low-delay compromise.

7.8 Experiment G: Bikes Early-Warning Lead Time

Dataset. The River Bikes stream records a long real-world bike-demand sequence with seasonal and weather-driven changes. We normalize the target stream to unit variance and use it as a second passive benchmark beyond ELEC2.

Protocol. As in ELEC2, Page-Hinkley events on the normalized stream define the warning reference. The same fixed-window score strategies ($n = 50, 100, 300$), the dynamic n_t^* rule, and ADWIN on the residual stream are evaluated under the same one-to-one future matching rule. The dynamic rule uses the same deployment recipe as ELEC2: C_K is calibrated on the initial prefix, n_t^* is clipped to a bounded range, and the warning threshold is held fixed rather than tuned against future failures. CUSUM, forgetting-factor RLS, and scalar Kalman baselines are evaluated on the normalized stream with the same thresholding discipline. The Fréchet baseline compares each current window to the healthy prefix in the same representation space.




Figure 11: Bikes early-warning diagnostic. *Top*: $\hat{\sigma}_P(t)$ (sliding window $n = 100$) with warning threshold 0.295 (dashed red). Vertical lines mark score warnings (red), ADWIN warnings (purple), and Page-Hinkley failures (grey). *Bottom*: residual signal fed to ADWIN (50-step rolling mean).

Results. On Bikes, ADWIN again outperforms the score-based diagnostics as a generic passive detector, while the dynamic window rule improves modestly over fixed $n = 100$ in precision but not in lead count. The new CUSUM, FF-RLS, Kalman, and Fréchet baselines remain weaker than ADWIN under this calibration, which confirms that the passive-stream caveat is not ELEC2-specific.

Table 11: Bikes early-warning comparison: fixed windows, dynamic n_t^* , generic passive baselines, and ADWIN. The same one-to-one future matching rule is used as in ELEC2.

Strategy	Warnings	Leads	Precision	Median lead	Min lead
Fixed $n = 50$	572	75	13%	123	0
Fixed $n = 100$	989	101	10%	87	0
Fixed $n = 300$	3,363	137	4%	58	0
Dynamic n_t^*	717	85	12%	127	0
CUSUM	5,928	384	6%	214	11
FF-RLS	15,687	327	2%	220	0
Kalman	13,985	319	2%	220	0
Fréchet	21	4	19%	154.5	65
ADWIN ($\delta = 0.03$)	477	250	52%	92	0

7.9 Passive and Active Diagnostics Are Complementary

The experiments support a clean division of labor. On passive streams such as ELEC2 and Bikes, the relevant anomaly is changepoint detection on exogenous data. In that regime, generic detectors such as ADWIN are the right default tools, and the score should be read as a structural diagnostic with tunable semantics rather than as a universal replacement for passive alarms.

On active or logged benchmarks, the target anomaly changes. In the particle tracker and KuaiRand, the problem is not only that the data distribution moves, but that intervention can keep apparent model-world fit high while a stale model is being protected by effort. In that regime, $\mathcal{C}^E$ and the masking gap are the relevant observables. KSWIN remains a competitive generic detector on KuaiRand, but it pays with false positives comparable to or higher than the effort-aware diagnostic, while ADWIN and PageHinkley contribute little under the same calibration. The practical conclusion is therefore complementary rather than adversarial: use passive drift detectors for passive streams, and use effort-aware coherence when the system can actively shape the observations it is judged against.

7.10 Discussion: Limitations of the Empirical Validation

Four limitations bound the scope of these results. First, the Gaussian-linear setting satisfies H-ISS-P/A/ Φ /Env by construction (example 4.4); non-Gaussian, non-linear systems require case-by-case hypothesis verification. Second, the Sinkhorn rate is validated only for the finite-memory, sliding-window estimator class analyzed here, and only in the variance-dominated regime ($\zeta \leq 0.001$, $n \leq 100$); high-dimensional or fast-drift systems require the debiased S_ε with lag correction. Third, the ELEC2 failure definition (Page-Hinkley events) is one operational choice; different failure definitions yield different lead-time statistics. Fourth, the dynamic n_t^* strategy’s precision depends critically on $\hat{\zeta}_t$ estimation quality; a more accurate real-time drift estimator would sharpen the precision-recall frontier.

The Gaussian tracker provides the reproducible, analytically tractable validation of the trade-off and diagnostic claims. High-dimensional, non-Gaussian, and non-linear extensions constitute open directions for future validation (section 10).

8 Relation to Algorithmic Regulation

Ruffini [2026] reformulates the good regulator principle in algorithmic terms: a regulator that materially simplifies the realized world trajectory is unlikely to be independent of that world. In Ruffini’s setting, regulation is read through compression of the regulated output. The question here is different: not how much compression a regulator achieves on a realized episode, but how long a compressed internal model remains valid when the world itself drifts.

Tracking under drift can be interpreted as *online compression of a non-stationary process*. A finite-memory model compresses recent observations into a compact state, but that compression ages. The cube-root law in proposition 6.12 quantifies the resulting trade-off between statistical efficiency and temporal validity.

The decomposition $(\mathcal{P}, \mathcal{A}, \Phi)$ then serves as an explicit model of the three channels that matter in closed-loop regulation: what the system currently represents, what it does, and how it updates. Ruffini studies the compression value of regulation on individual episodes, whereas the analysis here tracks the stability of that compression under drift and feedback.

8.1 Coercive Masking

The central point of contact with algorithmic regulation is the masking regime. A controller can keep the realized trajectory simple either because it tracks the world well or because it pushes the world toward a stale model. The latter mechanism is invisible to a diagnostic that looks only at model-world fit.

Definition 8.1 (Coercive Regime). An information system $I(S, t) = (\mathcal{P}, \mathcal{A}, \Phi)$ operates in the *coercive regime* if:

- (i) $\mathcal{A}_t$ is action-effective (definition 3.2) with coupling coefficient $\alpha > 0$: the transition kernel is $P(s_{t+1} \mid s_t, o) = \mathcal{N}(\mu^*(t) + \alpha(o - \mu^*(t)), \sigma_w^2)$;
- (ii) $\mathcal{P}$ is static (proposition 5.1): $\hat{\mu}_t \equiv \hat{\mu}_0$ and $\mathcal{A}_t(\mathcal{P}) \equiv \hat{\mu}_0$;
- (iii) the exogenous drift is bounded: $|\zeta| < \alpha \cdot |\hat{\mu}_0 - \mu_0|$.

In the coercive regime, the system’s stale actions continuously pull μ^* toward the frozen estimate $\hat{\mu}_0$: $\mu^*(t+1) = \mu^*(t) + \zeta + \alpha(\hat{\mu}_0 - \mu^*(t))$, which is a mean-reverting process with equilibrium $\hat{\mu}_0 + \zeta/\alpha$. When $|\zeta/\alpha|$ is small, μ^* is held near $\hat{\mu}_0$, suppressing V_P and making the system appear healthy under the standard score.

Corollary 8.2 (Coercive Masking). *In the coercive regime (definition 8.1), the standard score $\sigma_P(t) = 1/(1+V_P(t))$ can fail to detect FM-1: $\sigma_P(t) \approx 1$ for all $t \leq t_{\text{mask}}$, where $t_{\text{mask}} = O(\alpha^{-1} \log(\alpha/\zeta))$ is the coercive horizon. Beyond t_{mask} , exogenous drift overwhelms the coercive capacity and V_P diverges.*

Proof sketch. Under condition (iii) of definition 8.1, the gap $\delta_t = \mu^*(t) - \hat{\mu}_0$ satisfies $\delta_{t+1} = (1 - \alpha)\delta_t + \zeta$. The fixed point is $\delta^* = \zeta/\alpha$, so $V_P(t) \leq \frac{1}{2}(\zeta/\alpha)^2 + \sigma_w^2/(2\alpha)$ in expectation, masking FM-1 while the coercive condition holds. Once the drift accumulates beyond the available control authority, the frozen model can no longer hold the environment in place and the underlying mismatch reappears. $\square$

Remark 8.3 (Compression by adaptation vs. compression by forcing). From the algorithmic-regulation viewpoint, coercive masking matters because it can improve realized compressibility without improving the model. The external trajectory becomes simpler, but for the wrong reason: the controller is forcing the world into alignment rather than updating its internal state to match a changing world.

Definition 8.4 (Effort-Corrected Score). The *effort signal* at time t is any monotone proxy for the corrective actuation burden required to keep the world aligned with the system’s present output. A generic observable surrogate is

$$E_t^{\text{obs}} := |\mathcal{A}_t(\mathcal{P}) - \mathcal{A}_{t-1}(\mathcal{P})|,$$

the magnitude of change in the system’s actuation. In controlled synthetic settings, a sharper proxy is the realized coercive effort

$$E_t^{\text{sim}} := \alpha |a_t - s_t^{\text{free}}|,$$

where s_t^{free} is the state that would have arisen before applying control.

The *effort-corrected score* is

$$\mathcal{C}^E(I(S, t), t) := \min(\sigma_P(t) \cdot e^{-\lambda E_t/E_0}, \sigma_A(t), \sigma_\Phi(t)),$$

where $E_0 > 0$ is a reference effort scale and $\lambda > 0$ is a sensitivity parameter. $\mathcal{C}^E$ penalizes increasing actuation effort, correcting for action-induced confounding: it discounts apparent coherence when the environment is being held in place rather than genuinely estimated. In action-effective settings where coercive masking is plausible, $\mathcal{C}^E$ is the primary closed-loop diagnostic and $\mathcal{C}$ is the corresponding apparent-fit baseline.

Definition 8.5 (Masking Gap). The *masking gap* at time t is

$$M_t := \mathcal{C}(I(S, t), t) - \mathcal{C}^E(I(S, t), t) \geq 0.$$

It quantifies how much apparent coherence is lost once corrective effort is charged explicitly. Large M_t indicates that the observed fit is being maintained by intervention rather than by an up-to-date model.

Remark 8.6 (Calibration guidance for E_t , E_0 , and λ). The effort proxy should follow the strongest observable hierarchy available: E_t^{sim} in controlled simulators, E_t^{obs} from direct actuator burden in deployed systems, and E_t^{log} from logged exposure or policy divergence when only observational traces are available.

The reference scale E_0 should be set to the typical healthy effort level of the domain, not to a global constant. In practice this means using a nominal controller effort in simulation or the median healthy effort over a temporally held-out prefix in logged data. In the particle tracker we use the prior-scale effort floor, in KuaiRand we use the median healthy KL divergence from the random-exposure baseline, and in deployed control systems the right default is the nominal actuator burden under steady operation.

The penalty λ controls how aggressively apparent coherence is discounted. Across the particle masking grid and the KuaiRand sensitivity sweep, values in the range $\lambda \in [2, 3]$ provide the most stable trade-off, with $\lambda = 3$ serving as the paper

default: $\lambda < 1$ under-penalizes masking, while $\lambda = 4$ tends to be more aggressive than the added separation usually justifies.

These sweeps function as the paper’s current ablations for proxy strength and effort sensitivity: they show that the main conclusions are stable across a useful operating range, while still making clear that calibration is part of the diagnostic rather than a nuisance parameter hidden from the analysis.

The empirical evidence for coercive masking comes from the particle-tracker experiment (section 7.7). In the passive FM-1 regime, the standard score and $\mathcal{C}^E$ coincide because no realized forcing of the world occurs, so $M_t = 0$. Under moderate coupling ($\alpha = 0.3$, $\lambda = 3.0$), the same frozen model reaches $\mathcal{C} = 0.949 \pm 0.012$ while $\mathcal{C}^E = 0.697 \pm 0.046$, so the mean masking gap is $M_t \approx 0.252$, showing that the apparent coherence is being purchased by intervention rather than adaptation.

9 Related Work

Tracking under non-stationarity. The closest technical neighborhood is concept drift and streaming adaptation [Gama et al., 2014]. Classical detectors such as ADWIN, DDM, and Page-Hinkley identify changepoints or sustained mean shifts in a monitored signal [Bifet and Gavaldà, 2007, Gama et al., 2004, Harries, 1999]. On passive streams they are baseline changepoint detectors, whereas the coherence score is intended as a structural diagnostic tied to model fidelity and update quality rather than as a universal replacement for generic drift alarms. Related finite-memory results in nonstationary optimization and online least-squares establish the same cube-root trade-off between variation and memory, but in Euclidean loss or bandit settings rather than in a distributional W_2 /Sinkhorn geometry [Besbes et al., 2015, Cheung et al., 2019, Yuan and Lamperski, 2020, Mazzetto and Upfal, 2023]. The kernel-level lower bound added in this revision makes that comparison more precise for causal averaging estimators: the same bias-vs.-variance tension now appears directly in the finite-memory class studied here. Classical smoothing rules such as EWMA and Kalman filtering also inhabit this neighborhood: they trade responsiveness against variance through forgetting or effective memory, but they are usually analyzed as estimators rather than as a closed-loop diagnostic baseline. Our EWMA results are included precisely to show that the cube-root law is a finite-memory phenomenon rather than a sliding-window artifact.

Transport-based diagnostics. WATCH uses Wasserstein geometry to detect distribution shift in feature space [Faber et al., 2022b,a]. Our work is adjacent but different in emphasis. We use W_2 and its debiased Sinkhorn surrogate inside a finite-memory tracking model to derive a lag-variance trade-off and a coherence score for model coherence.

Filtering and closed-loop control. The representation–action–adaptation decomposition is motivated by standard state-estimation and control logic, not by a claim that every adaptive system must have three roles. ISS theory supplies the stability template [Sontag, 1989, Dashkovskiy et al., 2010, Mironchenko and Prieur, 2020], while classical filtering provides the concrete Gaussian realization used here

[Jazwinski, 1970]. Once actions affect future observations, state update must solve a credit-assignment problem rather than a passive filtering problem.

Belief update semantics. The distinction between revision and update in dynamic worlds [Katsuno and Mendelzon, 1992] gives a formal language for the same issue. Bonanno [2024] makes explicit that action-conditioned belief update must distinguish exogenous change from state changes caused by the agent itself. We use this observation only to justify the presence of the action argument in the closed-loop update.

Algorithmic regulation and compression. Ruffini [2026] interprets good regulation through algorithmic compression of realized trajectories. This work adds a temporal constraint absent from that framework: even a useful compressed model becomes stale under drift. The cube-root law therefore complements the algorithmic-regulation viewpoint by quantifying how finite memory limits the duration for which compression remains valid in a changing environment.

Action-coupled prediction. The coercive masking regime is related to broader work on systems that change the data-generating process they seek to model. The claim is narrower than a general performative-prediction analysis: we isolate a diagnostic failure mode in which low observed error can be sustained by intervention, and we propose an effort-aware correction for that case. This connects most directly to performative prediction and feedback-loop analyses in recommendation and dataset construction, including [Perdomo et al., 2020, Hardt and Mendler-D"unner, 2023, Taori and Hashimoto, 2023, Pagan et al., 2023, Krauth et al., 2022, Eilat and Rosenfeld, 2023]. The main distinction is that we are not trying to solve the full causal problem of performative feedback; we are proposing a diagnostic for when intervention effort masks stale models in active or logged systems.

10 Limitations and Open Directions

Finite-memory scope. The cube-root law of proposition 6.12 is derived for the finite-memory estimator class analyzed here. The experiments cover sliding windows and closely related effective-memory schemes, but they do not amount to a universal minimax result over all adaptive estimators.

Geometry and observability. The current formulation uses W_2 and its debiased Sinkhorn approximation to measure model drift. Extending the same trade-off cleanly to other geometries, such as f -divergences or kernel MMD, remains open. Likewise, $\mathcal{C}^E$ requires an observable proxy for intervention effort, and the preferred proxy is domain-dependent. This is especially relevant in performative or feedback-loop settings, where intervention and exposure are partially confounded [Perdomo et al., 2020, Hardt and Mendler-D"unner, 2023, Taori and Hashimoto, 2023, Krauth et al., 2022].

Benchmark dependence. The passive benchmarks use Page-Hinkley events as the operational reference for lead time, and different failure definitions would change

the resulting alarm statistics. On active and logged benchmarks, the usefulness of $\mathcal{C}^E$ still depends on calibration choices such as the effort penalty λ , the warning threshold, and the local drift estimator $\hat{\zeta}_t$.

Effort proxies and σ_A surrogates. The practical hierarchy of effort proxies is strongest in simulation, weaker in direct observation, and weakest in logged recommendation data. Likewise, σ_A is exact only in settings with replicated or fully specified action channels. In the Gaussian tracker it is constructive, in the particle tracker it is a controlled surrogate based on actuation gap, and in KuaiRand it is a normalized exposure-entropy proxy rather than a direct estimate of $e^{-H(\mathcal{A}|\mathcal{P})}$. The logged benchmark should therefore be read as operational evidence for effort-aware diagnostics, not as full causal identification of intervention dynamics.

Scaling beyond current validation. The Gaussian and particle trackers validate the theory in controlled Gaussian and non-Gaussian settings, and the real-data benchmarks show that the diagnostics remain useful outside the synthetic environment. High-dimensional state spaces, richer action channels, and direct measurements of intervention effort are the next natural tests of the framework.

Opaque adaptation maps. For black-box Φ_t , the paper currently provides only the template requirement that an ISS-Lyapunov functional V_Φ exist. The CEGIS-based recovery idea mentioned in remark 6.17 is not implemented or validated in this repository, so diagnosing σ_Φ for opaque adaptation subsystems remains an open engineering problem.

Missing adaptive baselines. The current passive comparisons emphasize ADWIN, Page-Hinkley, KSWIN, and the NoDrift sanity baseline, and now also include CUSUM, forgetting-factor recursive least squares, and scalar Kalman predictors on the real-world streams. That does not exhaust the baseline space: additional representation- space or multivariate passive detectors remain future work, and the present calibration should still be read as benchmark-specific rather than universal.

11 Conclusion

Two conclusions organize the paper. First, tracking under drift with the finite-memory averaging class obeys a sharp limit. When total error is written as a sum of estimation variance and drift-induced staleness, optimizing over memory yields the cube-root law $\mathcal{E}_{\min} = \frac{3}{2}C_K^{2/3}\zeta^{1/3}$ for the estimator class analyzed here. In that sense, tracking under drift is constrained by the same object that makes online compression useful: memory. More memory reduces noise, but it also increases the age of the compressed model.

Second, action-coupled settings admit a specific diagnostic failure mode: apparent coherence can be maintained by intervention rather than by model adaptation. The decomposition into representation, actuation, and adaptation roles $(\mathcal{P}, \mathcal{A}, \Phi)$ separates that regime from ordinary degradation. Within this framing, $\mathcal{C}$ measures apparent coherence, while $\mathcal{C}^E$ and the masking gap measure how much of that apparent fit survives once intervention effort is charged explicitly.

The empirical results support this split. The Gaussian tracker recovers the predicted U-curve and the $\zeta^{1/3}$ scaling law for both sliding windows and EWMA in the controlled setting. The particle tracker shows that feedback masking can keep the standard score high even when the underlying model is stale. On passive benchmarks, generic changepoint detectors such as ADWIN remain the stronger default alarms. On active and logged benchmarks, however, $\mathcal{C}^E$ becomes the right tool because the failure mode is coercive masking rather than passive drift alone. The logged KuaiRand benchmark makes this point directly: the effort-corrected score detects a mode of feedback-induced degradation that the passive baselines do not capture under the same calibration.

Finite-memory tracking under drift obeys a structural coherence-delay trade-off. Closed-loop systems can hide model failure through feedback, and coercive masking is precisely the regime where effort-aware diagnostics add information beyond raw fit and passive drift alarms. How far this decomposition extends beyond the settings studied here remains open.

Acknowledgements

With gratitude to everyone who contributed ideas, feedback, tools, encouragement, and care along the way. This work stands on the generosity of colleagues, collaborators, open-source maintainers, and the broader research communities that made it possible.

References

- Jason Altschuler, Jonathan Weed, and Philippe Rigollet. Near-linear time approximation algorithms for optimal transport via Sinkhorn iteration. In *Advances in Neural Information Processing Systems 30 (NeurIPS)*, pages 1964–1974, 2017.
- Luigi Ambrosio, Nicola Gigli, and Giuseppe Savaré. *Gradient Flows in Metric Spaces and in the Space of Probability Measures*. Lectures in Mathematics ETH Zürich. Birkhäuser, Basel, 2005. doi: 10.1007/978-3-7643-8722-8.
- Omar Besbes, Yonatan Gur, and Assaf Zeevi. Non-stationary stochastic optimization. *Operations Research*, 63(5):1227–1244, 2015. doi: 10.1287/OPRE.2015.1408.
- Albert Bifet and Ricard Gavaldà. Learning from time-changing data with adaptive windowing. In *Proceedings of the 2007 SIAM International Conference on Data Mining (SDM)*, pages 443–448, 2007. doi: 10.1137/1.9781611972771.42.
- Giacomo Bonanno. A Kripke-Lewis semantics for belief update and belief revision. *Artificial Intelligence*, 328:104065, 2024. doi: 10.1016/j.artint.2023.104065.
- Wang Chi Cheung, David Simchi-Levi, and Ruihao Zhu. Learning to optimize under non-stationarity. In *Proceedings of the Twenty-Second International Conference on Artificial Intelligence and Statistics*, volume 89 of *Proceedings of Machine Learning Research*, pages 1079–1087. PMLR, 2019. URL <https://proceedings.mlr.press/v89/cheung19b.html>.

- Marco Cuturi. Sinkhorn distances: Lightspeed computation of optimal transport. In *Advances in Neural Information Processing Systems 26 (NeurIPS)*, pages 2292–2300, 2013.
- Sergey N. Dashkovskiy, Björn S. Ruffner, and Fabian R. Wirth. Small gain theorems for large scale systems and construction of ISS lyapunov functions. *SIAM Journal on Control and Optimization*, 48(6):4089–4118, 2010. doi: 10.1137/090746483.
- Itay Eilat and Nir Rosenfeld. Performative recommendation: Diversifying content via strategic incentives. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 9082–9103. PMLR, 2023. URL <https://proceedings.mlr.press/v202/eilat23a.html>.
- Kamil Faber, Roberto Corizzo, Bartłomiej Sniezynski, Michael Baron, and Nathalie Japkowicz. LIFEWATCH: Lifelong wasserstein change point detection. In *International Joint Conference on Neural Networks*, pages 1–8. IEEE, 2022a. doi: 10.1109/IJCNN55064.2022.9892891.
- Malte Faber, Magda Sakr, and Annette Bauer. WATCH: Wasserstein change point detection for high-dimensional time series data. In *2022 IEEE International Conference on Big Data*, pages 3268–3277, 2022b. doi: 10.1109/BigData55660.2022.10020525.
- João Gama, Pedro Medas, Gladys Castillo, and Pedro Rodrigues. Learning with drift detection. In *Advances in Artificial Intelligence — SBIA 2004*, volume 3171 of *Lecture Notes in Computer Science*, pages 286–295. Springer, 2004. doi: 10.1007/978-3-540-28645-5_29.
- João Gama, Indrė Žliobaitė, Albert Bifet, Mykola Pechenizkiy, and Abdelhamid Bouchachia. A survey on concept drift adaptation. *ACM Computing Surveys*, 46(4):44:1–44:37, 2014. doi: 10.1145/2523813.
- Aude Genevay, Lénaïc Chizat, Francis Bach, Marco Cuturi, and Gabriel Peyré. Sample complexity of Sinkhorn divergences. In *Proceedings of the 22nd International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 89 of *Proceedings of Machine Learning Research*, pages 1574–1583, 2019.
- Moritz Hardt and Celestine Mendler-Dünner. Performative prediction: Past and future. *CoRR*, abs/2310.16608, 2023. doi: 10.48550/arXiv.2310.16608.
- Michael Harries. SPLICE-2 comparative evaluation: Electricity pricing. Technical report, University of New South Wales, 1999. Available at <https://www.openml.org/d/151>.
- Andrew H. Jazwinski. *Stochastic Processes and Filtering Theory*, volume 64 of *Mathematics in Science and Engineering*. Academic Press, New York, 1970.
- Leslie Pack Kaelbling, Michael L. Littman, and Anthony R. Cassandra. Planning and acting in partially observable stochastic domains. *Artificial Intelligence*, 101(1–2):99–134, 1998. doi: 10.1016/S0004-3702(98)00023-X.

- Hirofumi Katsuno and Alberto O. Mendelzon. On the difference between updating a knowledge base and revising it. In Peter Gärdenfors, editor, *Belief Revision*, volume 29 of *Cambridge Tracts in Theoretical Computer Science*, pages 183–203. Cambridge University Press, 1992.
- Karl Krauth, Yixin Wang, and Michael I. Jordan. Breaking feedback loops in recommender systems with causal inference. *CoRR*, abs/2207.01616, 2022. doi: 10.48550/arXiv.2207.01616.
- Alessio Mazzetto and Eli Upfal. Nonparametric density estimation under distribution drift. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 24251–24270. PMLR, 2023. URL <https://proceedings.mlr.press/v202/mazzetto23a.html>.
- Andrii Mironchenko and Christophe Prieur. Input-to-state stability of infinite-dimensional systems: Recent results and open questions. *SIAM Review*, 62(3): 529–614, 2020. doi: 10.1137/19M1291248.
- Nicolò Pagan, Joachim Baumann, Ezzat Elokda, Giulia De Pasquale, Saverio Bolognani, and Anikó Hannák. A classification of feedback loops and their relation to biases in automated decision-making systems. In *Proceedings of the 3rd ACM Conference on Equity and Access in Algorithms, Mechanisms, and Optimization*, pages 7:1–7:14. ACM, 2023. doi: 10.1145/3617694.3623227.
- Juan Perdomo, Tijana Zrnic, Celestine Mendler-Dünnner, and Moritz Hardt. Performative prediction. In *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pages 7599–7609. PMLR, 2020. URL <https://proceedings.mlr.press/v119/perdomo20a.html>.
- Giulio Ruffini. The algorithmic regulator. *Entropy*, 28(3):257, 2026. doi: 10.3390/e28030257.
- Eduardo D. Sontag. Smooth stabilization implies coprime factorization. *IEEE Transactions on Automatic Control*, 34(4):435–443, 1989. doi: 10.1109/9.28018.
- Eduardo D. Sontag. On the input-to-state stability property. *European Journal of Control*, 1(1):24–36, 1995. doi: 10.1016/S0947-3580(95)70005-X.
- Rohan Taori and Tatsunori Hashimoto. Data feedback loops: Model-driven amplification of dataset biases. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 33883–33920. PMLR, 2023. URL <https://proceedings.mlr.press/v202/taori23a.html>.
- Jianjun Yuan and Andrew G. Lamperski. Trading-off static and dynamic regret in online least-squares and beyond. In *The Thirty-Fourth AAAI Conference on Artificial Intelligence*, pages 6712–6719. AAAI Press, 2020. doi: 10.1609/AAAI.V34I04.6149.

A Code Repository

All implementations are publicly available at:

<https://github.com/jovemexausto/coherence-delay-tradeoff>

The repository contains:

- `experiments/run.py` — single entry point for all experiment domains;
- `experiments/gaussian/` and `experiments/particle/` — core experiment logic;
- `experiments/bikes/`, `experiments/elec2/`, and `experiments/kuairand/` — domain-specific benchmarks;
- `REPRODUCIBILITY.md` — setup and run instructions.

All experiments use `uv` for dependency management (`uv sync`) and are reproducible with fixed random seeds as specified in each script.

B KuaiRand Logged Benchmark Protocol

This protocol instantiates the real-data masking benchmark on top of the KuaiRand sequential recommendation logs. The goal is to test whether an effort-corrected coherence score can surface exposure concentration and filter-bubble-like drift better than passive detectors and a smoothed-score baseline.

Data and signals

The primary state signal is `watch_ratio`. Auxiliary feedback signals include `click` and `like`; `dwel`, `follow`, and `share` are reserved for secondary validation when available. For user u at time t , let $P_{u,t}$ denote the inferred preference state, $A_{u,t}$ the exposure policy, and $\Phi_{u,t}$ the composite feedback signal. All user windows use `window_size = 20`, and video tags are binned from `video_features_basic_pure.csv`. The benchmark keeps only users with at least 20 interactions in the healthy, bubble, and collapse segments, yielding the 367-user evaluation cohort used in the main text.

Operational score components

In KuaiRand, σ_P is a normalized deviation of the current watch-ratio window from the healthy watch-ratio baseline, σ_Φ is a stability term based on within-window watch-ratio dispersion, and σ_A is an operational surrogate based on normalized exposure-tag entropy. This last term should be read as a logged-data proxy for action reproducibility, not as a direct estimator of $e^{-H(\mathcal{A}|\mathcal{P})}$, because the benchmark has no replicated policy logs.

The proxy hierarchy matters: if direct action logs or replica logs exist, the Gaussian-style exact interpretation can be approximated; if not, the logged-data entropy surrogate is the intended operational choice. We keep the report honest by treating the logged benchmark as evidence for a useful proxy, not as a proof that the surrogate equals the unobserved conditional entropy. In the ladder used throughout the paper, KuaiRand is a proxy-level benchmark.

Effort proxy

The effort proxy is defined as the divergence between the current exposure distribution and the healthy random-exposure baseline distribution:

$$E_{u,t} := D_{\text{KL}} \left(\pi_{u,t}^{\text{current}} \parallel \pi_{u,t}^{\text{healthy}} \right).$$

This choice is directional, decomposable, and directly interpretable as the cost of keeping the observed trajectory close to a healthy baseline. In settings where only aggregate exposure counts are available, total variation distance may be used as a fallback, but KL is the preferred default. We use additive smoothing 10^{-6} in the KL computation.

Protocol

The benchmark is executed in three phases.

1. **Healthy.** Exposure follows the random baseline, diversity is unconstrained, and effort remains low.
2. **Bubble.** Exposure becomes progressively concentrated, so that $E_{u,t}$ rises while standard coherence may remain artificially high.
3. **Collapse.** The concentrated exposure regime stops tracking the healthy baseline, generic detectors begin to fire, and the effort-corrected index should lead the baselines.

Evaluation

Evaluation is carried out sequentially within each user session and then batched across users with strict temporal splits. Thresholds for $\mathcal{C}$, $\mathcal{C}^E$, and EWMA-smoothed $\mathcal{C}^E$ are calibrated on healthy windows only, using the 20th percentile of the healthy score distribution. The reference effort scale E_0 is the median healthy KL divergence. Reported metrics are:

- detection rate of the bubble/collapse phases,
- median detection delay,
- false positives in the healthy phase,
- sensitivity to λ in $\mathcal{C}^E$,
- comparison against ADWIN, PageHinkley, KSWIN, NoDrift, and a smoothed $\mathcal{C}^E$ baseline (EWMA).

All baseline detectors consume the same effort-adjusted scalar signal, $1 - \mathcal{C}^E$, so the comparison is not biased by detector-specific inputs. The baseline hyperparameters are fixed a priori: ADWIN $\delta = 0.03$; PageHinkley $\delta = 0.005$, threshold = 20, $\alpha = 0.9999$; KSWIN window = 30, stat = 10, $\alpha = 0.001$. When enabled, thresholds are calibrated on healthy windows only.

Controls

Required controls include shuffled labels, random phase boundaries, no- coercion ablations, and removal of auxiliary feedback signals. These controls ensure that the benchmark does not reduce to a hindsight-tuned logged-bandit exercise.