

Useful Memory under Drift

Temporal-Validity Horizons for Finite-Memory Distribution Tracking

Vinicius Parreira 

Independent Researcher

figueiredo@pm.me

May 2026

Abstract

We investigate the fundamental limits of tracking an evolving distribution under temporal drift. Longer memory reduces finite-sample estimation noise but also introduces staleness through age mismatch with the present state. This balance is formalized by a finite temporal-validity horizon, which induces a near-optimal useful-memory region around the scale $n^*(a, H) \asymp (C_K/\zeta)^{1/(a+H)}$; the finite-sample rate exponent a and the temporal roughness exponent H govern the expansion and contraction of that region. Under explicit compact-support and fixed-span assumptions, a tractable one-dimensional proof model yields a root- n finite-sample term, with verifiable sufficient conditions for the integrated Bahadur remainder on a smooth deformation class. A Gaussian lower bound shows that the boundary is structural rather than estimator-specific, and the corresponding Gaussian location model admits a minimax benchmark theorem on deterministic Hölder paths. Beyond one dimension, sliced or coordinate-averaged Wasserstein gives a multivariate root- n transfer route under uniform projected control. Finally, empirical evidence confirms the predicted interior memory optima, roughness-dependent scaling, the same retention signature on a real stream, and a measurable lag between the loss of temporal validity and the activation of standard changepoint alarms.

Keywords: tracking under drift; temporal validity; finite-memory tracking; distribution tracking; Wasserstein distance; Sinkhorn divergence.

MSC 2020: 93D05, 93D30, 68P99.

1 Introduction

In stationary statistics, memory is an unambiguous resource: larger samples reduce variance because the target distribution does not change. Under drift, that logic fails. In streaming estimation, monitoring, and online adaptation, the evidence retained in memory was generated by distributions that may no longer match the current state. More memory therefore does two things at once: it reduces finite-sample noise, and it increases the age of the evidence used to estimate the current state.

That tension creates a horizon problem. Short windows are fresh but noisy. Long windows are stable but stale.

The temporal-validity horizon asks how long retained evidence remains useful for estimating the present before drift-induced misalignment overtakes the remaining variance gain.

A general horizon law captures this problem. The law takes the form

$$\mathbb{E} d(\hat{P}_t^{(n)}, P_t) \leq C_K n^{-a} + C_S \zeta n^H,$$

where $a > 0$ is the finite-sample rate exponent of the chosen geometry and estimator, $H \in (0, 1]$ is the temporal roughness exponent of the drift path, and ζ is the roughness amplitude. Balancing these terms yields the temporal-validity horizon

$$n^*(a, H) \asymp (C_K/\zeta)^{1/(a+H)}.$$

The horizon depends jointly on finite-sample geometry and on the temporal regularity with which drift accumulates staleness.

At any fixed tolerance $\delta > 0$, the same envelope induces a useful-memory region: the memory lengths whose total tracking error remains within $(1 + \delta)$ times the optimum. The horizon fixes the scale of that region, and its normalized shape depends only on the exponents a and H .

The tractable one-dimensional bounded-support fixed-span triangular-array model establishes the law through the quantile representation of W_2^2 , a triangular-array Bahadur expansion, and a design-effect comparison. In that regime, the root- n finite-sample rate

$$\mathbb{E} W_2(\hat{P}_n^{\text{tri}}, \bar{P}_n) = O(n^{-1/2}).$$

under explicit compact-support, fixed-span, and integrated-remainder hypotheses. This establishes the horizon law in a tractable one-dimensional bounded-support fixed-span model.

The structural lower bound shows that the horizon is imposed by the drift geometry itself. A Gaussian lower-bound construction sharpens the constant-level geometry, and in the Gaussian location model the same object yields a full minimax benchmark on deterministic Hölder paths.

For fixed- ε Sinkhorn geometry on an embedded fixed-span model, the proved structural ingredients yield pointwise and bandwise one-dimensional fixed- ε inheritance results on compact support, while the embedded low-dimensional extension isolates the remaining operational frontier.

Broader regular distribution families and operational geometries support different finite-sample rates, and therefore different members of the horizon family. Memory under drift is governed by the interaction between finite-sample geometry and temporal roughness. The compact-support one-dimensional proof model, the structural Gaussian lower bound, and the Gaussian location benchmark close the core law, while sliced or coordinate-averaged Wasserstein gives a multivariate root- n transfer route under uniform projected control, and fixed- ε Sinkhorn supplies an operational geometry in which the same horizon logic remains visible.

Useful memory is the near-optimal finite-memory region induced by the temporal-validity horizon. It is the set of memory lengths whose total tracking error remains within a fixed multiplicative tolerance of the optimum.

The finite-memory U-curve, roughness-dependent horizon scaling, the fixed- ε Sinkhorn evidence, and the measurable lag between validity loss and the first detector alarm are visible in the experiments. Temporal validity and changepoint evidence are different statistical objects: retained evidence can stop representing the present before change becomes detectable. The neighboring literatures on regret, window adaptation, smoothing, filtering, and transport geometry isolate separate pieces of that story; the present work identifies the time scale on which retained evidence remains temporally valid before age becomes the dominant source of error.

2 Useful Memory and Temporal-Validity Horizons




Figure 1: Finite-sample-staleness trade-off under drift. The finite-sample term falls with horizon, the staleness term rises with age, and total error is minimized at the temporal-validity horizon n^* . The green band marks a near-optimal useful-memory region, and the red band marks pre-detection validity loss: retained evidence can stop representing the present before enough changepoint evidence accumulates to trigger an alarm.

For a finite-memory estimator with lag weights $w_0, \dots, w_{n-1}$, the natural age variable is the effective lag

$$A(w) := \sum_{j=0}^{n-1} j w_j.$$

Its corresponding staleness at time t is the weighted trans-

port misalignment

$$S_t(w) := \sum_{j=0}^{n-1} w_j W_2(P_{t-j}, P_t).$$

For a uniform window, $A(w) = (n-1)/2$, so window length and temporal age are the same object up to constants.

Non-uniform weights separate two summaries. The finite-sample term is governed by the effective sample size

$$n_{\text{eff}}(w) := \left(\sum_{j=0}^{n-1} w_j^2 \right)^{-1},$$

while the drift term is governed by weighted lag moments. For each $H \in (0, 1]$, define

$$M_H(w) := \sum_{j=0}^{n-1} j^H w_j.$$

Then any weighted horizon law of the form

$$\mathbb{E} d(\hat{P}_t^{(w)}, P_t) \leq C_K n_{\text{eff}}(w)^{-a} + C_S \zeta M_H(w)$$

induces a weighted useful-memory region by balancing the two summaries. On a linear ramp, $M_1(w) = A(w)$ exactly, so the effective lag recovers the exact staleness term.

Proposition 2.1 (Uniform weights maximize effective sample size). *Let $w_0, \dots, w_{n-1}$ be nonnegative weights with $\sum_{j=0}^{n-1} w_j = 1$. Then*

$$n_{\text{eff}}(w) \leq n,$$

with equality if and only if $w_j = 1/n$ for all j . Moreover, if the drift profile is an exact linear ramp in the sense that the lag- j contribution is proportional to j , then the weighted staleness equals the effective lag:

$$S_t(w) = c A(w)$$

for the corresponding linear coefficient c .

Proof. By Cauchy-Schwarz,

$$1 = \left(\sum_{j=0}^{n-1} w_j \right)^2 \leq n \sum_{j=0}^{n-1} w_j^2,$$

so $n_{\text{eff}}(w) = 1 / \sum_{j=0}^{n-1} w_j^2 \leq n$, with equality exactly when the weights are all equal. On an exact linear ramp, the lag- j contribution is proportional to j , so averaging with weights w_j gives the weighted mean lag $A(w)$ and hence $S_t(w) = c A(w)$. $\square$

Throughout, P_t denotes the present target law, $\bar{P}_t^{(n)}$ the window target obtained by averaging the last n target laws, and $\hat{P}_t^{(n)}$ the estimator built from the last n observations. The finite-sample term always refers to the estimation error relative to $\bar{P}_t^{(n)}$, while the staleness term refers to the discrepancy between $\bar{P}_t^{(n)}$ and P_t .

Finite-sample error relative to $\bar{P}_t^{(n)}$ and staleness relative to P_t combine into a horizon law. When both sides are controlled, the same scale induces a useful-memory region: the range of memory lengths whose tracking error remains near optimal.

2.1 General Horizon Law

The main theorem separates finite-sample behavior from drift-induced staleness.

Theorem 2.2 (Abstract upper law). *Let d be a probability metric satisfying the triangle inequality. At time t , let P_t be the present target law, let $\bar{P}_t^{(n)}$ be the window target, and let $\hat{P}_t^{(n)}$ be the estimator built from the last n observations. Assume that for some constants $C_K > 0$, $C_S > 0$, exponent $a > 0$, roughness index $H \in (0, 1]$, and amplitude $\zeta > 0$,*

$$\mathbb{E} d(\hat{P}_t^{(n)}, \bar{P}_t^{(n)}) \leq C_K n^{-a}$$

and

$$d(\bar{P}_t^{(n)}, P_t) \leq C_S \zeta n^H$$

hold for all n in the regime of interest. Then

$$\mathbb{E} d(\hat{P}_t^{(n)}, P_t) \leq C_K n^{-a} + C_S \zeta n^H.$$

Proof. By the triangle inequality,

$$d(\hat{P}_t^{(n)}, P_t) \leq d(\hat{P}_t^{(n)}, \bar{P}_t^{(n)}) + d(\bar{P}_t^{(n)}, P_t).$$

Taking expectations and applying the two assumptions yields the claim. $\square$

The theorem isolates the structure of the problem: the path class supplies the staleness term, while the chosen geometry and sample regime supply the finite-sample term.

Proposition 2.3 (Triangular-array transfer for Hadamard-differentiable functionals). *Let $\mathcal{T}$ be a functional on probability laws that is Hadamard differentiable at P_t tangentially to a linear subspace of signed measures, with derivative $\dot{\mathcal{T}}_{P_t}$ bounded on the relevant tangent class. Suppose that for the triangular-array estimator and its window target, the base metric satisfies*

$$\mathbb{E} d(\hat{P}_t^{(n)}, \bar{P}_t^{(n)}) \leq C_K n^{-a}$$

and the drift envelope satisfies

$$d(\bar{P}_t^{(n)}, P_t) \leq C_S \zeta n^H.$$

If the remainder in the Hadamard expansion is $o(n^{-a})$ uniformly on the relevant class, then there exist constants $\tilde{C}_K, \tilde{C}_S > 0$ such that

$$\mathbb{E} \|\mathcal{T}(\hat{P}_t^{(n)}) - \mathcal{T}(P_t)\| \leq \tilde{C}_K n^{-a} + \tilde{C}_S \zeta n^H.$$

Proof. Write the Hadamard expansion at P_t and apply the boundedness of the derivative to the centered triangular-array fluctuation. The drift term contributes through the same envelope as the base metric because $\mathcal{T}$ is locally Lipschitz on the class. The remainder is lower order by assumption, so the transformed law inherits the same horizon envelope. $\square$

Proposition 2.4 (Finite-dimensional smooth transfer). *Let $\phi_1, \dots, \phi_k$ be bounded measurable functions and let $Z(P) = (\mathbb{E}_P[\phi_1], \dots, \mathbb{E}_P[\phi_k])$. Let $g \in C^2(U)$ on an open*

set $U \subset \mathbb{R}^k$ containing the relevant range of Z . Define $T(P) = g(Z(P))$. Suppose that on the triangular-array window target and estimator,

$$\mathbb{E} \|Z(\hat{P}_t^{(n)}) - Z(\bar{P}_t^{(n)})\| \leq C_K n^{-a}$$

and

$$\|Z(\bar{P}_t^{(n)}) - Z(P_t)\| \leq C_S \zeta n^H.$$

If the Hessian of g is bounded by M on the relevant range, then

$$\mathbb{E} |T(\hat{P}_t^{(n)}) - T(P_t)| \leq \tilde{C}_K n^{-a} + \tilde{C}_S \zeta n^H,$$

and the Taylor remainder is quadratic in the moment error:

$$\begin{aligned} |T(P') - T(P) - \nabla g(Z(P))^\top (Z(P') - Z(P))| \\ \leq \frac{M}{2} \|Z(P') - Z(P)\|^2. \end{aligned}$$

Proof. Apply the multivariate Taylor formula to g around $Z(P_t)$. The bounded Hessian gives the stated quadratic remainder. The first-order term is controlled by the assumed finite-sample and drift envelopes for Z , and the same decomposition as in proposition 2.3 yields the horizon law for T . $\square$

The same decomposition extends to non-uniform weights once the two summaries $n_{\text{eff}}(w)$ and $M_H(w)$ are controlled. The theorem-level question is whether weighted forgetting enlarges the useful-memory region beyond a constant factor on richer path classes, or whether it only changes the leading constant.

Corollary 2.5 (Temporal-validity horizon). *Under the assumptions of theorem 2.2, the upper envelope*

$$\Phi(n) := C_K n^{-a} + C_S \zeta n^H$$

is minimized, up to constant factors depending only on a , H , and C_S , at

$$n^*(a, H) \asymp (C_K/\zeta)^{1/(a+H)},$$

and the corresponding error scale is

$$\Phi(n^*) \asymp C_K^{H/(a+H)} \zeta^{a/(a+H)}.$$

Proof. Balance the decreasing finite-sample term against the increasing staleness term:

$$C_K n^{-a} \asymp C_S \zeta n^H.$$

This gives $n^{a+H} \asymp C_K/\zeta$, hence the displayed horizon law. Substituting that scale back into either term gives the optimal error scale. $\square$

Write the continuous optimizer of Φ as

$$n_* := \left(\frac{a C_K}{H C_S \zeta} \right)^{1/(a+H)}.$$

This optimizer differs from the scale $n^*(a, H)$ only by constants depending on a , H , and C_S , so the asymptotic horizon law remains unchanged.

2.2 Useful-Memory Region

The horizon scale induces a near-optimal region of memory lengths. For any $\delta > 0$, define the useful-memory region by

$$\mathcal{U}_\delta := \{n > 0 : \Phi(n) \leq (1 + \delta)\Phi(n_\star)\}.$$

This object records not only where the envelope is minimized, but also how much memory remains acceptable up to a fixed tolerance.

Proposition 2.6 (Useful-memory region at tolerance δ). *Let $\Phi(n) = C_K n^{-a} + C_S \zeta n^H$ with $a, H > 0$, and let $n_\star$ be its continuous optimizer. Then the normalized profile*

$$\Psi(x) := \frac{\Phi(n_\star x)}{\Phi(n_\star)} = \frac{Hx^{-a} + ax^H}{a + H}$$

is strictly decreasing on $(0, 1]$ and strictly increasing on $[1, \infty)$. Consequently,

$$\mathcal{U}_\delta = [n_-, n_+] = n_\star[x_-(\delta; a, H), x_+(\delta; a, H)],$$

where $x_-(\delta; a, H) \leq 1 \leq x_+(\delta; a, H)$ are the two solutions of $\Psi(x) = 1 + \delta$. In particular, the useful-memory region scales linearly with the horizon scale $n_\star$, and its relative shape depends only on a , H , and δ .

Proof. Since $n_\star$ minimizes Φ , differentiation gives

$$aC_K n_\star^{-a} = HC_S \zeta n_\star^H.$$

Substituting $n = n_\star x$ into Φ and dividing by $\Phi(n_\star)$ yields the displayed identity for Ψ . Differentiation then gives

$$\Psi'(x) = \frac{aH}{a + H} (x^{H-1} - x^{-a-1}),$$

which is negative for $x < 1$ and positive for $x > 1$. Since $\Psi(1) = 1$ and $\Psi(x) \rightarrow \infty$ as $x \rightarrow 0$ or $x \rightarrow \infty$, the sublevel set $\{x > 0 : \Psi(x) \leq 1 + \delta\}$ is a compact interval containing 1, which proves the representation of $\mathcal{U}_\delta$. $\square$

For small tolerance, the useful-memory region has the local expansion

$$x_\pm(\delta; a, H) = 1 \pm \sqrt{\frac{2\delta}{aH}} + o(\sqrt{\delta}), \quad \delta \downarrow 0,$$

because $\Psi''(1) = aH$. The local curvature therefore controls how quickly the near-optimal region contracts around the horizon.

The horizon depends jointly on the finite-sample rate exponent and the roughness exponent. Different geometries and different drift regularities induce different memory scales.

2.3 Drift Regularity and Staleness Accumulation

The staleness side of the law is controlled by a deterministic path class. For $H \in (0, 1]$ and $\zeta > 0$, write

$$\mathfrak{P}_H(\zeta) := \{(P_s)_{s \in \mathbb{Z}} : W_2(P_{t-j}, P_t) \leq \zeta j^H \text{ for all } j \geq 0\}.$$

The case $H = 1$ is the local Lipschitz envelope; smaller H correspond to slower accumulation of staleness with lag.

Lemma 2.7 (Hölder staleness bound via averaged couplings). *Let*

$$\bar{P}_t^{(n)} := \frac{1}{n} \sum_{j=0}^{n-1} P_{t-j}.$$

If the target path belongs to $\mathfrak{P}_H(\zeta)$, then

$$W_2^2(\bar{P}_t^{(n)}, P_t) \leq \frac{\zeta^2}{n} \sum_{j=0}^{n-1} j^{2H},$$

and therefore, by Jensen's inequality,

$$W_2(\bar{P}_t^{(n)}, P_t) \leq \zeta \left(\frac{1}{n} \sum_{j=0}^{n-1} j^{2H} \right)^{1/2} = C_{H,n} \zeta n^H,$$

where

$$C_{H,n} := \left(n^{-(2H+1)} \sum_{j=0}^{n-1} j^{2H} \right)^{1/2}.$$

In particular, $C_{H,n} \rightarrow (2H + 1)^{-1/2}$ as $n \rightarrow \infty$.

Proof. For each lag j , let π_j be an optimal coupling between P_{t-j} and P_t . Their average

$$\bar{\pi} := \frac{1}{n} \sum_{j=0}^{n-1} \pi_j$$

couples $\bar{P}_t^{(n)}$ with P_t , so

$$\begin{aligned} W_2^2(\bar{P}_t^{(n)}, P_t) &\leq \int \|x - y\|^2 d\bar{\pi}(x, y) \\ &= \frac{1}{n} \sum_{j=0}^{n-1} W_2^2(P_{t-j}, P_t). \end{aligned}$$

Applying the Hölder envelope gives the displayed finite- n constant, and taking square roots completes the proof. $\square$

Combining lemma 2.7 with corollary 2.5 yields the horizon family

$$n^*(a, H) \asymp (C_K/\zeta)^{1/(a+H)}.$$

In the root- n finite-sample regime $a = 1/2$,

$$n^*(1/2, H) \asymp (C_K/\zeta)^{2/(1+2H)}.$$

The finite- n constant in lemma 2.7 is the constant used here for uniform windows; the simpler $(2H + 1)^{-1/2}$ expression is its large-window limit.

2.4 One-Dimensional Proof Model for the Root- n Regime

The bounded-support fixed-span one-dimensional triangular array provides a tractable rigorous finite-sample instantiation of the abstract law. The model is built on common compact support, uniformly regular window mixtures, and an integrated Bahadur remainder hypothesis.

Proposition 2.8 (Root- n finite-sample bound in the 1D proof model). *Let $X_{n,1}, \dots, X_{n,n}$ be independent one-dimensional observations with laws $P_{n,1}, \dots, P_{n,n}$, and define the window mixture*

$$\bar{P}_n = \frac{1}{n} \sum_{j=1}^n P_{n,j}.$$

Assume:

- (i) *all $P_{n,j}$ are supported on a common compact interval $[L, U]$;*
- (ii) *the within-window drift span is uniformly bounded in n ;*
- (iii) *each mixture law $\bar{P}_n$ has a density $\bar{f}_n$ on $[L, U]$ that is bounded below by a positive constant and uniformly Hölder;*
- (iv) *writing q_n for the quantile function of $\bar{P}_n$ and $\hat{q}_n$ for the empirical quantile function of $\hat{P}_n^{\text{tri}}$, one has the Bahadur representation*

$$\hat{q}_n(u) - q_n(u) = \frac{u - \hat{F}_n(q_n(u))}{\bar{f}_n(q_n(u))} + R_n(u),$$

$$u \in (0, 1),$$

with

$$\mathbb{E} \int_0^1 R_n(u)^2 du = o(n^{-1}).$$

Then

$$\mathbb{E} W_2^2(\hat{P}_n^{\text{tri}}, \bar{P}_n) = O(n^{-1}),$$

and therefore, by Jensen's inequality,

$$\mathbb{E} W_2(\hat{P}_n^{\text{tri}}, \bar{P}_n) = O(n^{-1/2}).$$

Proof. In one dimension,

$$W_2^2(\hat{P}_n^{\text{tri}}, \bar{P}_n) = \int_0^1 (\hat{q}_n(u) - q_n(u))^2 du,$$

so the problem reduces to the integrated L_2 error of the empirical quantile process. By item (iv),

$$\hat{q}_n(u) - q_n(u) = \frac{u - \hat{F}_n(q_n(u))}{\bar{f}_n(q_n(u))} + R_n(u).$$

Define

$$Z_n(u) := \frac{u - \hat{F}_n(q_n(u))}{\bar{f}_n(q_n(u))}.$$

Then

$$\mathbb{E} \int_0^1 (\hat{q}_n(u) - q_n(u))^2 du = \mathbb{E} \int_0^1 (Z_n(u) + R_n(u))^2 du.$$

Let $m > 0$ be a uniform lower bound for $\bar{f}_n$. Since

$$\hat{F}_n(x) = \frac{1}{n} \sum_{j=1}^n \mathbf{1}\{X_{n,j} \leq x\},$$

independence across the row gives

$$\text{Var}(\hat{F}_n(q_n(u))) = \frac{1}{n^2} \sum_{j=1}^n p_{n,j}(u)(1 - p_{n,j}(u)) \leq \frac{1}{4n},$$

where $p_{n,j}(u) = P_{n,j}((-\infty, q_n(u)])$. Therefore

$$\mathbb{E} \int_0^1 Z_n(u)^2 du \leq m^{-2} \int_0^1 \text{Var}(\hat{F}_n(q_n(u))) du \leq \frac{1}{4m^2 n}.$$

By item (iv),

$$\mathbb{E} \int_0^1 R_n(u)^2 du = o(n^{-1}).$$

The cross term is then $o(n^{-1})$ by Cauchy–Schwarz, since $\mathbb{E} \int_0^1 Z_n(u)^2 du = O(n^{-1})$ and $\mathbb{E} \int_0^1 R_n(u)^2 du = o(n^{-1})$. Hence the expected squared Wasserstein error is $O(n^{-1})$, and Jensen's inequality gives the root- n rate for $\mathbb{E} W_2(\hat{P}_n^{\text{tri}}, \bar{P}_n)$.

The main conclusion is the exponent. Under this bounded-support fixed-span regime, the triangular-array finite-sample term remains root- n . $\square$

2.5 Verifiable Sufficient Conditions for the 1D Proof Model

An explicit deformation class makes the proof-model hypotheses verifiable. Let

$$g_\delta(x) = x + \delta x(1 - x), \quad x \in [0, 1], \quad |\delta| \leq \Delta < 1.$$

Then g_δ is an increasing C^2 diffeomorphism of $[0, 1]$, with

$$1 - \Delta \leq g'_\delta(x) \leq 1 + \Delta, \quad |g''_\delta(x)| \leq 2\Delta,$$

so the transformed densities satisfy

$$\frac{1}{1 + \Delta} \leq f_\delta(y) \leq \frac{1}{1 - \Delta},$$

and the same support is preserved.

Corollary 2.9 (Verifiable Bahadur remainder for the smooth deformation class). *Assume that each row law has the form*

$$P_{n,j} = (g_{\delta_{n,j}})_\# \text{Unif}[0, 1], \quad |\delta_{n,j}| \leq \Delta < 1,$$

and that the within-window span is uniformly bounded in n . Then item (iv) of proposition 2.8 holds. More precisely, there exists a constant $C_\Delta < \infty$ such that

$$\mathbb{E} \int_0^1 R_n(u)^2 du \leq C_\Delta n^{-2}.$$

Consequently,

$$\mathbb{E} W_2(\hat{P}_n^{\text{tri}}, \bar{P}_n) = O(n^{-1/2})$$

on this deformation class.

Proof. Write $\bar{F}_n$ and q_n for the mixture cdf and quantile. Since each $g_{\delta_{n,j}}$ is an increasing diffeomorphism of $[0, 1]$, every mixture law $\bar{P}_n$ is supported on $[0, 1]$ and has density $\bar{f}_n = \frac{1}{n} \sum_{j=1}^n f_{\delta_{n,j}}$ satisfying the explicit bounds

$$\frac{1}{1 + \Delta} \leq \bar{f}_n(y) \leq \frac{1}{1 - \Delta}, \quad \|\bar{f}'_n\|_\infty \leq \frac{2\Delta}{(1 - \Delta)^3}.$$

Hence the inverse map $\bar{F}_n^{-1}$ is uniformly twice differentiable on $(0, 1)$, with second derivative bounded by a constant depending only on Δ . Applying the second-order Taylor expansion of $\bar{F}_n^{-1}$ around $u = \bar{F}_n(q_n(u))$ gives

$$\hat{q}_n(u) - q_n(u) = \frac{u - \hat{F}_n(q_n(u))}{\hat{f}_n(q_n(u))} + R_n(u),$$

with

$$|R_n(u)| \leq C_\Delta |\hat{F}_n(q_n(u)) - u|^2.$$

Now

$$\hat{F}_n(q_n(u)) = \frac{1}{n} \sum_{j=1}^n \mathbf{1}\{X_{n,j} \leq q_n(u)\}$$

is an average of independent Bernoulli variables with mean u . Uniformly in u , its fourth centered moment is $O(n^{-2})$. Therefore

$$\mathbb{E} R_n(u)^2 \leq C_\Delta \mathbb{E} |\hat{F}_n(q_n(u)) - u|^4 \leq C'_\Delta n^{-2},$$

and integrating over $u \in (0, 1)$ yields

$$\mathbb{E} \int_0^1 R_n(u)^2 du \leq C'_\Delta n^{-2} = o(n^{-1}).$$

The root- n claim then follows from proposition 2.8. The same smoothness package is consistent with the classical Bahadur theory of Bahadur [1966] and van der Vaart [1998, Chapter 21], but here the explicit deformation bounds make the integrated remainder quantitative. $\square$

Remark 2.10 (Weak rowwise dependence). The proof of proposition 2.8 uses independence within a row only through the $O(n^{-1})$ variance bound for the empirical cdf term and through item (iv). Consequently, the same root- n conclusion persists under any short-range dependent rowwise model for which the empirical cdf still has L_2 fluctuation of order $n^{-1/2}$ and the integrated Bahadur remainder remains $o(n^{-1})$. Short-range mixing conditions are one natural route to such a package, but the present paper does not develop a general dependent-row theorem.

On the tested smooth fixed-support family, the integrated residual decays at empirical rates between 1.40 and 1.54 for spans 0.2 and 0.4, and the scaled quantity $n \int_0^1 R_n(u)^2 du$ still decays at rates between 0.40 and 0.54; see appendix B. On the fixed-span triangular-array uniform-kernel model, the residual-to-MSE ratio drops from about 0.38 at $n = 25$ to about 0.11 at $n = 400$, with integrated residual rates 1.43 for the triangular sample and 1.49 for the i.i.d. mixture benchmark. These diagnostics do not replace a general verification theorem, but they now sit behind an explicit sufficient-condition corollary on the smooth deformation class and quantify span dependence on the classes currently used in the manuscript.

2.6 Structural Lower Bound in the Root- n Regime

The upper law would be substantially weaker if it were only an estimator-specific statement. The lower bound shows that, in the root- n finite-sample regime $a = 1/2$, the horizon is imposed by the drift geometry itself.




Figure 2: Gaussian lower-bound construction. The two paths share most of the retained history and differ only on a recent segment of width h , yet remain separated at the present. That recent segment already produces the same horizon exponent as the upper law in the root- n regime.

Proposition 2.11 (Structural lower bound in the root- n regime). *Fix $H \in (0, 1]$ and work on the Gaussian location subclass with known variance σ^2 and path pair*

$$\mu_{t-j}^{\pm,h} = \pm \beta(h-j)_+^H, \quad j = 0, 1, \dots, h,$$

where $\beta \leq \zeta$ enforces the roughness budget. Let $\mathcal{G}_{H,\zeta}^{\text{Gauss}}$ denote the corresponding Gaussian location subclass of target paths, and suppose the retained observations are independent with laws $N(\mu_{t-j}^{\pm,h}, \sigma^2)$. Then there exists a constant $c_H > 0$ depending only on H such that the endpoint tracking risk over this subclass satisfies

$$\inf_{\hat{P}} \sup_{P \in \mathcal{G}_{H,\zeta}^{\text{Gauss}}} \mathbb{E} W_2(\hat{P}_t, P_t) \geq c_H \sigma^{2H/(2H+1)} \zeta^{1/(2H+1)}.$$

Consequently, in the root- n finite-sample regime $a = 1/2$, the horizon scale

$$h^*(H) \asymp (\sigma/\zeta)^{2/(2H+1)}$$

is structural even on this restricted subclass.

Proof. Compare two hypotheses generated by the path pair $\mu_{t-j}^{+,h}$ and $\mu_{t-j}^{-,h}$. At the present time, the endpoint separation is of order βh^H . On the Gaussian location subclass, this is also the W_2 separation. The Kullback–Leibler divergence between the two hypotheses is proportional to

$$\beta^2 \sum_{r=1}^h r^{2H} / \sigma^2.$$

Choosing

$$\beta = \min \left\{ \zeta, \frac{\sigma}{2\sqrt{\sum_{r=1}^h r^{2H}}} \right\}$$

keeps the testing problem in the Le Cam regime while respecting the roughness budget. The resulting two-point lower bound is of order βh^H . Balancing the active roughness constraint against the testing constraint gives $h^*(H) \asymp (\sigma/\zeta)^{2/(2H+1)}$ and therefore the displayed lower law. The exact Gaussian two-point calculation used here is recorded in appendix A.

The lower bound is subclass based and not class-tight, but it is enough to show that the horizon is a structural feature of the problem and not an estimator artifact.

The proposition above is the lower-bound fact needed for the main horizon law. The same Gaussian construction also exposes the constant-level geometry. The following proposition records the sharp Gaussian constants available within that proof scheme.

2.7 Gaussian Lower-Bound Geometry

Let Φ and φ denote the standard normal distribution function and density.

Proposition 2.12 (Gaussian lower-bound constants). *Fix $H \in (0, 1]$ and set $p_H := 2H/(2H+1)$. Let $x_H > 0$ be the unique maximizer of $x^{p_H}\Phi(-x)$, equivalently the unique solution of*

$$p_H \Phi(-x_H) = x_H \varphi(x_H).$$

For the Gaussian ramp profile $\mu_{t-j}^{\pm, h} = \pm \zeta(h-j)_+^H$, let

$$L_h^{\text{ramp}}(\sigma, \zeta) := \zeta h^H \Phi \left(-\frac{\zeta}{\sigma} \sqrt{\sum_{r=1}^h r^{2H}} \right).$$

Then, as $\sigma/\zeta \rightarrow \infty$,

$$\sup_{h \geq 1} L_h^{\text{ramp}}(\sigma, \zeta) \sim C_H^{\text{ramp}} \sigma^{2H/(2H+1)} \zeta^{1/(2H+1)},$$

with

$$C_H^{\text{ramp}} = (2H+1)^{H/(2H+1)} x_H^{2H/(2H+1)} \Phi(-x_H).$$

Among all endpoint-saturating Hölder profiles, the profile

$$g_r^{\min} = h^H - (h-r)^H$$

minimizes the testing energy and has asymptotic energy constant

$$I_H = \int_0^1 (1 - (1-u)^H)^2 du = \frac{2H^2}{(H+1)(2H+1)}.$$

If

$$L_h^{\min}(\sigma, \zeta) := \zeta h^H \Phi \left(-\frac{\zeta}{\sigma} \sqrt{\sum_{r=1}^h (g_r^{\min})^2} \right),$$

then, as $\sigma/\zeta \rightarrow \infty$,

$$\sup_{h \geq 1} L_h^{\min}(\sigma, \zeta) \sim C_H^{\min} \sigma^{2H/(2H+1)} \zeta^{1/(2H+1)},$$

with

$$C_H^{\min} = I_H^{-H/(2H+1)} x_H^{2H/(2H+1)} \Phi(-x_H).$$

For $H < 1$, this strictly improves the ramp constant by the factor

$$\frac{C_H^{\min}}{C_H^{\text{ramp}}} = \left(\frac{H+1}{2H^2} \right)^{H/(2H+1)} > 1,$$

while at $H = 1$ the two profiles coincide.

Proof. The exact Gaussian testing problem reduces the two-point lower bound to the displayed quantities L_h^{ramp} and $L_h^{\min}$. After rescaling by $h = A(\sigma/\zeta)^{2/(2H+1)}$, the ramp energy and endpoint-minimal energy have limits

$(2H+1)^{-1}$ and I_H , respectively. The corresponding normalized payoff becomes a scalar multiple of $x^{p_H}\Phi(-x)$, whose unique maximizer is x_H . The endpoint-minimal profile minimizes testing energy among all endpoint-saturating Hölder profiles, so it yields the sharper constant. Full details are given in appendix A. $\square$

These constants sharpen the lower-bound picture without changing the horizon exponents or the status of the main results.

The same lower-bound geometry also yields a full class-level benchmark once the model is specialized to Gaussian location tracking. In that setting, the deterministic Hölder envelope, the finite-memory estimator, and the lower-bound construction all live in the same one-dimensional parameterization, so this model admits a direct minimax rate statement.

Theorem 2.13 (Gaussian location minimax benchmark on the deterministic Hölder class). *Fix $H \in (0, 1]$, $\zeta > 0$, and $\sigma > 0$. Let $\mathfrak{M}_H(\zeta)$ denote the class of deterministic mean paths $(\mu_s)_{s \in \mathbb{Z}}$ such that*

$$|\mu_{t-j} - \mu_t| \leq \zeta j^H \quad \text{for all } j \geq 0.$$

Suppose the observations are independent with

$$Y_{t-j} \sim N(\mu_{t-j}, \sigma^2), \quad j = 0, 1, \dots, n-1,$$

and let the target law at time t be $P_t = N(\mu_t, \sigma^2)$. Define the optimal finite-memory minimax risk

$$\mathcal{R}_H^*(\sigma, \zeta) := \inf_{n \geq 1} \inf_{\hat{P}_t^{(n)}} \sup_{\mu \in \mathfrak{M}_H(\zeta)} \mathbb{E} W_2(\hat{P}_t^{(n)}, P_t).$$

Then there exist constants $0 < c_H \leq C_H < \infty$, depending only on H , such that

$$\begin{aligned} c_H \sigma^{2H/(2H+1)} \zeta^{1/(2H+1)} &\leq \mathcal{R}_H^*(\sigma, \zeta) \\ &\leq C_H \sigma^{2H/(2H+1)} \zeta^{1/(2H+1)}. \end{aligned}$$

In particular, the deterministic Hölder class has horizon scale

$$n_H^* \asymp (\sigma/\zeta)^{2/(2H+1)}.$$

Moreover, the uniform-window mean estimator has an exact asymptotic upper constant on this benchmark. If $y_H > 0$ solves

$$Hy_H(2\Phi(y_H) - 1) = \varphi(y_H),$$

then its asymptotic horizon shape is

$$A_H^{\text{mean}} = ((H+1)y_H)^{2/(2H+1)}$$

and its asymptotic risk constant is

$$\begin{aligned} U_H^{\text{mean}} &= ((H+1)y_H)^{-1/(2H+1)} \\ &\quad \times \left(2\varphi(y_H) + y_H(2\Phi(y_H) - 1) \right). \end{aligned}$$

Proof. For the upper bound, estimate the present mean by the uniform-window average

$$\hat{\mu}_t^{(n)} := \frac{1}{n} \sum_{j=0}^{n-1} Y_{t-j}$$

and set $\hat{P}_t^{(n)} = N(\hat{\mu}_t^{(n)}, \sigma^2)$. Because equal-variance Gaussian laws satisfy

$$W_2(N(m_1, \sigma^2), N(m_2, \sigma^2)) = |m_1 - m_2|,$$

the risk reduces to mean estimation error. The noise term is $\sigma\sqrt{2/(\pi n)}$, while the deterministic bias is bounded by

$$\frac{\zeta}{n} \sum_{j=0}^{n-1} j^H,$$

so

$$\sup_{\mu \in \mathfrak{M}_H(\zeta)} \mathbb{E} W_2(\hat{P}_t^{(n)}, P_t) \leq \sigma\sqrt{\frac{2}{\pi n}} + \frac{\zeta}{n} \sum_{j=0}^{n-1} j^H.$$

Rescaling by $n = A(\sigma/\zeta)^{2/(2H+1)}$ yields a one-parameter normalized upper envelope. Differentiating its exact Gaussian-location form gives the stated root equation for y_H , and therefore the displayed asymptotic upper constant and horizon shape. In particular, this proves the stated upper rate and horizon scale.

For the lower bound, use the symmetric two-point sub-family

$$\mu_{t-j}^{\pm, h} = \pm \beta(h^H - j^H)_+, \quad j \geq 0,$$

with $\beta \leq \zeta$. This family lies in $\mathfrak{M}_H(\zeta)$ because

$$|\mu_t^{\pm, h} - \mu_{t-j}^{\pm, h}| = \beta \min(h^H, j^H) \leq \zeta j^H.$$

Moreover, for fixed endpoint value βh^H , it minimizes the Gaussian testing energy over all symmetric lag profiles allowed by the present-point Hölder envelope, since each lag coordinate is independently minimized by $\beta(h^H - j^H)_+$. The exact Gaussian two-point calculation from the lower-bound analysis in appendix A therefore applies on the full deterministic Hölder class and yields a lower bound of order

$$\sigma^{2H/(2H+1)} \zeta^{1/(2H+1)}.$$

Minimizing over h gives the same horizon scale $(\sigma/\zeta)^{2/(2H+1)}$. Combining the two sides proves the theorem. $\square$

2.8 Multivariate Extension via Sliced and Product Geometries

The horizon law also extends beyond one dimension when the multivariate geometry reduces to a uniformly controlled family of one-dimensional transport problems. This is the natural setting for sliced or coordinate-averaged Wasserstein constructions, where the finite-sample term is built by integrating projected one-dimensional errors.

Corollary 2.14 (Sliced-Wasserstein transfer under uniform projected control). *Let ρ be a probability measure on a family of one-dimensional projections θ , and define*

$$SW_{2,\rho}^2(P, Q) := \int W_2^2(P_\theta, Q_\theta) \rho(d\theta),$$

where P_θ denotes the pushforward of P under projection θ . Assume that for every θ in the support of ρ :

- (i) *the projected triangular-array window satisfies the hypotheses of proposition 2.8 with constants uniform in θ ;*
- (ii) *the projected drift path obeys the same Hölder staleness envelope with exponent H and amplitude at most ζ uniformly in θ .*

Then there exist constants $C_K, C_S < \infty$ such that

$$\mathbb{E} SW_{2,\rho}(\hat{P}_t^{(n)}, P_t) \leq C_K n^{-1/2} + C_S \zeta n^H.$$

Consequently,

$$n^*(H) \asymp (C_K/\zeta)^{1/(H+1/2)}.$$

Coordinate-averaged product constructions are included as the special case in which ρ is supported on coordinate projections.

Proof. For each projection θ , proposition 2.8 gives

$$\mathbb{E} W_2^2(\hat{P}_{t,\theta}^{(n)}, \bar{P}_{t,\theta}^{(n)}) \leq C n^{-1}$$

with C uniform in θ . Integrating over ρ yields

$$\mathbb{E} SW_{2,\rho}^2(\hat{P}_t^{(n)}, \bar{P}_t^{(n)}) \leq C n^{-1},$$

and Jensen gives the root- n bound for $\mathbb{E} SW_{2,\rho}(\hat{P}_t^{(n)}, \bar{P}_t^{(n)})$. The projected Hölder bound integrates in the same way, giving

$$SW_{2,\rho}(\bar{P}_t^{(n)}, P_t) \leq C_S \zeta n^H.$$

Applying theorem 2.2 with $d = SW_{2,\rho}$ yields the stated envelope and horizon scale. $\square$

2.9 Benchmark Theorems and Operational Extensions

The Gaussian benchmark and the fixed- ε Sinkhorn extension place the horizon law in two broader settings: a full path-class benchmark with theorem-level closure, and an operational transport geometry in which the one-dimensional compact-support inheritance theory is already explicit. The linear-Gaussian entropic-OT setting suggests an intermediate recursive regime between the i.i.d. fixed- ε theory and the embedded low-dimensional extension [Kalman, 1960, Jazwinski, 1970, Akyildiz et al., 2024].

Fixed- ε Sinkhorn setting. For fixed- ε debiased Sinkhorn, the central question is whether a practical geometry can supply a stable finite-sample term on the embedded fixed-span model. The target law is

$$\mathbb{E} S_\varepsilon(\hat{P}_{t,\text{tri}}^{(n)}, \bar{P}_t^{(n)}) \leq C_\varepsilon n^{-a_\varepsilon}$$

under bounded support, fixed span, and structural low-dimensionality, where S_ε denotes a fixed- ε Sinkhorn-type geometry and the exponent a_ε matches the corresponding i.i.d. mixture benchmark.

Definition 2.15 (Fixed- ε Sinkhorn setting). Fix a bounded embedded fixed-span model, an intrinsic complexity index k , a dual Hölder smoothness level α , and an operational band $\mathcal{B} \subset (0, \infty)$ of regularization values. We say that this model is in the fixed- ε Sinkhorn setting if:

- (i) the i.i.d. fixed- ε benchmark on the same support class admits a finite-sample bound with exponent a_ε on $\mathcal{B}$;
- (ii) the triangular-window and i.i.d.-mixture supports have identical operational covering complexity on $\mathcal{B}$;
- (iii) the dual class lies in the parametric adaptation region $\alpha > k/2$.

Proposition 2.16 (Structural ingredients for fixed- ε Sinkhorn). *On the bounded embedded fixed-span model with squared-Euclidean cost, the first three items in definition 2.15 are structural: the support-complexity inheritance is exact, and the dual-complexity condition is precisely the LCA threshold $\alpha > k/2$.*

Proposition 2.17 (Pointwise fixed- ε inheritance in the 1D proof model). *Assume the hypotheses of proposition 2.8, let $\hat{P}_n^{\text{iid}}$ be the empirical measure of n i.i.d. samples from $\bar{P}_n$, and fix $\varepsilon > 0$. For the debiased Sinkhorn divergence S_ε with squared-Euclidean cost on the common compact support, one has*

$$\mathbb{E} S_\varepsilon(\hat{P}_n^{\text{tri}}, \bar{P}_n) = O(n^{-1/2}).$$

Proof. Write $V_\varepsilon(\mu, \nu)$ for the regularized OT value so that

$$S_\varepsilon(\mu, \nu) = V_\varepsilon(\mu, \nu) - \frac{1}{2} V_\varepsilon(\mu, \mu) - \frac{1}{2} V_\varepsilon(\nu, \nu).$$

Because all measures are supported on one compact interval, they lie in a common bounded quadratic-moment class. By Eckstein and Nutz [2022, Thm. 3.7(ii), Ex. 3.4], for fixed ε there exists $L_\varepsilon < \infty$ depending only on that class and on the cost such that

$$|V_\varepsilon(\mu, \nu) - V_\varepsilon(\mu', \nu')| \leq L_\varepsilon (W_2(\mu, \mu') + W_2(\nu, \nu')).$$

Applying this bound to the cross term and both self terms gives

$$|S_\varepsilon(\mu, \nu) - S_\varepsilon(\mu', \nu')| \leq 2L_\varepsilon (W_2(\mu, \mu') + W_2(\nu, \nu')).$$

With $\nu = \nu' = \bar{P}_n$,

$$S_\varepsilon(\hat{P}_n^{\text{tri}}, \bar{P}_n) \leq S_\varepsilon(\hat{P}_n^{\text{iid}}, \bar{P}_n) + 2L_\varepsilon W_2(\hat{P}_n^{\text{tri}}, \hat{P}_n^{\text{iid}}).$$

Taking expectations and using the triangle inequality gives

$$\mathbb{E} W_2(\hat{P}_n^{\text{tri}}, \hat{P}_n^{\text{iid}}) \leq \mathbb{E} W_2(\hat{P}_n^{\text{tri}}, \bar{P}_n) + \mathbb{E} W_2(\hat{P}_n^{\text{iid}}, \bar{P}_n).$$

The first term is $O(n^{-1/2})$ by proposition 2.8. The second is the classical i.i.d. one-dimensional quantile-process rate on the same compact support class [Bahadur, 1966, van der Vaart, 1998], so it is also $O(n^{-1/2})$. Hence

$$\mathbb{E} W_2(\hat{P}_n^{\text{tri}}, \hat{P}_n^{\text{iid}}) = O(n^{-1/2}).$$

Finally, under the null comparison with fixed ε and compact support, the i.i.d. debiased Sinkhorn term satisfies

$$\mathbb{E} S_\varepsilon(\hat{P}_n^{\text{iid}}, \bar{P}_n) = O(n^{-1})$$

by del Barrio et al. [2023, Thm. 5.1, after fixed- ε rescaling]; see also the null-limit law in Goldfeld et al. [2024]. Combining the bounds yields the displayed $O(n^{-1/2})$ rate. $\square$

Proposition 2.18 (Bandwise fixed- ε inheritance in the 1D proof model). *Let $\mathcal{B} = [\varepsilon_{\min}, \varepsilon_{\max}] \subset (0, \infty)$ be a compact regularization band. Under the hypotheses of proposition 2.17 on a common compact support class, there exists $C_{\mathcal{B}} < \infty$ such that*

$$\sup_{\varepsilon \in \mathcal{B}} \mathbb{E} S_\varepsilon(\hat{P}_n^{\text{tri}}, \bar{P}_n) \leq C_{\mathcal{B}} n^{-1/2}.$$

Consequently, the induced horizon scale is uniform on $\mathcal{B}$:

$$\sup_{\varepsilon \in \mathcal{B}} n_\varepsilon^*(H) \asymp (C_{\mathcal{B}}/\zeta)^{1/(H+1/2)}.$$

Proof. Repeat the proof of proposition 2.17. The stability constant L_ε from Eckstein and Nutz [2022, Thm. 3.7(ii), Ex. 3.4] is bounded on $\mathcal{B}$ because all measures live in one compact support class and $\varepsilon \geq \varepsilon_{\min} > 0$. The i.i.d. debiased-Sinkhorn benchmark on compact support is also uniform on $\mathcal{B}$ by the fixed- ε asymptotics of Goldfeld et al. [2024] and the compact-support null bounds of del Barrio et al. [2023, Thm. 5.1, after fixed- ε rescaling]. Since the triangular-to-i.i.d. W_2 comparison is independent of ε , the constants in the pointwise proof can be chosen uniformly over $\mathcal{B}$. $\square$

Proposition 2.19 (Conditional embedded fixed- ε band inheritance). *Assume the fixed- ε Sinkhorn setting from definition 2.15. Suppose that on a compact band $\mathcal{B} = [\varepsilon_{\min}, \varepsilon_{\max}]$, the centered Sinkhorn scaling map admits a uniform linearization with remainder $o(n^{-1/2})$ and a uniform self-coupling inverse bound. Then there exists $C_{\mathcal{B}} < \infty$ such that*

$$\sup_{\varepsilon \in \mathcal{B}} \mathbb{E} S_\varepsilon(\hat{P}_{t,\text{tri}}^{(n)}, \bar{P}_t^{(n)}) \leq C_{\mathcal{B}} n^{-1/2}.$$

Consequently, on any such band the induced horizon scale satisfies

$$\sup_{\varepsilon \in \mathcal{B}} n_\varepsilon^*(H) \asymp (C_{\mathcal{B}}/\zeta)^{1/(H+1/2)}.$$

Proof. The linearization hypothesis reduces the triangular-array Sinkhorn divergence to its first-order influence term plus an $o(n^{-1/2})$ remainder. The uniform self-coupling inverse bound keeps the influence operator bounded across $\mathcal{B}$, while the triangular-array fluctuation of the empirical measure is root- n on the embedded fixed-span class. The same balance then gives the uniform $n^{-1/2}$ bound and the stated horizon scale. $\square$

Conjecture 2.20 (Bandwise fixed- ε Sinkhorn horizon inheritance). *Beyond the one-dimensional fixed- ε band result above, assume the fixed- ε Sinkhorn setting from definition 2.15. Then, on bands where the i.i.d. benchmark*

stays in the parametric Sinkhorn regime, the triangular-array finite-sample term should satisfy

$$\mathbb{E} S_\varepsilon \left(\hat{P}_{t,\text{tri}}^{(n)}, \bar{P}_t^{(n)} \right) \leq C_\varepsilon n^{-1/2}, \quad \varepsilon \in \mathcal{B},$$

and the corresponding horizon scale should satisfy

$$n_\varepsilon^*(H) \asymp (C_\varepsilon/\zeta)^{1/(H+1/2)}.$$

More generally, if the i.i.d. benchmark on $\mathcal{B}$ has exponent a_ε , then the operational horizon law should be

$$n_\varepsilon^*(a_\varepsilon, H) \asymp (C_\varepsilon/\zeta)^{1/(a_\varepsilon+H)}.$$

Evidence. The i.i.d. benchmark comes from existing fixed- ε entropic transport results, including fixed-marginal limit theory for the Sinkhorn divergence and intrinsic-complexity adaptation [Goldfeld et al., 2024, Stromme, 2023, Groppe and Hundrieser, 2024]. In the present embedded fixed-span model, support-complexity inheritance is exact and the dual smoothness threshold reduces to $\alpha > k/2$. Calibration then asks whether triangular and i.i.d. rate estimates remain close across the regularization band. Additional one-dimensional smooth fixed-support diagnostics show that triangular and i.i.d. fitted exponents can agree closely on moderate fixed- ε bands, while direct Jacobian diagnostics concentrate the remaining difficulty on uniform control of the self-coupling derivative blocks across nontrivial bands and higher-dimensional embedded classes.

Regular dominated families remain the natural family-level extension class. Local testing geometry and metric comparability suggest the same horizon logic with a modified finite-sample exponent, whereas support-changing and singular families fall outside that regime.

Once a finite-sample regime and a roughness level are specified, the law yields an explicit memory scale and an optimal error scale. Rougher drift and weaker finite-sample rates both force shorter memory and larger residual error.

3 Empirical Signatures

These experiments expose the finite-sample signatures of the temporal-validity horizon, the useful-memory region induced by that horizon, and evidence relevant to the fixed- ε Sinkhorn frontier.

The Gaussian U-curve uses 20 Monte Carlo seeds, the roughness-family sweep uses 12, the fixed- ε Sinkhorn fits log-log slopes over sample sizes (24, 48, 96, 160) on the pairs (8, 1), (8, 2), (12, 1), and (12, 2) using 24 seeds, and the validity-detection experiment uses a piecewise-ramped drift stream with 12 seeds, a fixed operating window, and both ADWIN and Page-Hinkley detectors. The real-stream diagnostic uses the ELEC2 stream and a short anchor block on the marginal distribution of `nswprice`. The empirical goal is qualitative and structural: persistence of shape, ordering, and sign across repeated runs, not precise constant estimation for every curve.

The empirical signatures are fivefold. Finite-memory risk exhibits an interior optimum together with a near-optimal useful-memory band. That optimum moves with

temporal roughness. The same shape appears on a real stream. The fixed- ε Sinkhorn frontier receives structured empirical support. Temporal validity can fail before changepoint evidence becomes statistically visible.

3.1 Interior Memory Optimum

The Gaussian horizon sweep isolates the variance-staleness trade-off directly. Small windows are variance dominated, large windows are stale, and the minimum lies at an interior horizon.




Figure 3: Interior risk optimum over memory length. Tracking error first falls with horizon and then rises once staleness overtakes the remaining variance gain. The optimum remains finite across drift strengths.

This is the basic empirical signature of temporal validity under drift. Under stationarity, more memory would continue to help. Under drift, the same horizon that reduces variance also ages the evidence. Around the minimizing horizon, the curve displays a flat near-optimal region, which is the empirical analogue of the useful-memory region defined from the envelope Φ .

3.2 Roughness-Dependent Horizon Scaling

The roughness-family experiment varies the roughness index directly. For each Hölder-type regime, the surrogate error obeys the same structural decomposition as the theory,

$$\mathcal{E}(n) \approx \sigma n^{-1/2} + \zeta n^H,$$

and the experiment records the empirically optimal horizon across a sweep of ζ values.




Figure 4: Roughness-dependent horizon scaling. Each path class exhibits a finite optimum, and the location of that optimum changes with H . The horizon is a function of drift regularity, not a fixed window rule.

The empirical optima separate by roughness class, matching the path-dependent horizon logic of corollary 2.5.

3.3 Real-Stream Retention Diagnostic

The same qualitative signature appears on a real non-stationary stream. Using ELEC2, consider the one-dimensional marginal law of `nswprice`. At each time, approximate the present target by a short anchor block of 48 observations and compare it, in W_1 , with the empirical law induced by the preceding memory window. Averaging this diagnostic across the stream yields a real-data analogue of the finite-memory U-curve.




Figure 5: ELEC2 real-stream diagnostic on the `nswprice` marginal. The mean anchor-block W_1 error exhibits a finite interior optimum and a near-optimal useful-memory band. The figure is diagnostic: it asks whether the temporal-validity signature appears on a real stream.

On ELEC2, the diagnostic optimum occurs at a finite memory scale, with a short near-optimal band around it. The real stream therefore exhibits the same basic retention pattern as the synthetic horizon experiments: very short windows are too noisy, while very long windows become stale relative to the present block.

3.4 Moderate-Band Sinkhorn Evidence

The fixed- ε Sinkhorn experiments address the bandwise question: when a practical geometry replaces raw W_2 , does the finite-sample term remain stable enough for the horizon law to stay informative in an embedded operational regime? The pointwise and bandwise one-dimensional fixed-support statements appear in propositions 2.17 and 2.18; the calibrated embedded fixed-span model extends the same question to a broader low-dimensional setting.

The fixed- ε Sinkhorn geometry on embedded low-intrinsic-dimensional support shows that the corresponding finite-sample term remains stable across the tested regularization values and that the triangular slope stays close to the i.i.d. benchmark, with no separate exponent emerging. Combined with the structural ingredients isolated in proposition 2.16, this identifies a stable embedded operating regime. In the calibrated fixed-span model with $\alpha = 2$ and span 0.25, the maximal observed stable bands are $\varepsilon_{\max} = 0.5$ for $(d, k) = (8, 1)$, $(8, 2)$, and $(12, 2)$, and $\varepsilon_{\max} = 0.2$ for $(12, 1)$.

The same synthetic family also separates two mechanisms. Raw W_2 comparison between the triangular and i.i.d. empirical measures slows once intrinsic dimension increases, so the embedded fixed- ε statement does not follow from the available W_2 transfer route. Direct Jacobian

probes of the centered Sinkhorn scaling map instead isolate a moderate-band regime: on $\varepsilon \in [0.2, 0.8]$, the cross-coupling block stays contractive, while the self-coupling blocks are the main source of instability but appear numerically to plateau with sample size.

The calibrated fixed-span model shows that the tested grid has a stable core and a small mixed boundary region. Calibration details and supporting tables are kept in appendix B. The strongest support appears on moderate bands bounded away from zero.

The certified-band report on the same calibrated grid is positive on the same stable pairs: it returns maximal certified bands of at least 0.5 for $(d, k) = (8, 1)$, $(8, 2)$, and $(12, 2)$, and 0.2 for $(12, 1)$. This does not prove the embedded theorem, but it identifies a concrete moderate-band regime in which the conditional stability assumptions are simultaneously satisfied by the current certificate.

Ambient / intrinsic	α	Max certified ε
$(8, 1)$	2.0	0.5
$(8, 2)$	2.0	0.5
$(12, 1)$	2.0	0.2
$(12, 2)$	2.0	0.5

Table 1: Calibrated moderate-band Sinkhorn certificates on the embedded fixed-span grid. The table records the largest regularization value for which the current certificate returns a certified band on the tested synthetic frontier.

To probe whether the frontier is controlled more by local intrinsic geometry than by the ambient coordinate count, we estimate the effective intrinsic dimension with TwoNN on the same synthetic frontier family. On the nine-pair, six- ε grid studied here, the resulting $\hat{k}$ predicts the stable-band cutoff better than ambient dimension and improves the binary cut at $\varepsilon = 0.2$. This is consistent with geometry-aware calibration: the relevant control variable is the estimated local geometry, not the ambient coordinate count. The estimator therefore enters here as a regime-selection diagnostic for the band analysis.

3.5 Validity Precedes Detection

Temporal validity and changepoint evidence are different clocks. To make that separation explicit, consider a stream whose local drift rate ramps upward over time. The horizon n_t^* contracts as the ramp steepens, while a long operating horizon remains fixed. Define

$$\Delta_{\text{val-det}} := t_{\text{detect}} - t_{\text{valid}},$$

where t_{valid} is the first time the operating horizon persistently exceeds the temporal-validity horizon, and t_{detect} is the first detector alarm that follows.

Across the reported runs, the measured lag remains positive. Retained evidence can already be invalid for the present while detector statistics remain below alarm threshold.

The positive lag is not a knife-edge feature of one tuning choice. Across the tested ADWIN thresholds, mean detection time remains strictly later than mean validity-expiry time, and the same qualitative lag persists under




Figure 6: Temporal validity can fail before detector alarm. As drift intensifies, the temporal-validity horizon contracts until the operating horizon becomes stale at t_{valid} ; the detector responds only later at t_{detect} , leaving a strictly positive validity-detection lag. The lower panel shows excess tracking error accumulating before the alarm appears.

a different detector family as well. The empirical claim is robustness of the sign of $\Delta_{\text{val-det}}$ on the tested detector/tuning family, not a detector-uniform delay law.

3.6 Empirical Signatures

The experiments support a structural U-curve with a near-optimal useful-memory region, roughness-dependent horizon scaling, the same retention pattern on the ELEC2 stream, structured support for the fixed- ϵ Sinkhorn frontier, and a measurable lag between validity loss and detector alarm.

4 Implications for Online Adaptation

The horizon law identifies the quantities that any adaptive memory rule must track. A policy that updates memory online needs a finite-sample regime, represented by the exponent a or by an operational exponent a_ϵ , together with a local roughness description of the recent path through the pair (H, ζ) . Once these objects are available, the structural scale $n_\star$ and the useful-memory region $\mathcal{U}_\delta$ follow directly from the horizon law.

In operational geometries, local support structure enters before roughness estimation. The embedded fixed- ϵ theory already shows that intrinsic complexity controls which finite-sample regime is available through the threshold $\alpha > k/2$. A local intrinsic-dimension statistic such as TwoNN therefore serves as a regime-selection statistic: it does not estimate the horizon by itself, but it selects the finite-sample law to which the horizon formula should be attached.

The roughness side is a separate statistical task. In deterministic Hölder paths, H and ζ govern how quickly staleness can accumulate with lag. A direct log-log fit to one short empirical sequence of transport discrepancies is too unstable to play that role reliably: single-path local motion, finite-sample noise, and short-lag discreteness all distort the envelope. The practical online object is therefore a stabilized roughness estimate built from aggregated

recent lags, together with an activity-style scale statistic for band selection.




Figure 7: Embedded geometry and online adaptation. Left: TwoNN recovers the intrinsic regime and aggregated lag statistics estimate H far more accurately than endpoint slopes. Middle: the structural controller improves decisively on the naive plug-in rule and beats the activity and static baselines, although the broader validation-based adaptive hybrid remains stronger. Right: the structural controller tracks the contraction and expansion of the oracle window on the default phase profile.

That comparison also distinguishes two controller architectures. A direct plug-in rule based on a local slope and intercept is too unstable to be a useful online prescription. A more viable rule is structural: estimate the local regime with intrinsic geometry, estimate H from aggregated lag discrepancies, use an activity statistic for the current scale of motion, and then map those quantities to an admissible useful-memory band. In the synthetic study, that structural controller improves substantially on the naive plug-in rule and on a pure activity baseline, even though a wider validation-based hybrid remains more accurate.

On any compact parameter range with H bounded away from zero and ζ bounded above and below, the map

$$(H, \zeta) \mapsto n_\star(H, \zeta) = \left(\frac{aC_K}{HC_S\zeta} \right)^{1/(a+H)}$$

is Lipschitz in $(H, \log \zeta)$ after taking logarithms. Small estimation error in the roughness class therefore produces only a controlled multiplicative error in the plug-in horizon. The remaining difficulty is the scale variable: a local activity statistic is good enough to produce a useful controller, but not yet stable enough to match the stronger validation-based hybrid.

The online problem is therefore structurally well posed. The horizon law supplies the target for adaptation; the remaining task is to estimate its ingredients with guarantees strong enough to keep a data-driven memory rule inside the useful-memory region.

5 Related Work

Nonstationary optimization and dynamic regret. Dynamic-regret and variation-budget analyses show that under drift, forgetting scales and window lengths must adapt to how quickly the target moves [Besbes et al., 2015, Cheung et al., 2019, Yuan and Lamperski, 2020, Mazzetto and Upfal, 2023, Wang et al., 2024]. That literature isolates the control importance of retention scale. The horizon law gives the corresponding statistical law for distribution tracking: finite memory has a temporal-validity horizon determined jointly by a finite-sample exponent and a drift roughness exponent.

This distinction is especially sharp relative to adaptive update policies such as [Mazzetto and Upfal \[2023\]](#). Those works optimize a decision or intervention rule under evolving data. The present object is earlier in the pipeline: it asks how long retained evidence continues to represent the present distribution before age becomes the dominant source of tracking error.

Adaptive windowing, drift detection, and update policies. Adaptive-window and concept-drift methods study when accumulated evidence is sufficient to react [[Bifet and Gavalda, 2007](#), [Gama et al., 2004, 2014](#)]. [Hinder et al. \[2023\]](#) emphasize that the window length itself is part of the problem. Resource-aware update scheduling under drift addresses a related decision problem in which update cost must be balanced against accumulated loss [[Piaseczny et al., 2025](#)], and finite-time analyses of sliding-window or discounted adaptation under switching environments make the same design variable explicit from the regret side [[Garivier and Moulines, 2011](#)]. Streaming OT and MMD detectors contribute the alarm-time perspective, emphasizing delay, power, and memory usage [[Shvetsov et al., 2020](#), [Kalinke et al., 2025](#)]. These works provide the reaction timescale. The horizon law adds the validity timescale: retained evidence can cease to represent the present before change statistics cross an alarm threshold.

Adaptive filtering and locally stationary smoothing. Classical adaptive filtering and state-space tracking tune forgetting factors, gain schedules, or process-noise levels to balance smoothing against responsiveness [[Kalman, 1960](#), [Jazwinski, 1970](#)]. Locally stationary smoothing and bandwidth selection solve the analogous bias-variance problem for evolving regression or spectral targets [[Dahlhaus, 1997](#), [Vogt, 2012](#)]. Recent Wasserstein-based results for locally stationary functional time series move that same local windowing logic directly into a transport metric [[Tinio et al., 2025](#)]. These literatures supply mechanisms for adaptation. The present work contributes the distributional law that such mechanisms must respect when the retained object is an empirical distribution rather than a latent state or a local regression surface.

Distribution tracking under transport geometry. Wasserstein distance and Sinkhorn-type estimators provide the finite-sample geometries that enter the horizon law. For empirical Wasserstein distances, [Fournier and Guillin \[2015\]](#), [Boissard and Le Gouic \[2014\]](#), and [Weed and Bach \[2019\]](#) show that the sampling exponent depends on geometric complexity even before drift enters. The horizon law takes that finite-sample exponent as one input and composes it with a staleness exponent induced by the drift path. In this sense, the transport literature provides the sampling side of the law, while the horizon law provides its temporal counterpart.

Entropic and smoothed transport as operational geometries. Regularized transport changes both statistical stability and computational behavior [[Genevay et al., 2019](#), [Mena and Niles-Weed, 2019](#), [del Barrio et al., 2023](#),

[Rigollet and Stromme, 2025](#), [Stromme, 2023](#), [Groppe and Hundrieser, 2024](#), [Eckstein and Nutz, 2022](#), [Ghosal et al., 2022](#)]. Fixed- ϵ Sinkhorn theory provides i.i.d. asymptotics, differentiability, null limit laws, and stability constants [[Goldfeld et al., 2024](#)]. Those results are the finite-sample anchor for the operational horizon law developed here. Recent linear-Gaussian entropic-OT analyses make the recursive structure of regularized transport explicit [[Akyildiz et al., 2024](#)], suggesting a constructive intermediate regime between fixed- ϵ i.i.d. asymptotics and fully drifted triangular-array tracking.

Smooth-functional and delta-method arguments provide the multivariate transfer mechanism behind this regime. Hadamard differentiability is the natural regularity class for lifting triangular-array fluctuations through a transport functional, and the fixed- ϵ Sinkhorn map sits in that category after gauge fixing and uniform stability control [[Goldfeld et al., 2024](#)]. That route is the right multivariate complement to the one-dimensional quantile-process proof model.

Intrinsic geometry and operational calibration. Geometry-aware calibration of regularization scales points toward intrinsic dimension as the relevant complexity variable in regularized transport [[Genans and Wintemberger, 2026](#)]. Local intrinsic-dimension diagnostics such as TwoNN provide a direct statistic for effective support geometry [[Facco et al., 2017](#)]. That perspective enters through the operational Sinkhorn frontier: intrinsic geometry helps identify which finite-sample regime is available, and therefore which horizon law is relevant.

Low-dimensional surrogate geometries. Sliced, max-sliced, subspace-robust, and projection-robust Wasserstein constructions replace ambient transport by one-dimensional or low-dimensional comparisons [[Kolouri et al., 2018](#), [Deshpande et al., 2019](#), [Paty and Cuturi, 2019](#), [Lin et al., 2020](#), [Nadjahi et al., 2020](#), [Nietert et al., 2022](#)]. These examples show concretely that geometry choice can change the finite-sample exponent. They therefore fit naturally into the horizon framework as alternative sampling geometries whose drift-induced useful-memory scales need not match raw Wasserstein behavior.

Wasserstein nonstationarity and Age of Information. Wasserstein nonstationarity measures already appear in online optimization and distributionally robust optimization under drift [[Jiang et al., 2020](#), [Keehan et al., 2025](#), [Amini and Vinas, 2026](#)]. The Age-of-Information literature treats age as a resource that must be balanced against communication or update costs [[Kaul et al., 2012](#), [Yates et al., 2021](#)]. Together these viewpoints make age and nonstationarity into operational variables. The horizon law gives them a distributional statistical interpretation: age is costly because it accumulates geometric mismatch with the present, and the roughness exponent determines how fast that cost grows.

[Keehan et al. \[2025\]](#) are particularly close in spirit: they derive concentration bounds for weighted empirical laws under sequential Wasserstein drift and optimize weights

for distributionally robust decisions. That question begins once a decision rule is fixed. The temporal-validity horizon addresses the preceding statistical question of how long evidence should be retained before staleness dominates the finite-sample gain. In this sense, their weighting problem and the present horizon problem are complementary rather than competing.

Weighted empirical-process and weighted-Wasserstein precursors are also relevant for the non-uniform-memory line. In particular, Mason [2016] studies weighted approximations to empirical Wasserstein fluctuations, providing a technical ancestor for the weighted-memory summaries $n_{\text{eff}}(w)$ and $M_H(w)$ used here.

Across these neighboring literatures, the shared theme is that nonstationarity makes retention scale matter. The present work identifies the statistical object that organizes that dependence for distribution tracking: a temporal-validity horizon, together with a useful-memory region and a distinct detectability timescale.

6 Limitations and Open Problems

The central structural object is the temporal-validity horizon for finite-memory tracking under drift, and the closed operational corollary is the induced useful-memory region. The theorem-level results are the general horizon law, the useful-memory region proposition, the uniform-window staleness bound, the one-dimensional root- n proof model, the structural Gaussian lower bound at the exponent level, and the Gaussian location minimax benchmark. The remaining claims are narrower.

Finite-sample scope. The theory is modular in the finite-sample exponent a , but the fully proved finite-sample instantiation established here is the one-dimensional root- n regime under bounded support and fixed span. The same balancing logic extends to other geometries only once an appropriate finite-sample theorem is supplied.

Operational scope. For fixed- ε Sinkhorn, the structural ingredients of the embedded fixed-span model are proved, and the compact-support one-dimensional pointwise and bandwise inheritance theorems are explicit. The remaining operational step is the embedded low-dimensional theorem that connects the triangular-array finite-sample term to the i.i.d. benchmark uniformly across a nontrivial regularization band. The numerical evidence is strongest on moderate compact bands bounded away from zero, where direct Jacobian probes remain stable.

Lower-bound scope. The structural lower bound is subclass based. The Gaussian construction is enough to show that, in the root- n regime, the horizon is structural rather than an artifact of a particular estimator. Class-tightness and constant-sharpness for broader deterministic Hölder path classes remain outside the present closure.

Sharp-constant and lower-theory scope. The exact Gaussian lower-bound constants are explicit within the

proof scheme, but they are not class-tight across broader path classes. Extending the lower theory beyond the root- n regime, to other finite-sample exponents $a \neq 1/2$, and to other parametric deformations such as Gaussian scale remains open.

Family-level scope. Regular dominated families are the natural extension class. The horizon logic is expected to carry over when local testing geometry and metric geometry remain comparable, but no theorem-level family law is proved here. Support-changing and singular families remain obstruction classes.

Weights and memory shape. Uniform windows are the main case. The open question is whether broader non-uniform weights can enlarge the useful-memory region or improve constants without changing the exponent across richer path classes. The minimal tapered-weight diagnostic is consistent with a modest constant-level gain on the default synthetic stream, but it is still only a single-stream signal and does not yet support a general non-uniform-memory claim. The theorem-oriented weighted summaries in the main text isolate the relevant quantities as $n_{\text{eff}}(w)$ and $M_H(w)$; the remaining question is whether those summaries support a robust constant improvement beyond the uniform window on more than one drift profile.

Online roughness estimation. The law is stated in terms of the roughness parameters H and ζ , but no fully online procedure is proposed here for estimating them or for keeping the memory inside the useful region under misspecification.

Empirical scope. The experiments are structurally motivated. Synthetic designs make the horizon law, useful-memory region, and validity-detection lag visible under controlled drift, while the ELEC2 diagnostic shows that the same retention signature appears on a real stream. Their role is to expose the object, not to demonstrate competitive performance on a benchmark suite.

Detectability versus validity. Validity loss before detector alarm is an empirical consequence of the theory, not the primary theorem object. Temporal validity can fail before detector alarms appear, but it does not prove a universal detection-delay theorem across detector classes.

Open problems. Five directions organize the remaining frontier: operational theorem closure, sharp lower theory, memory design, online roughness estimation, and online control together with validity-detection theory.

Problem 1: operational theorem closure for embedded Sinkhorn and smooth multivariate transfer. Embedded Sinkhorn has a conditional moderate-band theorem, a certified band report, and a self-coupling stability certificate. The remaining problem is to prove the self-coupling stability bound in the embedded fixed-span model, and to lift the smooth transfer proposition

to a transport-functional theorem with a verified Jacobian remainder bound.

Problem 2: sharp lower theory beyond the current benchmark family. The exact Gaussian constants are available within the proof scheme, but the class-tight lower-bound problem remains open for broader deterministic path classes and for alternative parametric deformations. The open problem is to extend the lower theory beyond the root- n regime, beyond the location benchmark, and beyond the current class-level sharp constants.

Problem 3: useful-memory design beyond uniform windows. The analysis is built around uniform windows, but the weighted-memory theorem shows that uniform weights maximize effective sample size and linear ramps recover exact lag. The open problem is to determine when non-uniform weights can improve constants or enlarge the useful-memory region on richer path classes, and whether the summaries $n_{\text{eff}}(w)$ and $M_H(w)$ support a robust optimization law of the form

$$\min_{w: \|w\|_1=1} [C_K n_{\text{eff}}(w)^{-a} + C_S \zeta M_H(w)].$$

This is the memory-design problem that links static window design to fully adaptive control.

Problem 4: online estimation of roughness and scale. A structural controller, an intrinsic-dimension statistic, and an aggregated-lag roughness estimate are available. The open problem is to derive convergence rates for online estimators of H and ζ when transport discrepancies are themselves observed with finite-sample error. A concrete target is the log-log estimator

$$\hat{H}_L \in \arg \min_H \sum_{j=1}^L (\log \hat{D}_j - \log \hat{\zeta} - H \log j)^2,$$

where $\hat{D}_j$ estimates the lag- j transport discrepancy. The open problem is to derive the rate of $\hat{H}_L$ as a function of L , the finite-sample exponent a , and the lag-noise level, and then to propagate that rate through the horizon map to obtain a high-probability statement of the form $\mathbb{P}(\hat{n}_t \in \mathcal{U}_\delta) \geq 1 - \eta_{L,n}$.

Problem 5: online control and validity-detection theory. The positive validity-detection lag suggests that a memory controller can, in principle, react before standard alarms fire. The open problem is to prove an excess-risk guarantee of the form

$$\mathbb{E} \Phi(\hat{n}_t) \leq (1 + \delta_t) \Phi(n_t^*)$$

for an adaptive window rule, and to derive detector-specific lower or upper laws that quantify when validity must fail before statistical detection. Under stylized ramp drift, a natural target is a detector-specific law for the expected lag $\Delta_{\text{val-det}}$, for example a scaling relation of the form $\mathbb{E} \Delta_{\text{val-det}} \asymp (\zeta/\sigma)^{-2/(2H+1)}$ up to detector-dependent constants.

7 Conclusion

Under drift, memory is never purely beneficial. It reduces finite-sample noise, but it also ages. The resulting trade-off makes finite memory a temporal-validity object with a finite horizon.

That idea is formalized through the horizon law

$$\mathbb{E} d(\hat{P}_t^{(n)}, P_t) \leq C_K n^{-a} + C_S \zeta n^H,$$

whose optimal scale is

$$n^*(a, H) \asymp (C_K/\zeta)^{1/(a+H)}.$$

Finite-memory tracking under drift is governed by this balance between finite-sample accuracy and drift-induced staleness, and the same balance induces a useful-memory region around the optimal scale.

A tractable one-dimensional proof model supplies a rigorous root- n finite-sample regime, and a structural Gaussian lower-bound construction shows that the corresponding horizon is not an artifact of one estimator. The same object yields a Gaussian location minimax benchmark on deterministic Hölder paths. In fixed- ε Sinkhorn geometry, the compact-support one-dimensional pointwise and bandwise inheritance theorems show that the same horizon logic remains visible in an operational transport geometry.

The same decomposition extends to non-uniform lag weights through effective sample size and weighted lag moment summaries. The theorem-level question is whether that weighted horizon law improves constants on richer drift classes or only reshapes the same balance.

The finite-memory U-curve, roughness-dependent horizon scaling, a useful-memory region around the optimal scale, and a strictly positive lag between validity loss and detector alarm all fit the same horizon law. The resulting message is both theoretical and operational: retained evidence has a path-dependent horizon, useful memory occupies a finite near-optimal region, and statistical detectability need not coincide with temporal validity.

Temporal validity of retained evidence is therefore a first-class statistical object under drift. The horizon law identifies its governing scale, the useful-memory region records its operational range, and the detectability lag shows that validity can fail before alarms appear. Those three statements define the contribution of the paper.

The five open problems in section 6 mark the precise frontier of this program. The first two close the current theorem line, the third isolates the design limits of memory itself, and the last two turn the horizon law into a fully online object with estimation and control guarantees. Together they define a coherent next stage for temporal validity under drift without changing the identity of the present work.

Acknowledgements

With gratitude to everyone who contributed ideas, feedback, tools, encouragement, and care along the way. This work stands on the generosity of colleagues, collaborators, open-source maintainers, and the broader research communities that made it possible.

A Proof Details

This appendix gives the detailed derivations for structural results whose full calculations would interrupt the main line.

A.1 Gaussian lower-bound constants

Let

$$S_h := \sum_{r=1}^h r^{2H} \quad \text{and} \quad G_h := \sum_{r=1}^h (h^H - (h-r)^H)^2.$$

For the Gaussian two-point model with equal priors and common covariance $\sigma^2 I_h$, the Bayes classification error between means $\pm m$ is $\Phi(-\|m\|/\sigma)$. Therefore the ramp and endpoint-minimal payoffs are exactly

$$L_h^{\text{ramp}}(\sigma, \zeta) = \zeta h^H \Phi\left(-\frac{\zeta}{\sigma} \sqrt{S_h}\right)$$

and

$$L_h^{\text{min}}(\sigma, \zeta) = \zeta h^H \Phi\left(-\frac{\zeta}{\sigma} \sqrt{G_h}\right).$$

Proof of proposition 2.12. Fix $H \in (0, 1]$ and write $p_H = 2H/(2H+1)$.

For the ramp profile,

$$S_h \sim \frac{h^{2H+1}}{2H+1}.$$

Set

$$h = A \left(\frac{\sigma}{\zeta}\right)^{2/(2H+1)}$$

with $A > 0$ fixed. Then

$$\frac{\zeta}{\sigma} \sqrt{S_h} \rightarrow \frac{A^{H+1/2}}{\sqrt{2H+1}}.$$

At the same time,

$$\zeta h^H = A^H \sigma^{2H/(2H+1)} \zeta^{1/(2H+1)}.$$

Hence

$$\frac{L_h^{\text{ramp}}(\sigma, \zeta)}{\sigma^{2H/(2H+1)} \zeta^{1/(2H+1)}} \rightarrow A^H \Phi\left(-\frac{A^{H+1/2}}{\sqrt{2H+1}}\right).$$

Define

$$x = \frac{A^{H+1/2}}{\sqrt{2H+1}}.$$

Then

$$A^H = (2H+1)^{H/(2H+1)} x^{p_H},$$

so the normalized limiting payoff becomes

$$(2H+1)^{H/(2H+1)} x^{p_H} \Phi(-x).$$

Its logarithmic derivative is

$$\frac{p_H}{x} - \frac{\varphi(x)}{\Phi(-x)}.$$

Since the inverse Mills ratio $\varphi(x)/\Phi(-x)$ is strictly increasing on $(0, \infty)$, this objective has a unique maximizer $x_H > 0$, characterized by

$$p_H \Phi(-x_H) = x_H \varphi(x_H).$$

This proves

$$\sup_{h \geq 1} L_h^{\text{ramp}}(\sigma, \zeta) \sim C_H^{\text{ramp}} \sigma^{2H/(2H+1)} \zeta^{1/(2H+1)}$$

with

$$C_H^{\text{ramp}} = (2H+1)^{H/(2H+1)} x_H^{2H/(2H+1)} \Phi(-x_H).$$

For the endpoint-minimal profile, the Hölder constraint and endpoint condition imply

$$g_r \geq g_h - (h-r)^H = h^H - (h-r)^H$$

for every feasible endpoint-saturating profile g . Therefore the profile

$$g_r^{\text{min}} = h^H - (h-r)^H$$

minimizes testing energy pointwise and hence minimizes $\sum_{r=1}^h g_r^2$. Its rescaled energy satisfies the Riemann-sum limit. Thus

$$h^{-(2H+1)} G_h \rightarrow I_H,$$

where

$$I_H = \int_0^1 (1 - (1-u)^H)^2 du = \frac{2H^2}{(H+1)(2H+1)}.$$

With the same scaling

$$h = A \left(\frac{\sigma}{\zeta}\right)^{2/(2H+1)},$$

one has

$$\frac{\zeta}{\sigma} \sqrt{G_h} \rightarrow A^{H+1/2} \sqrt{I_H}.$$

Setting $x = A^{H+1/2} \sqrt{I_H}$ gives

$$A^H = I_H^{-H/(2H+1)} x^{p_H},$$

so the normalized limiting payoff becomes

$$I_H^{-H/(2H+1)} x^{p_H} \Phi(-x).$$

The same maximizer x_H therefore yields

$$\sup_{h \geq 1} L_h^{\text{min}}(\sigma, \zeta) \sim C_H^{\text{min}} \sigma^{2H/(2H+1)} \zeta^{1/(2H+1)}$$

with

$$C_H^{\text{min}} = I_H^{-H/(2H+1)} x_H^{2H/(2H+1)} \Phi(-x_H).$$

Finally,

$$\frac{C_H^{\text{min}}}{C_H^{\text{ramp}}} = \left(\frac{I_H^{-1}}{2H+1}\right)^{H/(2H+1)} = \left(\frac{H+1}{2H^2}\right)^{H/(2H+1)}.$$

This factor is strictly larger than 1 for $H < 1$ and equal to 1 at $H = 1$. $\square$

A.2 Two-point lower bound used in proposition 2.11

The structural lower bound above is the exponent-level consequence of the exact Gaussian two-point calculation. For the ramp pair

$$\mu_{t-j}^{\pm, h} = \pm \zeta (h-j)_+^H,$$

the present-time separation is $2\zeta h^H$, so any estimator incurs risk at least of order $L_h^{\text{ramp}}(\sigma, \zeta)$. Minimizing over h and using the asymptotic just proved gives the lower scale

$$\sigma^{2H/(2H+1)} \zeta^{1/(2H+1)}$$

and therefore the horizon exponent $(2H+1)^{-1}$ for the root- n regime.

B Finite-Sample and Operational Diagnostics

This appendix gives finite-sample diagnostics for the theorem line and the remaining operational frontier. These diagnostics are not substitutes for theorem statements.

B.1 Finite-Sample Geometry Diagnostic

For the finite-sample term, we estimate empirical W_2 between two i.i.d. samples from the same distribution using exact assignment in low and moderate dimension, then fit a log-log slope in n . We also probe the fixed-span one-dimensional triangular window directly by comparing a triangular-array sample, with one draw from each drifted component, against an i.i.d. sample from the same window mixture. In that diagnostic the within-window span ζn^H is held fixed. Finally, we probe the Sinkhorn diagnostic at a few fixed values of ε .

Table 2: Empirical scaling diagnostic for the finite-sample term. The bounded-support i.i.d. W_2 experiment is close to the expected $n^{-1/2}$ rate in low dimension, the corresponding triangular-array diagnostic compares one sample from each drifted component against an i.i.d. sample from the window mixture, and the Sinkhorn diagnostic is reported at fixed ε . These rows are numerical support for the operational finite-sample discussion rather than theorem statements.

Setting	Estimated slope	Comment
$W_2, d = 1$	-0.48	close to $-1/2$
$W_2, d = 2$	-0.41	slower decay
$W_2, d = 4$	-0.25	slower decay
$W_2, d = 5$	-0.24	slower decay
W_2 , i.i.d. mixture, $d = 1, H = 0.5$	-0.50	fixed span $\zeta n^H = 0.25$
W_2 , triangular, $d = 1, H = 0.5$	-0.45	fixed span $\zeta n^H = 0.25$
W_2 , i.i.d. mixture, $d = 1, H = 1.0$	-0.47	fixed span $\zeta n^H = 0.25$
W_2 , triangular, $d = 1, H = 1.0$	-0.41	fixed span $\zeta n^H = 0.25$
Sinkhorn $\varepsilon = 0.50$	-0.21	stability constant changes
Sinkhorn $\varepsilon = 0.10$	-0.53	stability constant changes
Sinkhorn $\varepsilon = 0.05$	-0.67	stability constant changes

In low dimension, the raw i.i.d. Wasserstein term is compatible with a root- n rate, whereas higher-dimensional raw Wasserstein slows down. In the one-dimensional fixed-span triangular-array diagnostic, the recovered slopes stay near the corresponding i.i.d. mixture baseline, which is consistent with the proof-model and design-effect reading used in the 1D proof-model proposition.

B.2 Bahadur Remainder Diagnostics

Two deformation families make the proof-model remainder assumption concrete. The first is a smooth fixed-support monotone warp class $g_\delta(x) = x + \delta x(1 - x)$ with $|\delta| \leq \Delta < 1$. On the tested spans 0.2 and 0.4, the integrated remainder decays at rates between 1.40 and 1.54, and the scaled quantity $n \int_0^1 R_n(u)^2 du$ still decays at rates between 0.40 and 0.54. The second is the fixed-span triangular-array uniform-kernel model used in the proof line. There the integrated residual rate is about 1.43 for the triangular sample and 1.49 for the i.i.d. mixture benchmark, while the residual-to-MSE ratio drops from about 0.38 at $n = 25$ to about 0.11 at $n = 400$.

These diagnostics locate the role of drift span explicitly. Within the tested fixed-span range, the remainder stays lower-order and the empirical rate remains well above the n^{-1} threshold needed for item (iv). Rougher kernels and growing-span experiments remain less regular, which is why the manuscript keeps the general triangular-array verification problem open.

B.3 Extended and Operational Regimes

The extended rows compare triangular and i.i.d. mixture rates on embedded low-intrinsic-dimensional support across span. The operational rows compare fixed- ε Sinkhorn rates across regularization values on embedded low-intrinsic-dimensional support in ambient dimension eight.

Table 3: Representative rows for the extended and operational finite-sample regimes. The extended rows compare triangular and i.i.d. mixture rates across span, while the operational rows compare fixed- ε Sinkhorn across ε on embedded low-intrinsic-dimensional support.

Regime	Setting	Parameter	Rate exponent a
extended	triangular ambient $d=8$, intrinsic $k=1$, span=0.10	0.10	0.4652
extended	iid mixture ambient $d=8$, intrinsic $k=1$, span=0.10	0.10	0.5722
extended	triangular ambient $d=8$, intrinsic $k=1$, span=0.25	0.25	0.4737
extended	iid mixture ambient $d=8$, intrinsic $k=1$, span=0.25	0.25	0.5588
extended	triangular ambient $d=8$, intrinsic $k=1$, span=0.50	0.50	0.4792
extended	iid mixture ambient $d=8$, intrinsic $k=1$, span=0.50	0.50	0.5425
operational	Sinkhorn iid mixture ambient $d=8$, intrinsic $k=1$, $\varepsilon=0.50$	0.50	0.5599
operational	Sinkhorn triangular ambient $d=8$, intrinsic $k=1$, $\varepsilon=0.50$	0.50	0.5706
operational	Sinkhorn iid mixture ambient $d=8$, intrinsic $k=1$, $\varepsilon=0.20$	0.20	0.5466
operational	Sinkhorn triangular ambient $d=8$, intrinsic $k=1$, $\varepsilon=0.20$	0.20	0.6150
operational	Sinkhorn iid mixture ambient $d=8$, intrinsic $k=1$, $\varepsilon=0.10$	0.10	0.4646
operational	Sinkhorn triangular ambient $d=8$, intrinsic $k=1$, $\varepsilon=0.10$	0.10	0.5060
operational	Sinkhorn iid mixture ambient $d=8$, intrinsic $k=1$, $\varepsilon=0.05$	0.05	0.4339
operational	Sinkhorn triangular ambient $d=8$, intrinsic $k=1$, $\varepsilon=0.05$	0.05	0.5124

The extended-regime rows indicate that low intrinsic dimension preserves a rate near the root- n regime without a dramatic triangular-versus-benchmark gap. The operational-regime rows record empirical evidence relevant to the fixed- ε Sinkhorn setting in definition 2.15: together with exact support-complexity inheritance and the dual threshold $\alpha > k/2$, they align with the embedded bandwise statement. A full proof of triangular-array inheritance would close this loop.

Two complementary diagnostics sharpen the same boundary. First, direct probes of the centered Sinkhorn scaling Jacobian on the embedded fixed-span model show that bands approaching zero are the numerically unstable region, whereas on a moderate band such as $[0.2, 0.8]$ the cross-coupling block remains contractive and the self-coupling inverse norms appear to plateau with sample size. Second, on the supported synthetic frontier grid, intrinsic-dimension estimates predict the observed stable-band cutoff better than ambient dimension. These diagnostics identify a useful regime statistic and isolate the band where the remaining instability is concentrated. They do not upgrade the embedded fixed- ε statement from conjecture to theorem.

The current self-coupling certificate is positive on the calibrated (8, 2) and (12, 2) moderate-band pairs, where both self blocks keep positive spectral gap, worst spectral radius below 0.97, and largest-sample mean inverse norm below 3.0.

Pair	Coup.	Largest- n mean inv.	Proxy
(8, 2)	xx	2.714	yes
(8, 2)	yy	2.745	yes
(12, 2)	xx	2.759	yes
(12, 2)	yy	2.715	yes

Table 4: Self-coupling stability certificate on the calibrated moderate-band Sinkhorn grid. The certificate requires positive spectral gap, worst spectral radius below 0.97, and largest-sample mean inverse norm below 3.0.

B.4 Smooth Multivariate Transfer Diagnostics

The multivariate transfer diagnostic uses a drifting three-dimensional mean vector and a smooth log-sum-exp functional. The finite-sample error of the window estimator, the drift-induced bias of the window target, and the linearization residual are tracked separately across window lengths. The target behavior is the same as in the abstract multivariate transfer lemma: the finite-sample piece decreases with window length, the drift piece increases, and the linearization residual remains lower order than the finite-sample term on the tested windows.

Window	Tot.	Fin.	Drift	Resid.
8	0.2009	0.1988	0.0165	0.0239
16	0.1513	0.1542	0.0355	0.0162
32	0.1373	0.1363	0.0735	0.0066
64	0.1492	0.0807	0.1503	0.0051

Table 5: Smooth multivariate transfer diagnostic under triangular-array drift. The finite-sample term decreases with window length, the drift term increases, the total error has an interior minimum, and the linearization residual stays lower order across the tested windows.

References

O. Deniz Akyildiz, Pierre Del Moral, and Joaquín Miguez. Gaussian entropic optimal transport: Schrödinger bridges and the sinkhorn algorithm. *CoRR*, abs/2412.18432, 2024. doi: 10.48550/arXiv.2412.18432.

Arash A. Amini and Luciano Vinas. Wasserstein concentration of empirical measures for dependent data via the method of moments. *CoRR*, abs/2601.07228, 2026. doi: 10.48550/arXiv.2601.07228.

R. R. Bahadur. A note on quantiles in large samples. *The Annals of Mathematical Statistics*, 37(3):577–580, 1966.

Omar Besbes, Yonatan Gur, and Assaf Zeevi. Non-stationary stochastic optimization. *Operations Research*, 63(5):1227–1244, 2015. doi: 10.1287/OPRE.2015.1408.

Albert Bifet and Ricard Gavaldà. Learning from time-changing data with adaptive windowing. In *Proceedings of the 2007 SIAM International Conference on Data*

Mining (SDM), pages 443–448, 2007. doi: 10.1137/1.9781611972771.42.

Emmanuel Boissard and Thibaut Le Gouic. On the mean speed of convergence of empirical and occupation measures in wasserstein distance. *Annales de l’Institut Henri Poincaré, Probabilités et Statistiques*, 50(2):539–563, 2014. doi: 10.1214/12-AIHP517.

Wang Chi Cheung, David Simchi-Levi, and Ruihao Zhu. Learning to optimize under non-stationarity. In *Proceedings of the Twenty-Second International Conference on Artificial Intelligence and Statistics*, volume 89 of *Proceedings of Machine Learning Research*, pages 1079–1087. PMLR, 2019. URL <https://proceedings.mlr.press/v89/cheung19b.html>.

Rainer Dahlhaus. Fitting time series models to nonstationary processes. *The Annals of Statistics*, 25(1):1–37, 1997. doi: 10.1214/aos/1034276620.

Eustasio del Barrio, Alberto Gonzalez-Sanz, Jean-Michel Loubes, and Jonathan Niles-Weed. An improved central limit theorem and fast convergence rates for entropic transportation costs. *SIAM Journal on Mathematics of Data Science*, 5(3):639–669, 2023. doi: 10.1137/22M149260X.

Ishan Deshpande, Yuan-Ting Hu, Ruoyu Sun, Ayis Pyrros, Nasir Siddiqui, Sanmi Koyejo, Zhizhen Zhao, David Forsyth, and Alexander G. Schwing. Max-sliced wasserstein distance and its use for GANs. In *2019 IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 10648–10656, 2019. doi: 10.1109/CVPR.2019.01090.

Stephan Eckstein and Marcel Nutz. Quantitative stability of regularized optimal transport and convergence of sinkhorn’s algorithm. *SIAM Journal on Mathematical Analysis*, 54(6):5922–5948, 2022. doi: 10.1137/21M145505X.

Elena Facco, Maria d’Errico, Alex Rodriguez, and Alessandro Laio. Estimating the intrinsic dimension of datasets by a minimal neighborhood information. *Scientific Reports*, 7(1):12140, 2017. doi: 10.1038/s41598-017-11873-y. URL <https://doi.org/10.1038/s41598-017-11873-y>.

Nicolas Fournier and Arnaud Guillin. On the rate of convergence in Wasserstein distance of the empirical measure. *Probability Theory and Related Fields*, 162(3–4):707–738, 2015. doi: 10.1007/s00440-014-0583-7.

João Gama, Pedro Medas, Gladys Castillo, and Pedro Rodrigues. Learning with drift detection. In *Advances in Artificial Intelligence — SBIA 2004*, volume 3171 of *Lecture Notes in Computer Science*, pages 286–295. Springer, 2004. doi: 10.1007/978-3-540-28645-5_29.

João Gama, Indrė Žliobaitė, Albert Bifet, Mykola Pechenizkiy, and Abdelhamid Bouchachia. A survey on concept drift adaptation. *ACM Computing Surveys*, 46(4):44:1–44:37, 2014. doi: 10.1145/2523813.

- Aurélien Garivier and Éric Moulines. On upper-confidence bound policies for switching bandit problems. In *Algorithmic Learning Theory*, volume 6925 of *Lecture Notes in Computer Science*, pages 174–188. Springer, 2011. doi: 10.1007/978-3-642-24412-4_16.
- Ferdinand Genans and Olivier Wintenberger. Geometry-aware optimal transport: Fast intrinsic dimension and wasserstein distance estimation. *CoRR*, abs/2602.04335, 2026. doi: 10.48550/arXiv.2602.04335.
- Aude Genevay, Lénaïc Chizat, Francis Bach, Marco Cuturi, and Gabriel Peyré. Sample complexity of Sinkhorn divergences. In *Proceedings of the 22nd International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 89 of *Proceedings of Machine Learning Research*, pages 1574–1583, 2019.
- Promit Ghosal, Marcel Nutz, and Espen Bernton. Stability of entropic optimal transport and schrodinger bridges. *Journal of Functional Analysis*, 283(9):109622, 2022. doi: 10.1016/j.jfa.2022.109622.
- Ziv Goldfeld, Kengo Kato, Gabriel Rioux, and Ritwik Sadhu. Limit theorems for entropic optimal transport maps and the sinkhorn divergence. *Electronic Journal of Statistics*, 18(1):980–1041, 2024. doi: 10.1214/24-EJS2217.
- Michel Groppe and Shayan Hundrieser. Lower complexity adaptation for empirical entropic optimal transport. *Journal of Machine Learning Research*, 25(344):1–55, 2024. URL <https://jmlr.org/papers/v25/23-0671.html>.
- Fabian Hinder, Valerie Vaquet, and Barbara Hammer. The window dilemma: Why concept drift detection is ill-posed, 2023. Preprint.
- Andrew H. Jazwinski. *Stochastic Processes and Filtering Theory*, volume 64 of *Mathematics in Science and Engineering*. Academic Press, New York, 1970.
- Jiashuo Jiang, Xiaocheng Li, and Jiawei Zhang. Online stochastic optimization with wasserstein based non-stationarity, 2020. URL <https://arxiv.org/abs/2012.06961>.
- Florian Kalinke, Marco Heyden, Georg Gntuni, Edouard Fouché, and Klemens B"ohm. Maximum mean discrepancy on exponential windows for online change detection. *Transactions on Machine Learning Research*, 2025. doi: 10.48550/arXiv.2205.12706.
- Rudolf E. Kalman. A new approach to linear filtering and prediction problems. *Journal of Basic Engineering*, 82(1):35–45, 1960. doi: 10.1115/1.3662552.
- Sanjit Kaul, Roy Yates, and Marco Gruteser. Real-time status: How often should one update? In *Proceedings of IEEE INFOCOM*, pages 2731–2735, 2012. doi: 10.1109/INFCOM.2012.6195689.
- Dominic S. T. Keehan, Edward J. Anderson, and Wolfram Wiesemann. Don’t look back in anger: Wasserstein distributionally robust optimization with nonstationary data. *CoRR*, abs/2510.18566, 2025. doi: 10.48550/arXiv.2510.18566.
- Soheil Kolouri, Gustavo K. Rohde, and Heiko Hoffmann. Sliced wasserstein distance for learning gaussian mixture models. In *2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 3427–3436, 2018. doi: 10.1109/CVPR.2018.00361.
- Tianyi Lin, Chenyou Fan, Nhat Ho, Marco Cuturi, and Michael I. Jordan. Projection robust wasserstein distance and riemannian optimization. *Advances in Neural Information Processing Systems*, 33:9383–9397, 2020.
- David M. Mason. A weighted approximation approach to the study of the empirical wasserstein distance. In *High Dimensional Probability VII*, volume 71 of *Progress in Probability*, pages 137–154. Springer, 2016. doi: 10.1007/978-3-319-40519-3_6.
- Alessio Mazzetto and Eli Upfal. Nonparametric density estimation under distribution drift. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 24251–24270. PMLR, 2023. URL <https://proceedings.mlr.press/v202/mazzetto23a.html>.
- Gonzalo Mena and Jonathan Niles-Weed. Statistical bounds for entropic optimal transport: Sample complexity and the central limit theorem. *Advances in Neural Information Processing Systems*, 32:4543–4553, 2019.
- Kimia Nadjahi, Alain Durmus, Lenaic Chizat, Soheil Kolouri, Shahin Shahrampour, and Umut Simsekli. Statistical and topological properties of sliced probability divergences. *Advances in Neural Information Processing Systems*, 33:2085–2095, 2020.
- Sloan Nietert, Ziv Goldfeld, Ritwik Sadhu, and Kengo Kato. Statistical, robustness, and computational guarantees for sliced wasserstein distances. *Advances in Neural Information Processing Systems*, 35:28179–28193, 2022.
- Fran cois-Pierre Paty and Marco Cuturi. Subspace robust wasserstein distances. In *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pages 5072–5081, 2019.
- Adam Piaseczny, Md Kamran Chowdhury Shisher, Shiqiang Wang, and Christopher G. Brinton. Rcda: Adaptive model updates in the presence of concept drift under a constrained resource budget. *arXiv preprint arXiv:2505.24149*, 2025. doi: 10.48550/arXiv.2505.24149.
- Philippe Rigollet and Austin J. Stromme. On the sample complexity of entropic optimal transport. *The Annals of Statistics*, 53(1), 2025. doi: 10.1214/24-AOS2455.
- Nikolay Shvetsov, Nazar Buzun, and Dmitry V. Dylov. Unsupervised non-parametric change point detection in quasi-periodic signals. *arXiv preprint*

- arXiv:2002.02717*, 2020. doi: 10.1145/3400903.3400917.
- Austin J. Stromme. Minimum intrinsic dimension scaling for entropic optimal transport. *CoRR*, abs/2306.03398, 2023. doi: 10.48550/arXiv.2306.03398.
- Jan Nino G. Tinio, Mokhtar Z. Alaya, and Salim Bouzebda. Bounds in wasserstein distance for locally stationary functional time series. *CoRR*, abs/2504.06453, 2025. URL <https://arxiv.org/abs/2504.06453>.
- Aad W. van der Vaart. *Asymptotic Statistics*. Cambridge University Press, Cambridge, 1998.
- Michael Vogt. Nonparametric regression for locally stationary time series. Technical Report CWP22/12, cemmap, 2012. URL <http://www.cemmap.ac.uk/wps/cwp221212.pdf>.
- Zhiyong Wang, Jize Xie, Yi Chen, John C. S. Lui, and Dongruo Zhou. Variance-dependent regret bounds for non-stationary linear bandits, 2024. URL <https://arxiv.org/abs/2403.10732>.
- Jonathan Weed and Francis Bach. Sharp asymptotic and finite-sample rates of convergence of empirical measures in Wasserstein distance. *Bernoulli*, 25(4A):2620–2648, 2019. doi: 10.3150/18-BEJ1065.
- Roy D. Yates, Yin Sun, D. Richard Brown, Sanjit K. Kaul, Eytan Modiano, and Sennur Ulukus. Age of information: An introduction and survey. *IEEE Journal on Selected Areas in Communications*, 39(5):1183–1210, 2021. doi: 10.1109/JSAC.2021.3064980.
- Jianjun Yuan and Andrew G. Lamperski. Trading-off static and dynamic regret in online least-squares and beyond. In *The Thirty-Fourth AAAI Conference on Artificial Intelligence*, pages 6712–6719. AAAI Press, 2020. doi: 10.1609/AAAI.V34I04.6149.