

Useful Memory Has a Horizon

Structural Worst-Case Limits for Finite-Memory Tracking under Drift

Vinicius Parreira 

Independent Researcher

`figueiredo@pm.me`

May 2026

Abstract

In stationary settings, the law of large numbers makes memory an unambiguous statistical resource: more data improves estimation because the target distribution does not change over time. Under drift, this premise fails. Past evidence is generated by distributions that may no longer represent the present, so memory reduces variance while accumulating temporal misalignment. Once drift-induced misalignment dominates variance reduction, additional memory increases tracking error relative to the current distribution. Finite-memory tracking is therefore governed by a temporal-validity problem: how long can retained evidence continue to improve estimation of the present?

For drift path classes in which temporal misalignment grows as ζn^H , this variance-staleness trade-off induces a roughness-indexed useful-memory horizon $n_H^* \asymp (C_K/\zeta)^{2/(1+2H)}$ together with the corresponding minimax rates. The worst-case W_2 -Lipschitz envelope corresponds to the $H = 1$ member, recovering the cube-root horizon $n^* \asymp (C_K/\zeta)^{2/3}$. A Gaussian location lower bound shows that the same scaling defines a structural finite-memory floor.

This framework exposes a fundamental invalidity gap in non-stationary systems: retained evidence can lose temporal validity before a changepoint becomes statistically detectable. In finite samples, this structure appears through a finite optimum in horizon sweeps, roughness-dependent scaling across path classes, and measurable excess error when the operating horizon departs from the useful-memory scale.

Keywords: tracking under drift; useful memory; temporal validity; Wasserstein distance; Sinkhorn divergence; concept drift; finite-memory horizon.

MSC 2020: 93D05, 93D30, 68P99.

1 Introduction

Adaptive systems rarely track a fixed target. In streaming estimation, monitoring, and online decision-making, the data-generating distribution itself drifts over time. A learner with finite memory therefore faces a basic question: how much of the past should still count once the world that produced it has started to move away?

That question creates a structural trade-off. More past data reduces sampling noise, but it also increases the age of the evidence used to represent the present. Under drift, memory is simultaneously helpful and harmful: short windows are fresh but noisy, while long windows are stable but stale. The paper studies this tension as a horizon problem for temporal validity, and the main object is a roughness-indexed family of horizon laws.

Our main theorem package begins with a roughness-indexed trade-off between finite-sample variance and drift-induced staleness. In the worst-case W_2 -Lipschitz envelope, that family specializes to the cube-root law and a structural finite-memory floor. A Gaussian location subclass provides a matching lower bound establishing the same structural finite-memory floor.

The same object extends beyond the worst-case Lipschitz regime. For deterministic Holder-type path classes, temporal roughness changes how staleness accumulates, and the horizon law becomes a family indexed by roughness. This makes the cube-root law the $H = 1$ member of the family.

The practical conflict is between two clocks. One clock measures the age at which retained evidence stops being valid for the present. The other measures when accumulated data provide enough statistical evidence for a changepoint alarm. Those clocks need not move together. The central conceptual distinction is therefore simple: temporal validity is not changepoint evidence. A detector can remain statistically silent while retained memory is already stale for the present.

The contributions are fivefold:

- It formalizes useful memory under drift as a variance-staleness trade-off in Wasserstein geometry.
- It proves a worst-case finite-memory horizon law under local W_2 -Lipschitz drift.
- It shows a Gaussian-location lower bound establishing a structural tracking floor in the critical-window regime.
- It extends the same analysis to Holder-type path classes and their roughness-indexed horizon laws.
- It identifies detector-silent staleness and measures the invalidity gap between temporal validity and changepoint evidence.

The empirical evidence follows that logic in stages. It first recovers the predicted U -curve and the finite useful-memory optimum. It then shows the gap between temporal validity and changepoint evidence, the regime-dependent scaling across roughness classes, and the nonlinear cost of horizon misalignment. Real-stream evidence, when kept, serves only as weak confirmation of the same structural picture.

The theoretical results are stated in W_2 ; the Sinkhorn divergence is the computational measurement layer used in the synthetic experiments. The window size

result of [Hinder et al. \[2023\]](#) motivates the broader issue, but the contribution here is more specific: memory has a finite temporal validity horizon under worst-case drift, and that horizon law extends naturally to Holder-type path classes.

2 Temporal Validity Under Drift

Under drift, memory is both a statistical resource and a temporal liability. This section isolates the structural question behind useful memory: how long can retained evidence remain temporally valid before staleness overwhelms the benefit of variance reduction?




Figure 1: Two clocks of drift in normalized theorem units. The finite-sample term falls with horizon, the staleness term rises with age, and total error is minimized at the useful-memory horizon n^* . The shaded region marks detector-silent staleness: temporal invalidity can arrive before changepoint evidence.

The family of horizon laws comes first: temporal roughness controls how staleness accumulates, and the worst-case Lipschitz regime corresponds to the $H = 1$ member of that family.

2.1 Roughness-Indexed Horizon Family

Proposition 2.1 (Roughness-indexed horizon law). *Let a deterministic Holder-type path class have roughness index $H \in (0, 1]$ and staleness growth*

$$\text{Staleness}_H(n; \zeta) \asymp \zeta n^H.$$

If the finite-sample term remains of order $C_K n^{-1/2}$, then the error bound

$$\mathcal{E}_H(n) \leq C_K n^{-1/2} + c_H \zeta n^H$$

is minimized at

$$n_H^* \asymp (C_K/\zeta)^{2/(1+2H)}, \quad \mathcal{E}_{H,\min} \asymp C_K^{2H/(1+2H)} \zeta^{1/(1+2H)}.$$

In particular, $H = 1$ recovers the cube-root law.

Proof. Balance the decreasing statistical term against the increasing staleness term. The optimizer is determined by the point where the two terms have the same order. $\square$

Remark 2.2 (Family interpretation). This proposition is the roughness-indexed completion of the worst-case Lipschitz regime. As H decreases, staleness accumulates differently with lag and the useful-memory horizon shifts accordingly. The point is not that cube-root is special in itself; the point is that it is the $H = 1$ member of a broader geometry.

2.2 Worst-Case Lipschitz Regime

Theorem 2.3 (Lag-Inclusive Tracking Error Bound). *Let $X_{t-j} \sim P_{t-j}^*$ and let*

$$\hat{P}_t^{(n)} = \frac{1}{n} \sum_{j=0}^{n-1} \delta_{X_{t-j}}$$

be the uniform-window empirical estimator. Assume the target path is locally W_2 -Lipschitz with rate ζ , so that

$$W_2(P_{t-j}^*, P_t^*) \leq \zeta j \quad \text{for all } 0 \leq j < n,$$

and assume the finite-sample term obeys

$$\mathbb{E} \left[W_2 \left(\hat{P}_t^{(n)}, \bar{P}_t^{(n)} \right) \right] \leq C_K n^{-1/2},$$

where $\bar{P}_t^{(n)} := n^{-1} \sum_{j=0}^{n-1} P_{t-j}^$. Then the expected tracking error satisfies*

$$\mathcal{E}(n) \leq C_K n^{-1/2} + \frac{1}{2} \zeta n,$$

where the first term is the finite-sample statistical cost and the second term is the staleness cost induced by retaining old information.

Lemma 2.4 (Finite-sample Wasserstein term). *Let Q be a probability measure on $\mathbb{R}^{d_{\text{eff}}}$ with support contained in a ball of radius R and finite second moment. Let $X_1, \dots, X_n$ be i.i.d. samples from Q , and let $\hat{Q}_n := n^{-1} \sum_{i=1}^n \delta_{X_i}$. Then:*

- *If $d_{\text{eff}} < 4$, there exists a constant $C_{d_{\text{eff}}, R}$ such that*

$$\mathbb{E} W_2(\hat{Q}_n, Q) \leq C_{d_{\text{eff}}, R} n^{-1/2}.$$

- *If $d_{\text{eff}} = 4$, the same bound holds with an additional $(\log n)^{1/2}$ factor.*
- *If $d_{\text{eff}} > 4$, the raw Wasserstein rate is slower than $n^{-1/2}$, so the constant-only form above should not be used.*

In that regime, the statistical fluctuations do not follow the $n^{-1/2}$ carrier exponent, so the resulting horizon scaling stays outside the displayed cube-root regime. If the geometry is computed with debiased Sinkhorn divergence S_ε on a bounded domain of diameter D , then for fixed $\varepsilon > 0$ the empirical fluctuations remain $O(n^{-1/2})$ with a constant $L(D, \varepsilon)$ that is dimension-free at fixed ε and increases as $\varepsilon \downarrow 0$.

A short numerical sanity check for this regime appears in Appendix D; it is not used as a theorem, but it matches the intended low-dimensional interpretation of the finite-sample term.

2.3 Optimal Horizon in the Lipschitz Case

Remark 2.5 (Scope of the horizon law). The cube-root law follows when the finite-sample term behaves as $\mathbb{E} W_2(\hat{P}_t^{(n)}, \bar{P}_t^{(n)}) \lesssim C_K n^{-1/2}$. The main text uses only that regime. More generally, if the finite-sample term admits the scaling

$$\mathbb{E} W_2(\hat{P}_t^{(n)}, \bar{P}_t^{(n)}) \lesssim C_K n^{-a} \quad \text{for some } a > 0,$$

then balancing it against the drift-induced staleness cost $\frac{1}{2}\zeta n$ yields

$$n^*(a) = \left(\frac{2a C_K}{\zeta} \right)^{\frac{1}{a+1}}, \quad \mathcal{E}_{\min}(a) = \frac{a+1}{(2a)^{a/(a+1)}} C_K^{1/(a+1)} \zeta^{a/(a+1)}.$$

In the setting of lemma 2.4, the effective exponent is $a = 1/2$ for $d_{\text{eff}} < 4$, with an additional $(\log n)^{1/2}$ factor for $d_{\text{eff}} = 4$; for $d_{\text{eff}} > 4$ the raw Wasserstein rate is slower than $n^{-1/2}$. That case lies outside the displayed cube-root regime. The present result isolates the worst-case linear-staleness envelope induced by the local Lipschitz path bound.

The two terms in theorem 2.3 pull in opposite directions. Increasing n reduces variance but raises the age of the retained evidence. That tension is the basic variance-staleness trade-off that governs useful memory under drift.

Proposition 2.6 (Optimal Window Under Drift). *Under theorem 2.3, the error $\mathcal{E}(n)$ is minimised at*

$$n^* = \left(\frac{C_K}{\zeta} \right)^{2/3},$$

with minimum value

$$\mathcal{E}_{\min} = \frac{3}{2} C_K^{2/3} \zeta^{1/3}.$$

Remark 2.7 (Worst-case Lipschitz regime). The cube-root law is the worst-case Lipschitz member of a broader horizon family.

Proof. Apply the triangle inequality for W_2 :

$$W_2(\hat{P}_t^{(n)}, P_t^*) \leq W_2(\hat{P}_t^{(n)}, \bar{P}_t^{(n)}) + W_2(\bar{P}_t^{(n)}, P_t^*).$$

The first term is controlled by the assumed finite-sample bound. For the second term, let π_j be an optimal coupling between P_{t-j}^* and P_t^* . Then the averaged coupling $\bar{\pi} := n^{-1} \sum_{j=0}^{n-1} \pi_j$ couples $\bar{P}_t^{(n)}$ with P_t^* , so the transport cost is at most the average of the individual costs. Combined with the W_2 -Lipschitz drift bound, this gives

$$W_2(\bar{P}_t^{(n)}, P_t^*) \leq \frac{1}{n} \sum_{j=0}^{n-1} W_2(P_{t-j}^*, P_t^*) \leq \frac{1}{n} \sum_{j=0}^{n-1} \zeta j \leq \frac{1}{2} \zeta n.$$

Combining the two bounds yields the stated inequality. $\square$

Remark 2.8 (Measurement layer). The theorem is stated in W_2 , so ε does not enter the mathematical drift law itself. In the synthetic measurement layer, however, the observable geometry is often a debiased Sinkhorn divergence with fixed ε , which makes the finite-sample constant calibration-dependent. In the synthetic benchmarks, $\hat{\zeta}_t$ is estimated from short-window block differences using an EMA update and C_K is calibrated from a stationary prefix; Appendix E records the sensitivity check and the prequential calibration rule.

For fixed ε on bounded domains, the debiased Sinkhorn divergence converges at rate $n^{-1/2}$ with ε -dependent constants [Genevay et al., 2019], while the approximation gap to W_2^2 is $O(\varepsilon)$ and vanishes as $\varepsilon \rightarrow 0$. We therefore use Sinkhorn as a calibration proxy in the measurement layer.

Corollary 2.9 (EWMA Analogue). *For exponential weights $w_j = \alpha(1 - \alpha)^j$, the effective lag is $\tau_\alpha = (1 - \alpha)/\alpha$ and the effective sample size scales as $n_{\text{eff}} \asymp 1/\alpha$. Balancing the same statistical and staleness terms therefore yields the same cube-root exponent in the effective memory scale. This gives the exponentially weighted analogue of the Lipschitz case.*

2.4 A Structural Finite-Memory Floor




Figure 2: Lower-bound witness on the Gaussian location subclass. The two paths share most of the retained history and differ only on a short recent segment of width h , yet they remain separated at the present. That unresolved recent segment is enough to create a structural finite-memory floor.

Proposition 2.10 (Finite-Memory Floor via Gaussian Location Subclass). *Let $\mathcal{P}_{\zeta,m}$ be the class of length- m distribution paths for which the following hold:*

1. *the path is W_2 -Lipschitz with rate at most ζ ;*
2. *the lower-bound construction is carried out on the Gaussian location subclass $X_t \sim \mathcal{N}(\mu_t, \sigma^2 I_d)$ with $\|\mu_t - \mu_s\|_2 \leq \zeta|t - s|$;*
3. *the estimator at time m only uses the last m samples.*

Measure risk by the endpoint W_2 error, which coincides with Euclidean endpoint error in the Gaussian location subclass. Then, in the critical-window regime, the minimax tracking error over that class cannot beat the same variance-plus-staleness scaling:

$$\inf_{\hat{\mathcal{P}}_m} \sup_{\mathcal{P}^*} \mathbb{E} d(\hat{\mathcal{P}}_m, \mathcal{P}^*) \gtrsim \max\{C_K m^{-1/2}, c \sigma^{2/3} \zeta^{1/3}\}.$$

Consequently, the cube-root horizon is not merely the optimum of one estimator; Consequently, the cube-root horizon is the structural floor of finite memory on this drift class.

Proof. Write $R_m^* := \inf_{\delta_m} R_m(\delta_m)$. We exhibit two independent obstructions. The first is the ordinary finite-sample estimation floor, obtained by restricting to the constant-path subclass $\mu_s \equiv \theta$, which is contained in the admissible drift class. For the two hypotheses $\theta_{\pm} = \pm \sigma / (2\sqrt{m})$, the induced product-Gaussian laws have Kullback–Leibler divergence

$$\text{KL}(P_+, P_-) = \frac{1}{2\sigma^2} \sum_{j=1}^m (\theta_+ - \theta_-)^2 = \frac{1}{2\sigma^2} m \left(\frac{\sigma}{\sqrt{m}} \right)^2 = \frac{1}{2}.$$

For any estimator δ_m , the plug-in test $\psi = \mathbf{1}\{\delta_m \geq 0\}$ must err whenever the estimate falls on the wrong side of the origin; on that event the endpoint error is at least $\sigma / (2\sqrt{m})$. By Le Cam’s two-point lemma and Pinsker’s inequality, the average testing error is at least

$$\frac{1 - \text{TV}(P_+, P_-)}{2} \geq \frac{1 - \sqrt{\text{KL}(P_+, P_-)/2}}{2} \geq \frac{1}{4}.$$

Hence

$$R_m^* \geq \frac{\sigma}{8\sqrt{m}},$$

so one valid explicit choice is $c_0 = 1/8$.

The second obstruction is local drift. Fix any $1 \leq h \leq m$ and define

$$\beta_h := \min\left\{\zeta, \frac{\sigma}{4h^{3/2}}\right\}, \quad \mu_{t-j}^{\pm, h} := \pm \beta_h (h - j)_+, \quad 0 \leq j \leq m - 1,$$

where $(u)_+ := \max\{u, 0\}$. Each path belongs to the admissible class because its slope magnitude is at most $\beta_h \leq \zeta$. The endpoint gap is

$$|\mu_t^{+, h} - \mu_t^{-, h}| = 2\beta_h h.$$

The corresponding Gaussian product laws satisfy

$$\text{KL}(P_{+, h}, P_{-, h}) = \frac{1}{2\sigma^2} \sum_{j=0}^{m-1} (2\beta_h (h - j)_+)^2 \leq \frac{2\beta_h^2}{\sigma^2} \sum_{r=1}^h r^2 \leq \frac{2\beta_h^2}{\sigma^2} h^3 \leq \frac{1}{8}.$$

Applying the same reduction from estimation to testing, now with endpoint error threshold $\beta_h h$, and using Le Cam together with Pinsker gives average testing error at least

$$\frac{1 - \text{TV}(P_{+, h}, P_{-, h})}{2} \geq \frac{1 - \sqrt{\text{KL}(P_{+, h}, P_{-, h})/2}}{2} \geq \frac{3}{8}.$$

Therefore

$$R_m^* \geq \frac{3}{8} \beta_h h \geq \frac{3}{32} \min\{\zeta h, \sigma h^{-1/2}\}$$

which is again a fully explicit numerical lower bound. Since h was arbitrary,

$$R_m^* \geq \frac{3}{32} \sup_{1 \leq h \leq m} \min\{\zeta h, \sigma h^{-1/2}\}.$$

Taking the maximum of the constant-path and ramp bounds proves the displayed lower bound. For the cube-root regime, assume $1 \leq h_* := (\sigma/\zeta)^{2/3} \leq m$. Choose an integer $\bar{h} \in [h_*/2, h_*]$. Then

$$\zeta \bar{h} \geq \frac{1}{2} \zeta h_* = \frac{1}{2} \sigma^{2/3} \zeta^{1/3},$$

and

$$\sigma \bar{h}^{-1/2} \geq \sigma h_*^{-1/2} = \sigma^{2/3} \zeta^{1/3}.$$

Therefore

$$\min\{\sigma \bar{h}^{-1/2}, \zeta \bar{h}\} \geq \frac{1}{2} \sigma^{2/3} \zeta^{1/3},$$

and the ramp bound gives

$$R_m^* \geq \frac{3}{64} \sigma^{2/3} \zeta^{1/3}.$$

So one admissible explicit constant in the critical-window regime is $c = 3/64$.

Remark 2.11 (Scope of the lower bound). The lower bound is parametric only in the testing construction: it uses a Gaussian location subclass whose mean path is itself W_2 -Lipschitz, and the distance in the lower bound is the endpoint W_2 error (Euclidean on that subclass). That is enough to make the result minimax for any larger class containing that subclass, including the ambient W_2 -Lipschitz path class. So the proposition is not claiming that every path is Gaussian; it is claiming that the worst case already appears inside the Gaussian mean-shift subproblem. Equivalently, since the Gaussian location family is a subfamily of the ambient W_2 -Lipschitz class at the same drift rate ζ , any lower bound proven on that subclass is automatically a lower bound for the larger class.

2.5 Holder Path Classes and Horizon Laws

The preceding proposition already states the family law. This subsection keeps the Holder terminology explicit for the reader and emphasizes that the Lipschitz case is the $H = 1$ member of the same family.

2.6 Temporal Validity and Memory Age

Definition 2.12 (Effective Memory Age). Let an estimator place weights $w_0, \dots, w_{n-1}$ on the most recent observations, with $\sum_j w_j = 1$. Its *effective memory age* is

$$A(w) = \sum_{j=0}^{n-1} j w_j,$$

and its corresponding freshness proxy is the monotone transform $F(w) = 1/(1 + A(w))$. The role of $F(w)$ is purely geometric: it records how old the retained evidence is before any local drift scale is specified. For a uniform window, $A(w) = (n-1)/2$.

Definition 2.13 (Temporal Validity). For an estimator with weights w at time t , define its *operational staleness*

$$S_t(w) = \sum_{j=0}^{n-1} w_j W_2(\mathcal{P}^*(S, t-j), \mathcal{P}^*(S, t)).$$

Its *temporal validity* is the inverse alignment score

$$C_t(w) = \frac{1}{1 + S_t(w)}.$$

If the target path is locally W_2 -Lipschitz at rate ζ_t , in the sense that $W_2(\mathcal{P}^*(S, t-j), \mathcal{P}^*(S, t)) \leq \zeta_t j$ for the relevant lags, then

$$S_t(w) \leq \zeta_t A(w), \quad C_t(w) \geq \frac{1}{1 + \zeta_t A(w)}.$$

In that sense, $F(w)$ is the rate-free freshness proxy, while $C_t(w)$ measures operational validity with the current distribution.

The quantity $A(w)$ gives a compact description of the temporal age of a finite-memory estimator. Short windows preserve temporal validity at the cost of variance; long windows reduce variance at the cost of staleness. In the worst-case Lipschitz regime, the cube-root horizon marks the point where that exchange stops improving. For a uniform window under approximately constant local drift, the induced staleness scales as $S_t(w) \approx \zeta(n-1)/2$, which is the same age term that appears in theorem 2.3 up to constants.

Remark 2.14 (Measurement layer). Debiased Sinkhorn divergence serves as a practical measurement layer in the synthetic experiments when the underlying geometry must be estimated from samples. That computational choice supports the tracking law, and the law itself concerns the horizon of useful memory.

3 Empirical Evidence

The experiments show that under continuous drift, memory has a finite useful-memory horizon and can become temporally invalid before changepoint evidence appears. The evidence separates layers that are often collapsed: the useful-memory horizon itself, the gap between temporal validity and detector evidence, the resulting stale-memory regime, and the predictive cost of horizon misalignment.

In the experiments, UMR is a backend-agnostic conservative horizon cap driven by a smoothed local drift proxy. It is evaluated in two roles. First, it acts as a runtime horizon policy for estimators whose effective memory can be adjusted, such as capped sliding-window, EWMA, and multi-horizon backends. Second, it is wrapped around ADWIN as one detector-side instance of the same temporal-validity rule. Fixed-window methods remain uncapped baselines. The resulting comparisons expose detector-silent staleness.

3.1 Structural U-Curve

We begin with the Gaussian sweep used to validate the finite-memory trade-off. The setup produces the characteristic U-curve: for small windows, error is dominated by variance; for large windows, error is dominated by staleness. The empirical minimum remains finite and moves in the expected direction as drift strengthens, and EWMA follows the same finite-memory pattern on its effective memory scale.




Figure 3: Finite-memory law validation on the Gaussian sweep. Left: tracking error as a function of window size, showing the predicted U-shape for multiple drift rates. Right: the minimum achievable error decreases as drift weakens, consistent with the variance-staleness trade-off. The figure shows that the useful-memory optimum is finite and stable across drift rates.

The same pattern is visible from the lag-variance frontier in fig. 4, which exposes the frontier knee where extra memory stops paying off. Bootstrap intervals in Appendix C support the same ordering.

3.2 Temporal Validity Is Not Changept Evidence

The detector asks whether there is enough statistical evidence of a change. The horizon asks whether the retained evidence is still temporally valid for the present. Those questions need not be synchronized.

3.3 The Cap-Only Regime

The gap becomes concrete on the continuous-drift benchmark. We compare a short fixed window, a long fixed window, EWMA, ADWIN, and an ADWIN wrapper equipped with the UMR horizon cap. ADWIN can retain far more memory than the useful-memory scale, while a conservative cap contracts that stale window earlier.

3.4 Recovery Asymmetry Under Alternating Drift

The next benchmark makes local horizon instability visible. The stream alternates between slow drift, fast bursts, recovery plateaus, and micro-bursts with irregular




Figure 4: Lag-variance frontier for useful memory. The fixed-window sweep traces the empirical frontier: short horizons are noisy, long horizons are stale, and the knee marks the useful-memory optimum. The frontier is asymmetric because the gain from more memory saturates while the cost of age keeps accumulating.

Table 1: Continuous-drift benchmark. The key gap is between statistical evidence and temporal validity: ADWIN keeps a much wider window than the UMR-capped horizon scale, while the long fixed window collapses under lag.

Method	Tail MAE	Mean horizon	Detector events	Cap events
Fixed-100	0.0849	100.0	—	—
Fixed-500	0.2388	500.0	—	—
EWMA	0.1291	39.0	—	—
ADWIN	0.1757	335.8	3.9	0.0
ADWIN + UMR	0.0978	88.9	0.0	1374.7

durations, so the local oracle horizon need not be globally stable. In that setting, a single fixed window can look reasonable on average while remaining locally misaligned for long stretches.

The benchmark makes two things visible. First, horizon validity can vary even when aggregate prediction error stays competitive. Second, the adjustment is asymmetric: contraction of the useful horizon is easier to trigger than re-expansion after a low-drift stretch. Proposition 3.1 formalizes the same asymmetry in the idealized EMA-to-horizon map. The transition-window table and the EMA-sensitivity sweep appear in Appendix C.

Proposition 3.1 (Asymmetric Horizon Recovery). *Let the drift estimator satisfy the EMA recursion*

$$\hat{\zeta}_t = \alpha \Delta_t + q \hat{\zeta}_{t-1}, \quad q := 1 - \alpha \in (0, 1),$$




Figure 5: Cap leads detection under continuous drift. The top panel compares the first ADWIN+UMR cap event with the first ADWIN detection, and the bottom panel shows the resulting lead time. Under slow drift, memory becomes stale well before changepoint evidence is available; under faster drift, the gap narrows but remains positive.

with $\Delta_t = \delta_{\text{slow}}$ in the slow regime and $\Delta_t = M$ in the burst regime. Assume the horizon map is unclipped on $[\delta_{\text{slow}}, M]$ so that

$$n^*(\delta) = \left(\frac{C_K}{\delta} \right)^{2/3}$$

throughout the transition. Write $n_b := n^*(M)$ and $n_s := n^*(\delta_{\text{slow}})$. For any $0 < \varepsilon < n_s - n_b$, define the inverse thresholds

$$\delta_b(\varepsilon) := C_K (n_b + \varepsilon)^{-3/2}, \quad \delta_s(\varepsilon) := C_K (n_s - \varepsilon)^{-3/2}.$$

Then the exact recovery time from the slow regime to the burst regime is

$$T_{\text{contract}}(\varepsilon) = \left\lceil \frac{\log((M - \delta_{\text{slow}})/(M - \delta_b(\varepsilon)))}{\log(1/q)} \right\rceil,$$

while, after a burst of duration T_{burst} steps, the exact recovery time back to the slow regime is

$$T_{\text{expand}}(\varepsilon) = \left\lceil \frac{\log((\hat{\zeta}_{\text{end}} - \delta_{\text{slow}})/(\delta_s(\varepsilon) - \delta_{\text{slow}}))}{\log(1/q)} \right\rceil,$$



Figure 6: Cap-only regime on the continuous-drift benchmark. Fixed-500 is the stale-memory baseline, plain ADWIN can retain far more memory than the useful-memory scale, and ADWIN+ UMR contracts the effective horizon earlier. The gain comes from capping stale memory before detector evidence is strong enough to fire.

where

$$\hat{\zeta}_{\text{end}} = M - (M - \delta_{\text{slow}})q^{T_{\text{burst}}}$$

is the burst-end drift estimate. If

$$T_{\text{burst}} \geq T_0(\varepsilon) := \left\lceil \frac{\log(1 - \rho_\varepsilon)}{\log q} \right\rceil, \quad \rho_\varepsilon := \frac{\delta_s(\varepsilon) - \delta_{\text{slow}}}{M - \delta_b(\varepsilon)},$$

then $T_{\text{expand}}(\varepsilon) > T_{\text{contract}}(\varepsilon)$. Moreover, for small ε ,

$$M - \delta_b(\varepsilon) = \frac{3}{2}C_K n_b^{-5/2} \varepsilon + O(\varepsilon^2), \quad \delta_s(\varepsilon) - \delta_{\text{slow}} = \frac{3}{2}C_K n_s^{-5/2} \varepsilon + O(\varepsilon^2),$$

so the slow-side threshold distance is smaller by a factor $(\delta_{\text{slow}}/M)^{5/3}(1 + o(1))$.

Proof. The EMA recursion has the closed form

$$\hat{\zeta}_{t_0+k} = \Delta^* + (\hat{\zeta}_{t_0} - \Delta^*)q^k$$

whenever the regime is constant with fixed point Δ^* . Hence after the transition from δ_{slow} to M ,

$$\hat{\zeta}_{t_0+k} = M - (M - \delta_{\text{slow}})q^k,$$

and after the return to the slow regime,

$$\hat{\zeta}_{t_0+T_{\text{burst}}+k} = \delta_{\text{slow}} + (\hat{\zeta}_{\text{end}} - \delta_{\text{slow}})q^k.$$




Figure 7: Alternating-timescales benchmark. The top panel shows the drift schedule, the middle panel compares the local oracle horizon with fixed windows and ADWIN+UMR, and the bottom panel shows horizon regret over time. Contraction is easier to trigger under bursts than re-expansion after a low-drift stretch.

Because $n^*(\delta)$ is monotone decreasing on the unclipped interval and its inverse is $g(n) = C_K n^{-3/2}$, the condition $|n^*(\hat{\zeta}_{t_0+k}) - n_b| \leq \varepsilon$ is equivalent to $\hat{\zeta}_{t_0+k} \geq \delta_b(\varepsilon)$, which gives the displayed formula for $T_{\text{contract}}(\varepsilon)$. Likewise, $|n^*(\hat{\zeta}_{t_0+T_{\text{burst}}+k}) - n_s| \leq \varepsilon$ is equivalent to $\hat{\zeta}_{t_0+T_{\text{burst}}+k} \leq \delta_s(\varepsilon)$, which yields $T_{\text{expand}}(\varepsilon)$.

The inequality $T_{\text{expand}}(\varepsilon) > T_{\text{contract}}(\varepsilon)$ holds whenever the burst lasts long enough that

$$1 - q^{T_{\text{burst}}} \geq \rho_\varepsilon,$$

since this makes the numerator of T_{expand} larger than that of T_{contract} by the exact threshold relation above. The small- ε expansions follow by Taylor expansion of $n \mapsto C_K n^{-3/2}$ around n_b and n_s . $\square$

Remark 3.2 (Effect of clipping). If clipping at $n_{\min}$ is active near the burst side, then contraction can only be faster. The observed asymmetry is therefore conservative: clipping weakens the theoretical barrier in the burst regime, but it cannot eliminate the longer slow- side recovery. In streams with gradual bursts, the observed ratio can exceed the step-burst threshold because $\hat{\zeta}_{\text{end}}$ does not reach M .

3.5 Horizon Misalignment Carries Cost

The next question is when horizon misalignment becomes operationally expensive. For each backend, we compare its instantaneous horizon n_t with the local horizon target n_t^* and plot the excess error $E_t - E_t^*$ against both the absolute gap $|n_t - n_t^*|$ and the relative gap $|\log(n_t/n_t^*)|$. Here E_t^* is the best error achieved at time t by the fixed-window oracle grid, and the curve is binned over all seeds and time points.




Figure 8: Excess error as a function of horizon misalignment. The relation is backend dependent: fixed-100 stays inside a low-cost band, whereas fixed-500 and ADWIN pay substantially more once the gap widens. ADWIN+ UMR sits between these extremes.

These results expose a measurable tolerance band around the useful-memory scale, together with backend-specific degradation once that band is exceeded.

3.6 Representative Real-Stream Case Study

The synthetic experiments carry the clearest identification, but ELEC2 provides the same qualitative phenomenon on public data. Coarse horizon adaptation still changes the operating regime outside the synthetic setting, although the effect is smaller and more backend dependent. A Bikes summary is given in appendix F.

Additional stability and sensitivity results appear in the appendix: bootstrap intervals and EMA-sensitivity in Appendix C, phase-wise recovery in Appendix B, the Bikes summary in Appendix F, and the controlled downstream check in Appendix G.

3.7 What the Evidence Shows

Taken together, the evidence supports five claims. First, the finite-memory optimum is real. Second, temporal validity and changepoint evidence are distinct objects. Third, the cap-only regime is operationally visible. Fourth, horizon recovery is asymmetric under alternating timescales. Fifth, horizon control has operational value, but its predictive payoff depends on the backend and on the drift regime.




Figure 9: Public-stream case study on ELEC2. The real-stream effect is smaller and more backend dependent than in the synthetic benchmarks, but coarse horizon adaptation still changes the operating regime on public data.

4 Related Work

Tracking under nonstationarity. Dynamic regret, online least-squares, and adaptive filtering ask how much past evidence should remain relevant once the target is moving [Besbes et al., 2015, Cheung et al., 2019, Yuan and Lamperski, 2020, Mazzetto and Upfal, 2023]. Here the same question is kept in distributional form, with the memory horizon as the main object of analysis.

Variation budgets and horizon selection. The $\zeta^{1/3}$ exponent also appears in non-stationary stochastic optimization and bandits [Besbes et al., 2015, Cheung et al., 2019, Wang et al., 2024, Yuan and Lamperski, 2020]. Those results concern gradient-based methods on convex losses in Euclidean space. Here, drift is W_2 -Lipschitz on measures, the estimator class is causal averaging kernels, and Proposition 2.10 makes the $\zeta^{1/3}$ rate a minimax lower bound for arbitrary window-restricted estimators.

The same cube-root tuning shows up in sliding-window analyses of drifting data [Bifet and Gavaldà, 2007]: the window length is the control variable, but the quantity being protected is temporal validity. The same perspective appears here, but with the geometry changed from scalar losses to distributional tracking.

Jiang, Li, and Zhang study online stochastic optimization under a Wasserstein nonstationarity measure and obtain distribution-aware regret guarantees [Jiang

et al., 2020]. Keehan, Anderson, and Wiesemann analyze a related drift-aware weighting problem in Wasserstein distributionally robust optimization [Keehan et al., 2025]. The present focus is narrower: when memory itself is finite, how long can retained evidence remain useful before age becomes the dominant source of error?

Transport geometry. Wasserstein distance and Sinkhorn-type estimators are natural because they respect the geometry of the signal space [Genevay et al., 2019, Niles-Weed and Bach, 2019, Goldfeld et al., 2024, Akyildiz et al., 2024]. In our setting, geometry quantifies how much old evidence can be retained before it becomes stale.

Adaptive forgetting and detection. EWMA, forgetting-factor RLS, and Kalman filtering all balance variance reduction against responsiveness [Jazwinski, 1970]. The cube-root law makes temporal validity explicit. The asymmetry result is adjacent to changepoint detection and EMA-based control, but the contribution here is the gap between those questions: once drift evidence is mapped into a horizon rule, re-expansion can be slower than contraction.

This also connects to the concept-drift literature more broadly [Gama et al., 2004, 2014]. In particular, the “Window Dilemma” argument of Hinder et al. [2023] emphasizes that drift detection is ill-posed when the window scale is itself uncertain. The same issue appears here in a different form: temporal validity is not fixed in advance, so the horizon becomes part of the object that must be estimated.

Public-stream baselines. Window-Dilemma-style and MELO-style baselines widen the empirical comparison. The comparison is broader, and the main claim stays the same: the cube-root regime describes the worst-case horizon law for the class studied here.

Multi-horizon ensembles attack the same practical issue from another angle by maintaining several historical windows and combining their predictions through uncertainty or expert weighting. The public-stream experiments use that family as a neighboring design point in the comparison suite [Rajput et al., 2025].

Freshness and Age of Information. The Age-of-Information literature asks how old the last received update is and how update policies should balance communication cost against information staleness [Kaul et al., 2012, Yates et al., 2021]. The distributional analogue is the same question under W_2 drift: useful memory has a finite worst-case horizon, and the cube-root law quantifies that worst-case scale for the class studied here.

5 Limitations and Scope

Scope of the law. The cube-root law is exact for the finite-memory averaging class studied here, and the lower bound is stated for window-restricted estimators on a W_2 -Lipschitz drift class. The lower bound itself is proved through a Gaussian location subclass and endpoint W_2 error, so its structural claim is worst-case and subclass-based.

More refined drift classes may admit tighter staleness growth and different horizon exponents. The analysis here isolates the linear-staleness law induced by the W_2 -Lipschitz path bound.

Dependence on local drift estimation. The operational regulator depends on a local estimate of drift. The current benchmark suggests that the remaining gap to the oracle comes mostly from drift estimation. At the same time, a short observable calibration prefix can materially improve the regulator, so drift estimation is part of the method and enters the operational picture directly. Better local geometry estimators may therefore sharpen horizon recovery without changing the law.

Geometry beyond the current metric. The formulation uses Wasserstein geometry and a debiased Sinkhorn measurement layer when samples must be compared computationally. Other geometries may admit analogous trade-offs, but the constants and the calibration of useful memory may change.

What the experiments show. The experiments show the operational distinction between temporal validity and changepoint evidence, together with the cap-only regime that distinction creates. The public-stream results provide supporting evidence with weaker identification than the synthetic benchmarks.

Scope of the comparison suite. The benchmark suite now includes a Window-Dilemma-style switcher and a MELO-style multi-horizon hedge, but it does not cover the full space of adaptive tracking methods. The analysis remains centered on the worst-case cube-root law and on the temporal-validity interpretation of memory.

The drift proxy itself has a practical resolution limit: very small block sizes can grossly overestimate $\hat{\zeta}_t$, so the calibration prefix needs blocks large enough to smooth local noise. In the current sweep, $d = 50$ is acceptable, while $d = 100$ is effectively stable.

On the current ELEC2 and Bikes arenas, adding UMR to the multi-horizon hedge changes results only modestly and in a benchmark-dependent way; it acts as a horizon prior within those experiments.

Path-class refinement. The linear staleness term comes from the local W_2 -Lipschitz path bound and is therefore a worst-case statement for that class. Characterizing narrower drift classes in which temporal misalignment accumulates more slowly, together with matching lower bounds, remains open.

6 Conclusion

How long can memory remain useful when the target distribution keeps moving? Under the drift class studied here, useful memory has a finite temporal horizon, and in the worst-case linear-staleness regime that horizon obeys a cube-root law. The governing object is a variance-staleness trade-off: preserving temporally valid evidence requires limiting the lag induced by old memory.

The theoretical contribution is a decomposition of tracking error into a variance term and a staleness term, together with a lower bound showing that the resulting

finite-memory floor is structural. The operational contribution is a horizon regulator: UMR turns the worst-case law into a temporal-validity cap, which keeps the effective horizon from drifting far into stale-memory territory.

The strongest empirical lesson is that temporal validity and changepoint evidence are distinct objects. The signature regime is *cap-only*: the memory horizon becomes operationally stale before there is enough statistical evidence for a changepoint alarm. Horizon misalignment carries a nonlinear operational cost, but the transfer from gap to predictive error depends on the backend and on the drift regime. The alternating-timescales benchmark adds a secondary lesson: contraction is easier to induce than re-expansion, so EMA-driven horizon regulators have a recoverability gap.

The real-stream evidence is intentionally secondary. ELEC2, with a short Bikes summary in the appendix, shows that the temporal-validity gap is not synthetic-only, but the main claim of the paper remains worst-case and structural.

Finite-memory systems are best judged by how long their retained evidence remains current enough to matter. The design consequence is simple: treat memory as a resource with a finite horizon, separate temporal validity from changepoint evidence, and use conservative horizon regulation to avoid stale-memory carryover.

The cube-root exponent is tied to the worst-case linear-staleness envelope studied here, and the lower bound is established through a Gaussian location subclass. The next step is correspondingly disciplined: characterize narrower path classes and better local drift estimators without diluting the current worst-case statement.

Code and reproducibility artifacts are available in the project repository [Parreira, 2026].

Acknowledgements

With gratitude to everyone who contributed ideas, feedback, tools, encouragement, and care along the way. This work stands on the generosity of colleagues, collaborators, open-source maintainers, and the broader research communities that made it possible.

A Backend + UMR Protocol

When a backend is paired with the useful-memory regulator (UMR), the runtime protocol is a shared backend-agnostic horizon cap: it estimates a local drift scale, smooths that estimate with an EMA, and converts it into a runtime horizon bound. The shared drift proxy and horizon rule are implemented in `drift/umr.py`. The canonical manuscript-facing entry points are the generation scripts in `code/`, which write final outputs to `artifacts/figures/`, `artifacts/csv/`, and `artifacts/tables/`.

A.1 Shared Regulator

At each time step, the regulator maintains a local drift estimate and a target horizon:

$$\Delta_t = \frac{|\bar{x}_t^{(d)} - \bar{x}_{t-d}^{(d)}|}{d}, \quad \hat{\zeta}_t = \alpha \Delta_t + (1 - \alpha) \hat{\zeta}_{t-1},$$

$$n_t^* = \text{clip} \left(\left(\frac{C_K}{\max(\hat{\zeta}_t, 10^{-9})} \right)^{2/3}, n_{\min}, n_{\max} \right).$$

There is no explicit cooldown parameter. The only stabilization is clipping through $n_{\min}$ and $n_{\max}$, plus the EMA smoothing in $\hat{\zeta}_t$.

A.2 Horizon-Capped Backends

For estimators whose natural state is an adjustable window size or effective horizon (sliding-window mean, EWMA, Window-Dilemma-style switching, MELO-style expert mixing), the runtime horizon is simply clipped by the regulator:

$$h_t^{\text{eff}} = \min(h_t^{\text{backend}}, n_t^*).$$

The backend then computes its estimate using h_t^{eff} exactly as it would have used h_t^{backend} .

Uncapped fixed-window baselines are separate controls: they hold their window length constant throughout the run and do not receive the regulator.

A.3 Detector Wrapper Instance

One detector-side instantiation uses ADWIN on a scalar signal. The wrapper maintains a rolling buffer of the last $n_{\max}$ observations and applies the same cap logic to the detector width:

1. Update the drift proxy and compute n_t^* .
2. Feed the new observation to ADWIN.
3. Read the current ADWIN width w_t .
4. If $w_t > n_t^*$, record a cap event, reset ADWIN, and re-feed the most recent n_t^* buffered observations to the fresh detector.
5. If the detector then reports drift, record a warning.

The cap event is therefore a horizon event, not a detector event. The detector warning is still ADWIN’s own decision after the reset path. This wrapper is one instance of the shared UMR rule, not the defining form of the regulator.

A.4 Hyperparameters

Three parameter families are used.

Synthetic core suite.

Table 2: Synthetic benchmark defaults used by `code/generate_cuberoot_adwin.py`.

Scenario	Default hyperparameters
Constant	<code>drift = 0.001</code> , <code>fixed_window = 100</code> , <code>fixed_long_window = 500</code> , <code>ewma_alpha = 0.05</code> , <code>adwin_delta = 0.002</code> , <code>Ck = 1.0</code> , <code>drift_window = 100</code>
Piecewise	<code>piecewise_drifts = (0.0005, 0.003, 0.001)</code> with the same window, EWMA, ADWIN, and C_K settings as above
Alternating	<code>piecewise_drifts = (0.0002, 0.0045, 0.00015, 0.007, 0.00025)</code> , <code>piecewise_lengths = (900, 140, 700, 70, 1190)</code> , <code>fixed_long_window = 200</code> , and the same remaining defaults
Common	$n_{\min} = 10$, $n_{\max} = 500$

Real-stream arenas. The ELEC2 and Bikes arenas are generated by `code/generate_elec2.py` and `code/generate_bikes.py` with:

- `fixed_window = 100`, `fixed_short_window = 50`, `fixed_long_window = 300`;
- `ewma_alpha = 0.05`;
- `block_size = 24`, `ema_alpha = 0.03`;
- `min_window = 30`, `max_window = 300`, `scale = 1.25`;
- `baseline_window = 100`, `prefix_length = 2000`;
- `adwin_delta = 0.03`;
- Page-Hinkley event detector: $\delta = 0.01$, threshold 84.0, $\alpha = 0.9999$;
- Window-Dilemma windows (50, 100, 300) with quantiles (0.33, 0.67);
- MELO-style expert windows (30, 50, 100, 200, 300) with learning rate 6.0 and discount 0.995.

Drift-proxy calibration. When C_K is not fixed externally, the code estimates it from a stationary prefix via `calibrate_umr_constant`, i.e. standard deviation times the square root of prefix length. Appendix E additionally reports the prefix-validated (d, α) sensitivity study, generated by `code/generate_zeta_support.py` and `code/generate_zeta_self_calibration.py`; that robustness analysis does not replace the fixed-regulator core.

A.5 Algorithm

The scripts in `code/` instantiate this protocol and write the resulting artifacts to the canonical output directories under `artifacts/`.

Algorithm 1. UMR update rule.

1. **Input:** stream $x_{1:T}$, backend B , prefix length m , block size d , EMA coefficient α , bounds $[n_{\min}, n_{\max}]$, constant C_K .
2. For each $t = 1, \dots, T$:
 - (a) Update Δ_t from the last two blocks of length d .
 - (b) Update $\hat{\zeta}_t = \alpha\Delta_t + (1 - \alpha)\hat{\zeta}_{t-1}$.
 - (c) Compute $n_t^* = \text{clip}((C_K / \max(\hat{\zeta}_t, 10^{-9}))^{2/3}, n_{\min}, n_{\max})$.
 - (d) If B is horizon-based, set the effective horizon to $\min(h_t^B, n_t^*)$ and evaluate the backend with that horizon.
 - (e) If B is a detector wrapper, update the detector; if its internal width exceeds n_t^* , reset it and re-feed the last n_t^* buffered observations; then record any detector warning.
3. **Output:** backend estimate, warnings, n_t^* , drift proxy, and cap events.

Figure 10: UMR update rule.

B Phase-Wise Recovery Details

The phase-wise calibration check for the UMR-regulated ADWIN benchmark shows that the same calibration stays in the right operational range when the drift regime changes discretely.

Table 3: Oracle horizon recovery on piecewise drift. The same calibrated cap tracks the oracle scale across regimes: it is close in the fast-drift phase and remains within the right order of magnitude in the smoother phases.

Drift	Oracle window	Regulated n^*	Oracle MAE	ADWIN + UMR MAE
0.0005	200	191.2	0.0741	0.0874
0.0030	50	58.1	0.1258	0.1381
0.0010	100	89.7	0.0876	0.1033

C Supplementary Empirical Checks

Bootstrap intervals, transition-window summaries, and EMA sensitivity checks support the same qualitative conclusions as the primary empirical results.




Figure 11: Piecewise-drift recovery. The same regulator calibration follows the phase-dependent oracle window closely enough to preserve the right operational scale in each regime.

Table 4: Seed-wise bootstrap summary for the main synthetic benchmark. The first column block reports tail-MAE means with 95% bootstrap intervals across the 20 seeds; the second block reports paired differences versus ADWIN+ UMR.

Method	Tail MAE [95% CI]	Δ vs ADWIN+ UMR [95% CI]
Fixed-100	0.0849 [0.0815, 0.0884]	-0.0129 [-0.0143, -0.0115]
Fixed-500	0.2388 [0.2275, 0.2500]	+0.1409 [+0.1319, +0.1502]
EWMA	0.1291 [0.1263, 0.1319]	+0.0313 [+0.0278, +0.0350]
ADWIN	0.1757 [0.1658, 0.1852]	+0.0779 [+0.0706, +0.0857]
ADWIN + UMR	0.0978 [0.0943, 0.1016]	—

Table 5: Mean horizon regret over the first 50 steps after each regime change. Contraction averages over the two fast-drift transitions, and expansion averages over the two slow-drift recoveries. Lower is better.

Method	Contraction	Expansion
fixed-100	68.0	203.1
fixed-200	168.0	103.1
ADWIN	789.5	477.7
ADWIN + UMR	67.1	248.1

C.1 Bootstrap Stability Check

C.2 Transition-Window Asymmetry Table

C.3 Drift-EMA Sensitivity

The recovery asymmetry persists when the drift estimator itself is varied. The next sweep changes the EMA coefficient α in the drift estimator while keeping the alternating-timescales benchmark fixed. Larger α makes the drift estimate more reactive, reducing contraction regret while expansion regret stays persistently larger; the expansion-to-contraction ratio remains above one throughout the sweep and is higher at the larger α values.




Figure 12: EMA sensitivity of horizon recovery on alternating timescales. The drift-EMA coefficient changes how quickly the regulator reacts to bursts, but the asymmetry remains: expansion regret stays above contraction regret across the sweep.

D Finite-Sample Geometry Check

For the finite-sample term in lemma 2.4, we estimate empirical W_2 between two i.i.d. samples from the same distribution using exact assignment in low and moderate dimension, then fit a log-log slope in n . We also check the Sinkhorn proxy in one dimension at a few fixed values of ε .

The numerical check is only meant to support the modeling choice behind lemma 2.4. In low dimension, the raw Wasserstein term is compatible with an $n^{-1/2}$ fluctuation scale, whereas higher-dimensional raw Wasserstein decays more slowly. For fixed ε , the Sinkhorn proxy acts as a measurement layer whose calibration constant changes with regularization and ambient geometry even when the horizon argument keeps the same formal shape.

Table 6: Empirical scaling check for the finite-sample term. The bounded-support W_2 experiment is close to the expected $n^{-1/2}$ rate in low dimension, and the slope degrades as intrinsic dimension increases. The Sinkhorn proxy is reported at fixed ε , where the measured slope and constant both vary with regularization.

Setting	Estimated slope	Comment
$W_2, d = 1$	-0.48	close to $-1/2$
$W_2, d = 2$	-0.41	slower decay
$W_2, d = 4$	-0.25	slower decay
$W_2, d = 5$	-0.24	slower decay
Sinkhorn $\varepsilon = 0.50$	-0.21	proxy constant changes
Sinkhorn $\varepsilon = 0.10$	-0.53	proxy constant changes
Sinkhorn $\varepsilon = 0.05$	-0.67	proxy constant changes

E Online Drift Estimation and Sensitivity

The dynamic horizon regulator uses a short-window drift proxy built from block differences. For block length d , let $\bar{x}_t^{(d)}$ denote the mean of the most recent block and $\bar{x}_{t-d}^{(d)}$ the mean of the previous block. The instantaneous drift proxy is

$$\Delta_t = \frac{|\bar{x}_t^{(d)} - \bar{x}_{t-d}^{(d)}|}{d},$$

and the smoothed estimate follows the EMA recursion

$$\hat{\zeta}_t = \alpha \Delta_t + (1 - \alpha) \hat{\zeta}_{t-1}.$$

The regulator then sets

$$n_t^* = \text{clip} \left(\left(\frac{C_K}{\max(\hat{\zeta}_t, 10^{-9})} \right)^{2/3}, n_{\min}, n_{\max} \right).$$

Table 7 reports the sensitivity sweep on the alternating-timescales benchmark. Larger α makes the horizon react faster to bursts, which lowers contraction regret but leaves expansion regret larger; the expansion-to-contraction ratio stays above one throughout.

The concrete estimator used in the benchmarks and the accompanying empirical sweep show that the asymmetry persists as the drift filter is retuned.

E.1 Calibration Cost of C_K Misestimation

Let the regulator use an estimated constant $\hat{C}_K = rC_K$ while the true tracking law is governed by C_K . Since the selected horizon is $n^*(\hat{C}_K) = (\hat{C}_K/\zeta)^{2/3}$, the excess tracking error relative to the oracle choice is

Proposition E.1 (Calibration Cost of C_K Misestimation).

$$\Delta E(r) = \mathcal{E}(n^*(\hat{C}_K); C_K, \zeta) - \mathcal{E}(n^*(C_K); C_K, \zeta) = C_K^{2/3} \zeta^{1/3} \left(r^{-1/3} + \frac{1}{2} r^{2/3} - \frac{3}{2} \right).$$

Table 7: Drift-EMA sensitivity on the alternating-timescales benchmark. The ratio remains above one across the sweep, so re-expansion is slower than contraction for every tested α .

α	Contraction	Expansion	Ratio	Tail MAE
0.01	65.73	234.18	3.56	0.0938
0.02	68.51	241.03	3.52	0.0923
0.05	67.10	248.12	3.70	0.0925
0.1	61.86	251.57	4.07	0.0934
0.2	56.48	253.70	4.49	0.0940

Moreover, for $r = 1 + \delta$ with $|\delta| \ll 1$,

$$\Delta E(r) = \frac{1}{9} E_{\min} \delta^2 + O(\delta^3),$$

so small calibration errors are locally quadratic and symmetric to second order.

Table 8: Local calibration-cost sensitivity for the horizon constant. The excess error is normalized by $E_{\min}$, so the table is scale-free.

$r = \hat{C}_K / C_K$	Exact $\Delta E / E_{\min}$	Comment
0.90	0.0018	modest underestimation already moves the horizon
0.95	0.0004	local error is very small near the optimum
1.00	0.0000	oracle calibration
1.05	0.0004	local error is very small near the optimum
1.10	0.0015	moderate overestimation remains mild
1.25	0.0085	larger bias becomes visible
1.50	0.0288	large bias moves the operating point

The practical implication is simple: moderate calibration error in C_K is benign to first order, but large systematic bias still moves the operating point materially. The calibration-prefix machinery in Appendix A and the block-size sensitivity table below quantify how this shows up in practice. A short numerical check of the proposition is generated by `code/generate_zeta_support.py`.

E.2 Calibration Sensitivity Summary

Two short sensitivity checks are enough to capture the practical behavior of the regulator. First, the Sinkhorn regularization level changes the constant but not the exponent: across the tested values, the observed $\zeta^{1/3}$ slope stays at approximately 0.333, while the implied C_K moves from about σ at large ε to roughly 3.3σ at small ε . Second, the block length must not be too small: $d = 10$ produces a very large

positive bias in the drift proxy, $d = 50$ is already usable, and $d = 100$ is effectively stable.

The practical behavior is governed by two sensitivities. First, the Sinkhorn regularization level changes the calibration constant C_K without altering the cube-root scaling argument. Second, the block length must not be too small: overly short blocks make the drift proxy noisy and bias $\hat{\zeta}_t$ upward, which in turn pushes the regulator toward under-memory.

E.3 Robustness of the EMA Drift Proxy

Let the block-difference proxy satisfy $\Delta_t = \zeta_t + \xi_t$, where ξ_t is a zero-mean noise term with $\text{Var}(\xi_t) \leq \sigma_\Delta^2$ and the EMA recursion is

$$\hat{\zeta}_t = \alpha \Delta_t + (1 - \alpha) \hat{\zeta}_{t-1}.$$

If $\zeta_t \equiv \zeta$ is constant, then

$$\mathbb{E}[\hat{\zeta}_t] = \zeta + (1 - \alpha)^t (\hat{\zeta}_0 - \zeta), \quad \text{Var}(\hat{\zeta}_t) \leq \frac{\alpha}{2 - \alpha} \sigma_\Delta^2.$$

So smaller α reduces variance but increases lag, while larger α reacts faster but amplifies noise.

Since the regulator uses

$$n_t^* = \left(\frac{C_K}{\hat{\zeta}_t} \right)^{2/3},$$

the relative horizon error is first-order sensitive to relative drift-estimation error:

$$\frac{\delta n_t^*}{n^*} \approx -\frac{2}{3} \frac{\delta \hat{\zeta}_t}{\zeta}.$$

Thus persistent overestimation of ζ induces under-memory, and persistent underestimation induces over-memory. If the proxy has multiplicative bias $\mathbb{E}[\Delta_t] = (1 + b)\zeta$ with $b > 0$, then asymptotically

$$\frac{\hat{n}^*}{n^*} = (1 + b)^{-2/3},$$

so a 10% overestimate of drift shrinks the target horizon by about 6.5%.

For persistent overestimation, the same calculation gives systematic under-memory. For example, with a 10% positive bias the asymptotic horizon ratio is $1.1^{-2/3} \approx 0.94$, and with a 25% bias it is about 0.86.

E.4 Prefix-Validated Auto-Calibration

The operational fix is to choose (d, α) on a short calibration prefix by minimizing an *observable* prequential score over a small grid. For each candidate pair (d, α) , run the regulated estimator on the prefix, generate the one-step-ahead forecast $\hat{x}_{t|t-1}^{(d, \alpha)}$, and score it by

$$S(d, \alpha) = \frac{1}{T_0 - 1} \sum_{t=1}^{T_0-1} (x_{t+1} - \hat{x}_{t|t-1}^{(d, \alpha)})^2.$$

Table 9: EMA robustness check for a constant drift proxy with additive noise. The empirical mean matches the target drift, the empirical variance matches the closed-form $\alpha/(2 - \alpha)$ law, and the implied horizon distortion stays close to one when the proxy is unbiased.

α	$\mathbb{E}[\hat{\zeta}_t]$	$\text{Var}(\hat{\zeta}_t)$	Theory	$\mathbb{E}[\hat{n}^*]/n^*$
0.01	1.0002	0.00032	0.00031	1.0000
0.02	0.9999	0.00062	0.00063	1.0004
0.05	1.0004	0.00161	0.00160	1.0006
0.1	1.0001	0.00330	0.00329	1.0018
0.2	1.0000	0.00693	0.00694	1.0039

Table 10: Prefix-validated auto-calibration on the synthetic benchmarks. The calibration prefix chooses (d, α) by observable one-step-ahead prequential error, then the selected parameters are run on the full stream.

Scenario	Mean α	Mean d	Tail MAE (default)	Tail MAE (calibrated)
Constant drift	0.07	83.8	0.0978	0.1015
Alternating drift	0.03	93.8	0.0925	0.0952




Figure 13: Prefix-validated self-calibration of the drift proxy. The selected (d, α) shifts the operating point of the regulator; on these scenarios it does not improve tail MAE relative to the fixed default.

The selected pair minimizes this prequential MSE on the prefix and is then used for the full run. This is implementable online because it only uses observed data and the regulator’s own forecasts.

The failure mode is also visible: if the proxy is overreactive, the selected window contracts too aggressively and under-memory appears in the tail; if it is too inertial, the horizon lags the drift and over-memory follows. The calibration rule makes that sensitivity explicit by choosing the grid point that best predicts the next-step observations on the prefix.

The implemented self-calibrating rule is therefore simple: keep a small grid of block scales and EMA coefficients, score each candidate by prequential one-step

MSE on a short prefix, then deploy the best candidate for the remainder of the stream. For non-Lipschitz local behavior, the grid should include at least one short and one longer block scale so that bursty changes do not force a single timescale to carry the entire burden. This is an operational refinement, not a replacement for the fixed-regulator results above.

F Additional Public-Stream Summary

Bikes provides a second public-stream summary and shows that the same temporal-validity gap is not unique to one stream.

Table 11: Bikes public-stream summary. Precision is matched-warning precision against Page-Hinkley events; median lead is the median warning lead in steps.

Method	Precision	Median lead	Leads
EWMA	0.146	132.0	59
EWMA + UMR	0.146	132.0	59
Window Dilemma	0.127	108.5	76
Window Dilemma + UMR	0.128	125.0	76
MELO-style	0.141	147.5	76
MELO-style + UMR	0.129	147.0	71
ADWIN	0.526	92.0	251
ADWIN + UMR	0.520	82.0	141

The qualitative picture is the same as in ELEC2: coarse horizon adaptation still matters on public data, but the payoff remains backend-dependent.

G Controlled Downstream Check

The controlled synthetic downstream benchmark focuses on the phase-2 cap-only window: UMR has already contracted, while ADWIN remains silent.

Table 12: Controlled downstream benchmark summary. The reported delta is UMR minus Fixed-400 within the cap-only interval, averaged across six seeds.

Method	Global acc.	Cap-only acc.	Δ cap-only	Mean cap-only width
Fixed-400	0.9304	0.9040	+0.00	293.3
Fixed-100	0.9598	0.9576	+5.36	293.3
ADWIN	0.9340	0.9040	+0.00	293.3
UMR	0.9415	0.9205	+1.66	293.3




Figure 14: Controlled downstream benchmark. Top: true drift rate. Middle: effective horizons and the cap-only interval. Bottom: rolling prequential accuracy. The representative seed shows a 3.12 point gain for UMR over Fixed-400 within the cap-only window.

References

- O. Deniz Akyildiz, Pierre Del Moral, and Joaquín Míguez. Gaussian entropic optimal transport: Schrödinger bridges and the sinkhorn algorithm. *CoRR*, abs/2412.18432, 2024. doi: 10.48550/arXiv.2412.18432.
- Omar Besbes, Yonatan Gur, and Assaf Zeevi. Non-stationary stochastic optimization. *Operations Research*, 63(5):1227–1244, 2015. doi: 10.1287/OPRE.2015.1408.
- Albert Bifet and Ricard Gavaldà. Learning from time-changing data with adaptive windowing. In *Proceedings of the 2007 SIAM International Conference on Data Mining (SDM)*, pages 443–448, 2007. doi: 10.1137/1.9781611972771.42.
- Wang Chi Cheung, David Simchi-Levi, and Ruihao Zhu. Learning to optimize under non-stationarity. In *Proceedings of the Twenty-Second International Conference on Artificial Intelligence and Statistics*, volume 89 of *Proceedings of Machine Learning Research*, pages 1079–1087. PMLR, 2019. URL <https://proceedings.mlr.press/v89/cheung19b.html>.
- João Gama, Pedro Medas, Gladys Castillo, and Pedro Rodrigues. Learning with drift detection. In *Advances in Artificial Intelligence — SBIA 2004*, volume 3171

- of *Lecture Notes in Computer Science*, pages 286–295. Springer, 2004. doi: 10.1007/978-3-540-28645-5_29.
- João Gama, Indrė Žliobaitė, Albert Bifet, Mykola Pechenizkiy, and Abdelhamid Bouchachia. A survey on concept drift adaptation. *ACM Computing Surveys*, 46(4):44:1–44:37, 2014. doi: 10.1145/2523813.
- Aude Genevay, Lénaïc Chizat, Francis Bach, Marco Cuturi, and Gabriel Peyré. Sample complexity of Sinkhorn divergences. In *Proceedings of the 22nd International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 89 of *Proceedings of Machine Learning Research*, pages 1574–1583, 2019.
- Ziv Goldfeld, Kengo Kato, Gabriel Rioux, and Ritwik Sadhu. Limit theorems for entropic optimal transport maps and the sinkhorn divergence. *Electronic Journal of Statistics*, 18(1):980–1041, 2024. doi: 10.1214/24-EJS2217.
- Fabian Hinder, Valerie Vaquet, and Barbara Hammer. The window dilemma: Why concept drift detection is ill-posed, 2023. Preprint.
- Andrew H. Jazwinski. *Stochastic Processes and Filtering Theory*, volume 64 of *Mathematics in Science and Engineering*. Academic Press, New York, 1970.
- Jiashuo Jiang, Xiaocheng Li, and Jiawei Zhang. Online stochastic optimization with wasserstein based non-stationarity, 2020. URL <https://arxiv.org/abs/2012.06961>.
- Sanjit Kaul, Roy Yates, and Marco Gruteser. Real-time status: How often should one update? In *Proceedings of IEEE INFOCOM*, pages 2731–2735, 2012. doi: 10.1109/INFCOM.2012.6195689.
- Dominic S. T. Keehan, Edward J. Anderson, and Wolfram Wiesemann. Don’t look back in anger: Wasserstein distributionally robust optimization with nonstationary data. *CoRR*, abs/2510.18566, 2025. doi: 10.48550/arXiv.2510.18566.
- Alessio Mazzetto and Eli Upfal. Nonparametric density estimation under distribution drift. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 24251–24270. PMLR, 2023. URL <https://proceedings.mlr.press/v202/mazzetto23a.html>.
- Jonathan Niles-Weed and Francis Bach. Sharp asymptotic and finite-sample rates of convergence of empirical measures in Wasserstein distance. *Bernoulli*, 25(4A): 2620–2648, 2019. doi: 10.3150/18-BEJ1077.
- Vinicius Parreira. Useful memory has a horizon: Structural limits of drift tracking. <https://github.com/jovemexausto/useful-memory-horizon>, 2026. Accessed May 2026.
- Kishansingh Rajput, Malachi Schram, Brian Sammulu, and Sen Lin. Uncertainty guided online ensemble for non-stationary data streams in fusion science, 2025. URL <https://arxiv.org/abs/2511.02092>.

- Zhiyong Wang, Jize Xie, Yi Chen, John C. S. Lui, and Dongruo Zhou. Variance-dependent regret bounds for non-stationary linear bandits, 2024. URL <https://arxiv.org/abs/2403.10732>.
- Roy D. Yates, Yin Sun, D. Richard Brown, Sanjit K. Kaul, Eytan Modiano, and Sennur Ulukus. Age of information: An introduction and survey. *IEEE Journal on Selected Areas in Communications*, 39(5):1183–1210, 2021. doi: 10.1109/JSAC.2021.3064980.
- Jianjun Yuan and Andrew G. Lamperski. Trading-off static and dynamic regret in online least-squares and beyond. In *The Thirty-Fourth AAAI Conference on Artificial Intelligence*, pages 6712–6719. AAAI Press, 2020. doi: 10.1609/AAAI.V34I04.6149.